

Subsymbolic AI

Ibrahim Nasser*

Last updated on May 19, 2026

Contents

1	Introduction	1
2	Probability Theory	3
3	Bayesian Networks	15
4	Decision Networks	21
5	Hidden Markov Models	28
6	Markov Decision Processes	41
7	Introduction to Machine Learning	56

Appendices **62**

A	Mathematical Vernacular and Induction	62
A.1	Mathematical Language	62
A.2	Unary Natural Numbers	67
B	Set Theory and Foundations of Mathematics	69
B.1	Sets	69
B.2	Relations	71
B.3	Functions	75
B.4	Equivalence Classes	81
B.5	Algebraic Structures	83
C	Topics in Computer Science	85
C.1	Graph Theory	85
C.2	Complexity Theory	88
C.3	Formal Languages and Grammars	92

*This document and its source files are licensed under Creative Commons Attribution–NonCommercial–NoDerivatives 4.0 International License (CC BY-NC-ND 4.0).

1 Introduction

Def 1.1. Artificial intelligence (AI): The capability of computational systems to perform tasks typically associated with human intelligence, such as learning, reasoning, problem-solving, perception, and decision-making

Def 1.2. Symbolic AI: A subfield of [Artificial Intelligence \(AI\)](#) based on the assumption that many aspects of intelligence can be achieved by the manipulation of symbols, combining them into meaningful structures (expressions) and manipulating them (using processes) to produce new expressions.

Def 1.3. Statistical AI: Remedies the two shortcomings of [symbolic AI](#) approaches: that all concepts represented by symbols are crisply defined, and that all aspects of the world are knowable/representable in principle. Statistical AI adopts sophisticated mathematical models of uncertainty and uses them to create more accurate world models and reason about them.

Def 1.4. Subsymbolic AI (a.k.a connectionism/neural AI): A subfield of [AI](#) that posits that intelligence is inherently tied to brains, where information is represented by a simple sequence pulses that are processed in parallel via simple calculations realized by neurons, and thus concentrates on neural computing.

Def 1.5. Embodied AI: Posits that intelligence cannot be achieved by [reasoning](#) about the state of the world ([symbolically](#), [statistically](#), or [sub-symbolically](#)), but must be embodied i.e. situated in the world, equipped with a "body" that can interact with it via sensors and actuators. Here, the main method for realizing intelligent behavior is by learning from the world.

Def 1.6. Reasoning: The process of producing valid arguments and predictive world models. There are three forms of reasoning:

- **deductive reasoning** to produce new knowledge from existing knowledge. *Example:* All humans are mortal. Socrates is a human. Therefore, Socrates is mortal.
- **inductive reasoning** to produce knowledge from perception. *Example:* The sun has risen every morning in recorded history. Therefore, the sun will rise tomorrow.
- **abductive reasoning** to produce explanations for observations and given knowledge. *Example:* The grass is wet this morning. If it rained last night, that would explain it. Therefore, it probably rained.

Def 1.7. Inference: The act or process of reaching a **conclusion** about something from known facts or evidence (jointly called **premises**)

Def 1.8. Formal Logic: The science of [deductively](#) valid inferences, meaning arguments whose [conclusions](#) necessarily follow from their [premises](#) by virtue of their form (their structure), regardless of the topic.

Def 1.9. Agent: An agent is a [structure](#) $A := \langle \mathcal{P}, \mathcal{A}, f \rangle$ where:

- $\mathcal{P}$ is a [set](#) of percepts
- $\mathcal{A}$ is a [set](#) of actions
- f is a [function](#) $f : \mathcal{P}^* \rightarrow \mathcal{A}$ that maps from percepts to actions.

In other words, an agent is anything that perceives its environment via sensors and acts on it with actuators.

Def 1.10. Performance Measure: A [function](#) that evaluates a sequence of environments.

Def 1.11. Rational: An [agent](#) is called **rational**, if it chooses whichever action maximizes the expected value of the [performance measure](#) given the percept sequence to date.

Def 1.12. PEAS: To design [rational agents](#), we must specify the PEAS components: [Performance Measure](#), Environment, Actuators, and Sensors.

Def 1.13. Environment Types: For an [agent](#) a we classify the environment e of a by its **type**, which is one of the following. We call e

1. *fully observable*, iff the a 's sensors give it access to the complete state of the environment at any point in time; otherwise, we call it *partially observable*.
2. *deterministic*, iff the next state of the environment is completely determined by the current state and a 's actions; otherwise, *stochastic*.
3. *episodic*, iff a 's experience is divided into atomic episodes, where it perceives and then performs a single action, and the next episode does not depend on previous ones. Otherwise, *sequential*.
4. *dynamic*, iff the environment can change without an action performed by a ; otherwise, *static*. If the environment does not change but a 's performance measure does, we call e *semidynamic*.

5. *discrete*, iff the **sets** of e 's state and a 's actions are countable; otherwise, *continuous*.

6. *single-agent*, iff only a acts on e ; otherwise *multi-agent*.

Def 1.14. Reflex: An **agent** $\langle \mathcal{P}, \mathcal{A}, f \rangle$ is called a **reflex agent**, iff it only takes the last percept into account when choosing an action, i.e. $f(p_1, \dots, p_k) = f(p_k) \forall p_1 \dots p_k \in \mathcal{P}$

Def 1.15. Model-Based Agent: A model-based agent $\langle \mathcal{P}, \mathcal{A}, \mathcal{S}, \mathcal{T}, s_0, S, a \rangle$ is an **agent** $\langle \mathcal{P}, \mathcal{A}, f \rangle$ whose actions depend on:

- a **world model**: a **set** $\mathcal{S}$ of possible *states*, and a start state $s_0 \in \mathcal{S}$.
- a **transition model** $\mathcal{T}$ that predicts a new state $\mathcal{T}(s, a)$ from a state s and an action a .
- a **sensor model** S that given a state s and a percept p determines a new state $S(s, p)$.
- an **action function** $a : \mathcal{S} \rightarrow \mathcal{A}$ that given a state $s \in \mathcal{S}$ selects the next action $a \in \mathcal{A}$.

Note

If the agent is in state s then it took action a and now perceives p , then the agent state will become $s' = S(p, \mathcal{T}(s, a))$ and accordingly take action $a' = a(s')$.

Def 1.16. Goal Based Agent: A goal based agent $\langle \mathcal{P}, \mathcal{A}, \mathcal{S}, \mathcal{T}, s_0, S, a, \mathcal{G} \rangle$ is a **model based agent** with an explicit set of goals. It consists of:

- a set of internal states $\mathcal{S}$ and an initial state $s_0 \in \mathcal{S}$,
- a transition model $\mathcal{T} : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$,
- a state update function $S : \mathcal{P} \times \mathcal{S} \rightarrow \mathcal{S}$,
- a set of goals $\mathcal{G}$,
- a goal conditioned action function $a : \mathcal{S} \times \mathcal{G} \rightarrow \mathcal{A}$ selecting an action given a state and the goals.

Def 1.17. Utility-Based Agent: A utility-based agent uses a world model along with a utility function that models its preferences among the states of that world. It chooses the action that maximizes the expected utility.

Def 1.18. Learning Agent: A learning agent is an **agent** that can improve its own behavior through experience. It is composed of four main components:

- the performance element, which selects actions based on the agent's current percepts and represents its existing knowledge or behavior,
- the learning element, which improves the performance element over time by analyzing feedback and identifying how to make better decisions,
- the critic, which evaluates the agent's performance according to a given performance standard and provides feedback to the learning element,
- the problem generator, which suggests new and informative actions or experiences that help the agent explore and learn more effectively.

Def 1.19. State Representation: We call a state representation **atomic**, iff it has no internal structure (black box). However, iff each state is characterized by attributes and their values then the representation is **factored**. A **structured** state representation is when include representations of objects, their properties and relationships.

2 Probability Theory

Def 2.1. Sample Space: We say that a **set** Ω is a **sample space** of a random experiment, iff Ω contains all possible outcomes of that experiment. Each individual outcome $\omega \in \Omega$ is called a **sample point** or an **elementary outcome**.

Example 2.1. Throwing two dice $\rightsquigarrow \Omega = \{\langle 1, 1 \rangle, \langle 1, 2 \rangle, \dots, \langle 6, 6 \rangle\}$ with $\Gamma(\Omega) = 36$

Example 2.2. Number of tosses until the first head $\rightsquigarrow \Omega = \{1, 2, 3, \dots\} = \mathbb{N}$ with $\Gamma(\Omega) = \aleph_0$

Example 2.3. A random point in the interval $[0, 1] \rightsquigarrow \Omega = [0, 1]$ with $\Gamma(\Omega) = \mathfrak{c}$

Def 2.2. Event: Given a sample space Ω , we say that E is an **event** iff $E \subseteq \Omega$. That is, an event is a **subset** of the sample space. We say an event E occurs iff the outcome ω of the experiment satisfies $\omega \in E$.

Note

Because **events** are **sets**, we use Set Theory to describe relations between them:

- **Complement** $\bar{E}$: the **event** E that does not occur
- **Union** $A \cup B$: the **event** that A or B or both occurs
- **Intersection** $A \cap B$: the **event** that both A and B occur
- **Disjoint** / Mutually Exclusive: two **events** A and B are **disjoint** iff $A \cap B = \emptyset$

Note

We usually borrow propositional logic syntax to express **union** and **intersection**:

$$P(A \cup B) \rightsquigarrow P(A \vee B) \quad P(A \cap B) \rightsquigarrow P(A \wedge B)$$

We also use “,” to express **joint probability**

$$P(A \cap B) \rightsquigarrow P(A \wedge B) \rightsquigarrow P(A, B)$$

Example 2.4. The **sample space** for rolling one die is $\Omega = \{1, 2, \dots, 6\}$, **table 1** shows some **events** of interest.

Event Description	Set Representation
A : Even Number	$\{2, 4, 6\}$
B : Prime Number	$\{2, 3, 5\}$
$A \cap B$ (Even and Prime)	$\{2\}$
$A \cup B$ (Even or Prime)	$\{2, 3, 4, 5, 6\}$
$\bar{A}$ (Not Even)	$\{1, 3, 5\}$

Table 1: Examples of Events

Def 2.3. σ -algebra: Given a **sample space** Ω , a collection of **subsets** $\mathcal{F} \subseteq \mathcal{P}(\Omega)$ is called a **σ -algebra** iff it satisfies the following three axioms:

1. **Non-emptiness:** The sample space is included: $\Omega \in \mathcal{F}$. (which implies $\emptyset \in \mathcal{F}$ via complementation)
2. **Closure under complementation:** If $A \in \mathcal{F}$, then $\bar{A} \in \mathcal{F}$.
3. **Closure under countable unions:** If $A_1, A_2, \dots \in \mathcal{F}$ is a sequence of events, then $\bigcup_{i=1}^{\infty} A_i \in \mathcal{F}$.

The pair $(\Omega, \mathcal{F})$ is called a **measurable space**.

Example 2.5. Let the **sample space** be $\Omega = \{1, 2, 3\}$.

1. Consider the collection $\mathcal{F} = \{\emptyset, \{1\}, \{2, 3\}, \Omega\}$. This is a **σ -algebra** because:

- $\Omega \in \mathcal{F}$ (and $\emptyset \in \mathcal{F}$).
- It is closed under complements: $\bar{\{1\}} = \{2, 3\} \in \mathcal{F}$ and $\bar{\emptyset} = \Omega \in \mathcal{F}$.
- It is closed under unions: $\{1\} \cup \{2, 3\} = \Omega \in \mathcal{F}$.

2. Consider the collection $\mathcal{G} = \{\emptyset, \Omega, \{1\}, \{2\}\}$. This fails to be a σ -algebra for two reasons:

- No closure under Complements: $\{1\} \in \mathcal{G}$, but its complement $\{\bar{1}\} = \{2, 3\}$ is not in $\mathcal{G}$.
- No closure under Unions: $\{1\} \in \mathcal{G}$ and $\{2\} \in \mathcal{G}$, but their union $\{1, 2\}$ is not in $\mathcal{G}$.

Def 2.4. Probability Measure: A **probability measure** P is a function $P : \mathcal{F} \rightarrow [0, 1]$ that assigns a real number to every **event** $A \in \mathcal{F}$, satisfying the following **Kolmogorov Axioms**:

1. **Non-negativity:** $P(A) \geq 0$ for all $A \in \mathcal{F}$.
2. **Normalization:** $P(\Omega) = 1$.
3. **Countable Additivity:** For any sequence of **disjoint events** $A_1, A_2, \dots \in \mathcal{F}$:

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i)$$

Def 2.5. Probability Space: A **probability space** is the triplet $\langle \Omega, \mathcal{F}, P \rangle$, where:

- Ω is the **sample space**.
- $\mathcal{F}$ is a σ -algebra on Ω .
- P is a **probability measure** defined on $(\Omega, \mathcal{F})$.

Example 2.6. For a fair 6-sided die roll, the **probability space** is defined as follows:

- $\Omega = \{1, 2, 3, 4, 5, 6\}$
- $\mathcal{F} = \mathcal{P}(\Omega)$ (the power set)
- $P(A) = \frac{\Gamma(A)}{\Gamma(\Omega)} = \frac{\Gamma(A)}{6}$

Note

From the **Kolmogorov Axioms**, we can derive the following fundamental properties of **probability measures**:

- **Complement Rule:** For any **event** A , $P(\bar{A}) = 1 - P(A)$.
- **Monotonicity:** If $A \subseteq B$, then $P(A) \leq P(B)$.
- **Inclusion-Exclusion Principle:** For any two **events** A and B :

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

- **Set Partition Identity:** For any two **events** A and B , B can be split into parts that are in A and parts that are not:

$$P(B) = P(A \cap B) + P(\bar{A} \cap B)$$

- **Impossible Event:** $P(\emptyset) = 0$.

Def 2.6. Random Variable: Given a **probability space** $\langle \Omega, \mathcal{F}, P \rangle$, a **random variable** X is a **function** $X : \Omega \rightarrow D$ that maps each **sample point** $\omega \in \Omega$ to a numerical value in a **set** $D \subseteq \mathbb{R}$. We call D the **domain** of X , denoted by $\text{dom}(X)$.

Example 2.7. In a coin flip experiment where $\Omega = \{\text{Head}, \text{Tail}\}$, we can define a **random variable** X such that:

- $X(\text{Head}) = 1$
- $X(\text{Tail}) = 0$

In this case, $\text{dom}(X) = \{0, 1\}$.

Def 2.7. Probability of a Random Variable: For a **random variable** X and a specific value $x \in \text{dom}(X)$, the probability that X takes the value x is defined as the measure of the **event** consisting of all outcomes ω that map to x :

$$P(X = x) := P(\{\omega \in \Omega \mid X(\omega) = x\})$$

Example 2.8. Consider rolling a fair die where $\Omega = \{1, 2, \dots, 6\}$. If I define a **random variable** Y as the square of the outcome, then the probability of the **event** $Y = 4$ is:

$$P(Y = 4) = P(\{\omega \in \Omega \mid \omega^2 = 4\}) = P(\{2\}) = \frac{1}{6}$$

Def 2.8. Probability of Negation: For a **random variable** X and $x \in \text{dom}(X)$, we define the probability of the logical negation as:

$$P(X \neq x) := P(\{\omega \in \Omega \mid X(\omega) \neq x\}) = 1 - P(X = x)$$

Def 2.9. Probability of Disjunction: For **random variables** X_1, X_2 , we define the probability of the logical disjunction ("or") as:

$$P(X_1 = x_1 \vee X_2 = x_2) := P(\{\omega \in \Omega \mid X_1(\omega) = x_1\} \cup \{\omega \in \Omega \mid X_2(\omega) = x_2\})$$

Def 2.10. Joint Probability: Given two **random variables** X and Y defined on the same **sample space** Ω , the **joint probability** of $x \in \text{dom}(X)$ and $y \in \text{dom}(Y)$ is defined as the **probability** of the **intersection** of their individual **events**:

$$P(X = x, Y = y) := P(\{\omega \in \Omega \mid X(\omega) = x\} \cap \{\omega \in \Omega \mid Y(\omega) = y\})$$

Example 2.9. In a roll of two dice where X is the first die and Y is the second die: The joint probability $P(X = 3, Y = 4)$ is the probability of the **sample point** $(3, 4) \in \Omega$:

$$P(X = 3, Y = 4) := P(\{(3, 4)\}) := \frac{1}{36}$$

Example 2.10. To find the probability that the sum of two dice is 7 but neither is 3, we define the following **random variables** on the **sample space** $\Omega = \{1, \dots, 6\}^2$:

- X_1, X_2 : outcomes of the first and second die respectively.
- S : the sum, where $S(\omega) = X_1(\omega) + X_2(\omega)$.

We are interested in the **event** E :

$$E = \{\omega \in \Omega \mid S(\omega) = 7\} \cap \{\omega \in \Omega \mid X_1(\omega) \neq 3\} \cap \{\omega \in \Omega \mid X_2(\omega) \neq 3\}$$

Step 1: Identify outcomes where $S = 7$:

$$A = \{(1, 6), (2, 5), (3, 4), (4, 3), (5, 2), (6, 1)\}$$

Step 2: Apply the constraints $X_1 \neq 3$ and $X_2 \neq 3$:

- $\omega \in A$ and $X_1(\omega) = 3 \implies \omega = (3, 4)$
- $\omega \in A$ and $X_2(\omega) = 3 \implies \omega = (4, 3)$

Step 3: Calculate the **cardinality** of the filtered set:

$$E = A \setminus \{(3, 4), (4, 3)\} = \{(1, 6), (2, 5), (5, 2), (6, 1)\} \implies \Gamma(E) = 4$$

Step 4: Use the **probability measure**:

$$P(E) = \frac{\Gamma(E)}{\Gamma(\Omega)} = \frac{4}{36} = \frac{1}{9}$$

Def 2.11. Marginalization: For two **random variables** X and Y , the **marginal probability** of X is calculated by summing the **joint probabilities** over all $y \in \text{dom}(Y)$:

$$P(X = x) := \sum_{y \in \text{dom}(Y)} P(X = x, Y = y)$$

This is a direct application of the axiom of countable additivity to the partition of the **event** $\{X = x\}$ into smaller, mutually exclusive joint events.

Example 2.11. Suppose I have a **joint probability** table for two **random variables** X and Y representing a simplified weather/activity model:

	$Y = 0$ (Stay)	$Y = 1$ (Walk)	Total
$X = 0$ (Rain)	0.3	0.1	$P(X = 0) = \mathbf{0.4}$
$X = 1$ (Sun)	0.1	0.5	$P(X = 1) = \mathbf{0.6}$

Table 2: Marginalization Example

To find $P(X = 0)$, I marginalize over Y :

$$P(X = 0) = P(X = 0, Y = 0) + P(X = 0, Y = 1) = 0.3 + 0.1 = 0.4$$

Def 2.12. Conditional Probability: Let A and B be **events** in a **probability space** where $P(B) > 0$. The **conditional probability** of A given B is defined as:

$$P(A | B) := \frac{P(A \cap B)}{P(B)}$$

In this context, we refer to:

- $P(A)$ as the **prior** probability.
- $P(B)$ as the **evidence**.
- $P(A | B)$ as the **posterior** probability.

Example 2.12. Suppose I roll a fair 6-sided die, so $\Omega = \{1, 2, 3, 4, 5, 6\}$. Let A be the **event** that the outcome is 2, and B be the event that the outcome is even.

- $P(A) = P(\{2\}) = 1/6$
- $P(B) = P(\{2, 4, 6\}) = 3/6 = 1/2$
- $P(A \cap B) = P(\{2\}) = 1/6$

The probability that I roll a 2 **given** I know the result is even is:

$$P(A | B) = \frac{P(A \cap B)}{P(B)} = \frac{1/6}{1/2} = \frac{1}{3}$$

Theorem 2.1. $P(A | C) = 1 - P(\bar{A} | C)$

Proof.

1. Partition: $C = (A \cap C) \cup (\bar{A} \cap C)$.
2. Note that $(A \cap C)$ and $(\bar{A} \cap C)$ are **disjoint** because $A \cap \bar{A} = \emptyset$.
3. Apply the **Axiom of Additivity**: $P(C) = P(A \cap C) + P(\bar{A} \cap C)$.
4. Rearrange: $P(\bar{A} \cap C) = P(C) - P(A \cap C)$.
5. **Conditional probability**:

$$P(\bar{A} | C) = \frac{P(\bar{A} \cap C)}{P(C)}$$

6. Substitute:

$$P(\bar{A} | C) = \frac{P(C) - P(A \cap C)}{P(C)}$$

7. Simplify:

$$P(\bar{A} | C) = \frac{P(C)}{P(C)} - \frac{P(A \cap C)}{P(C)} = 1 - P(A | C)$$

8. Rearrange the terms:

$$P(A | C) = 1 - P(\bar{A} | C)$$

□

Theorem 2.2. $P(A | C) = P(A \cap B | C) + P(A \cap \bar{B} | C)$

Proof.

1. Partition: $(A \cap C) = (A \cap C \cap B) \cup (A \cap C \cap \bar{B})$
2. Note that $(A \cap C \cap B)$ and $(A \cap C \cap \bar{B})$ are **disjoint** because $B \cap \bar{B} = \emptyset$.
3. **Axiom of Additivity:** $P(A \cap C) = P(A \cap B \cap C) + P(A \cap \bar{B} \cap C)$
4. **Conditional probability:**

$$P(A | C) = \frac{P(A \cap C)}{P(C)}$$

5. Substitute:

$$P(A | C) = \frac{P(A \cap B \cap C) + P(A \cap \bar{B} \cap C)}{P(C)}$$

6. Distribute:

$$P(A | C) = \frac{P(A \cap B \cap C)}{P(C)} + \frac{P(A \cap \bar{B} \cap C)}{P(C)}$$

7. Finally:

$$P(A | C) = P(A \cap B | C) + P(A \cap \bar{B} | C)$$

□

Theorem 2.3. $P(A \cup B | C) = P(A | C) + P(B | C) - P(A \cap B | C)$

Proof.

1. **Conditional probability**

$$P(A \cup B | C) = \frac{P((A \cup B) \cap C)}{P(C)}$$

2. Distribute: $(A \cup B) \cap C = (A \cap C) \cup (B \cap C)$
3. Inclusion-Exclusion:

$$P((A \cap C) \cup (B \cap C)) = P(A \cap C) + P(B \cap C) - P((A \cap C) \cap (B \cap C))$$

4. Simplify: $(A \cap C) \cap (B \cap C) = A \cap B \cap C$
5. Substitute:

$$P(A \cup B | C) = \frac{P(A \cap C) + P(B \cap C) - P(A \cap B \cap C)}{P(C)}$$

6. Simplify:

$$P(A \cup B | C) = \frac{P(A \cap C)}{P(C)} + \frac{P(B \cap C)}{P(C)} - \frac{P(A \cap B \cap C)}{P(C)}$$

7. Finally:

$$P(A \cup B | C) = P(A | C) + P(B | C) - P(A \cap B | C)$$

□

Theorem 2.4.

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$

(Bayes' Theorem)

Proof.

1. **conditional probability:**

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

2. Multiply both sides by $P(B)$ to express the **joint probability**:

$$P(A \cap B) = P(A | B) \cdot P(B)$$

3. **conditional probability**:

$$P(B | A) = \frac{P(B \cap A)}{P(A)}$$

4. Multiply both sides by $P(A)$ to get a second expression for the **joint probability**:

$$P(B \cap A) = P(B | A) \cdot P(A)$$

5. **intersection** is commutative $\rightsquigarrow P(A \cap B) = P(B \cap A)$

6. Equate the right-hand sides of the expressions from Step 2 and Step 4:

$$P(A | B) \cdot P(B) = P(B | A) \cdot P(A)$$

7. Divide both sides by $P(B)$ (assuming $P(B) > 0$) to isolate the **posterior** probability:

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$

□

Def 2.13. Bayesian Terminology: Within the context of **Bayes' Theorem**, we define the following nomenclature:

- **Posterior:** $P(A | B)$ is the probability of the hypothesis A after observing the evidence B .
- **Likelihood:** $P(B | A)$ is the probability that the evidence B would be observed if the hypothesis A were true.
- **Prior:** $P(A)$ is the probability of the hypothesis A before the evidence B is taken into account.
- **Evidence:** $P(B)$ is the total probability of observing the evidence under all possible hypotheses.

Def 2.14. The Product Rule: For any two **events** A and B where the **conditional probabilities** are well-defined ($P(A) > 0, P(B) > 0$), the probability of their **joint occurrence** is:

$$P(A \cap B) = P(A | B) \cdot P(B)$$

$$P(A \cap B) = P(B | A) \cdot P(A)$$

This rule establishes that the probability of "A and B" is the probability of one event, scaled by the probability of the other event occurring given that the first has occurred.

Example 2.13. Suppose I draw two cards from a deck without replacement. Let A be the event the first card is an Ace, and B be the event the second card is an Ace.

- $P(A) = 4/52$
- $P(B | A) = 3/51$ (since one Ace is gone)

The probability of drawing two Aces is:

$$P(A \cap B) = P(B | A) \cdot P(A) = \frac{3}{51} \cdot \frac{4}{52}$$

Def 2.15. Independence: Two **events** A and B are **independent** iff $P(A \cap B) = P(A) \cdot P(B)$

Theorem 2.5. If A and B are **independent**, then the **posterior** equals the **prior**:

$$P(A | B) = P(A) \quad \text{and} \quad P(B | A) = P(B)$$

Proof.

1. **Conditional probability:**

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

2. **Independence:**

$$P(A | B) = \frac{P(A) \cdot P(B)}{P(B)} = P(A)$$

3. The proof for $P(B | A) = P(B)$ follows identical logic by symmetry.

□

Def 2.16. Pairwise Independence: A collection of **events** $A_1, A_2, \dots, A_n$ is **pairwise independent** iff every pair of distinct events is **independent**. That is, for all $i \neq j$:

$$P(A_i \cap A_j) = P(A_i) \cdot P(A_j)$$

Def 2.17. Mutual Independence: A collection of **events** $A_1, A_2, \dots, A_n$ is **mutually independent** iff for every finite subset of indices $I \subseteq \{1, 2, \dots, n\}$, the probability of their **intersection** is the product of their individual probabilities:

$$P\left(\bigcap_{i \in I} A_i\right) = \prod_{i \in I} P(A_i)$$

Note

Mutual independence strictly implies **pairwise independence**, but **pairwise independence** does not guarantee **mutual independence**.

Def 2.18. Conditional Independence: Given events A, B, C with $P(C) \neq 0$, then A and B are **conditionally independent** given C iff $P(A | B \cap C) := P(A | C)$, or $P(B | A \cap C) := P(B | C)$:

Theorem 2.6. If A and B are **conditionally independent** given C , then:

$$P(A \cap B | C) := P(A | C) \cdot P(B | C)$$

Proof.

1. By **definition**:

$$P(A \cap B | C) = \frac{P(A \cap B \cap C)}{P(C)}$$

2. By the **product rule**:

$$P(A \cap B \cap C) = P(A | B \cap C) \cdot P(B \cap C)$$

3. Expanding $P(B \cap C)$ further:

$$P(A \cap B \cap C) = P(A | B \cap C) \cdot P(B | C) \cdot P(C)$$

4. Substituting Step 3 into Step 1:

$$P(A \cap B | C) = \frac{P(A | B \cap C) \cdot P(B | C) \cdot P(C)}{P(C)}$$

5. Simplifying:

$$P(A \cap B | C) = P(A | B \cap C) \cdot P(B | C)$$

6. Since A and B are **conditionally independent** given C , $P(A | B \cap C) = P(A | C)$, yielding:

$$P(A \cap B | C) = P(A | C) \cdot P(B | C)$$

□

Def 2.19. Probability Distribution: A **probability distribution** is a vector v with values in $[0, 1]$ such that $\sum_i v_i = 1$.

Note

The **probability distribution** of a **random variable** X (written as $\mathbf{P}(X)$) is a complete description of the **probabilities** $P(X = x)$ for all values $x \in \text{dom}(X)$. It completely characterizes the behavior of the random variable by mapping each possible outcome to its corresponding probability, effectively describing how the total probability of 1 (from the **Normalization axiom**) is distributed across the domain of X .

Def 2.20. Full Joint Probability Distribution (JPD): Given a set of **random variables** $X_1, X_2, \dots, X_n$, a **full joint probability distribution** $\mathbf{P}(X_1, X_2, \dots, X_n)$ is a **probability distribution** that assigns a probability to every possible, mutually exclusive **joint probability** assignment for all variables. It can be represented as an n -dimensional tensor where each entry $P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) \in [0, 1]$ corresponds to a single atomic event, and the sum over all possible combinations of values is 1:

$$\sum_{x_1 \in \text{dom}(X_1)} \cdots \sum_{x_n \in \text{dom}(X_n)} P(X_1 = x_1, \dots, X_n = x_n) = 1$$

Example 2.14. Consider two boolean **random variables**: W (Weather: Sunny=1, Rainy=0) and C (Commute: Walk=1, Drive=0). The **full joint probability distribution** is given by the table:

	$C = \text{Walk}$	$C = \text{Drive}$	Total
$W = \text{Sunny}$	0.6	0.1	0.7
$W = \text{Rainy}$	0.05	0.25	0.3
Total	0.65	0.35	1.0

Note that the sum of all internal cells ($0.6 + 0.1 + 0.05 + 0.25 = 1.0$) satisfies the normalization requirement.

Def 2.21. Conditional Probability Distribution (CPD): Given **random variables** X and Y , the **conditional probability distribution** $\mathbf{P}(X \mid Y)$ represents a complete table or matrix of **probabilities**, providing the distribution of X for every possible observed value $y \in \text{dom}(Y)$. It consists of the **conditional probabilities** $P(X = x \mid Y = y)$ for all $x \in \text{dom}(X)$ and all $y \in \text{dom}(Y)$. For any fixed y , the distribution over X satisfies the normalization constraint:

$$\sum_{x \in \text{dom}(X)} P(X = x \mid Y = y) = 1$$

Example 2.15. Continuing from the previous example, we can form the **conditional probability distribution** $\mathbf{P}(C \mid W)$ by dividing each joint probability by the marginal probability of W . Note that for any fixed value of W (each column), the probabilities sum to 1:

$\mathbf{P}(C \mid W)$	$W = \text{Sunny}$	$W = \text{Rainy}$
$C = \text{Walk}$	$0.6/0.7 = 6/7$	$0.05/0.3 = 1/6$
$C = \text{Drive}$	$0.1/0.7 = 1/7$	$0.25/0.3 = 5/6$
Total	1.0	1.0

Def 2.22. A **Naive Bayes Model** (Bayesian Classifier) consists of:

- **Random variables** $E_1 \cdots E_n$ (effects)
- **Random variable** C (the cause)
- The **probability distribution** $\mathbf{P}(C)$
- The **CPD** $\mathbf{P}(E_i \mid C)$

Where all $E_1, \dots, E_n$ are **conditionally independent** given C .

Lemma 2.7. Chain Rule:

$$\mathbf{P}(X_1, \dots, X_n) = \mathbf{P}(X_1 \mid X_2, \dots, X_n) \cdot \mathbf{P}(X_2 \mid X_3, \dots, X_n) \cdots \mathbf{P}(X_{n-1} \mid X_n) \cdot \mathbf{P}(X_n)$$

Proof. via generalization of the **product rule** for e.g. $X_1, \dots, X_4$:

- $\mathbf{P}(X_1, X_2, X_3, X_4) = \mathbf{P}(X_1 \mid X_2, X_3, X_4) \cdot \mathbf{P}(X_2, X_3, X_4)$
- $\mathbf{P}(X_2, X_3, X_4) = \mathbf{P}(X_2 \mid X_3, X_4) \cdot \mathbf{P}(X_3, X_4)$
- $\mathbf{P}(X_3, X_4) = \mathbf{P}(X_3 \mid X_4) \cdot \mathbf{P}(X_4)$
- Hence, $\mathbf{P}(X_1, X_2, X_3, X_4) = \mathbf{P}(X_1 \mid X_2, X_3, X_4) \cdot \mathbf{P}(X_2 \mid X_3, X_4) \cdot \mathbf{P}(X_3 \mid X_4) \cdot \mathbf{P}(X_4)$

□

Question

Given a **Naive Bayes Model**, can we compute the Full **JPD** $\mathbf{P}(C, E_1, \dots, E_n)$ using only $\mathbf{P}(C)$ and $\mathbf{P}(E_i \mid C)$?

Answer

Yes. Because of the **conditional independence** assumption in a **Naive Bayes Model**, the **JPD** simplifies to:

$$\mathbf{P}(C, E_1, \dots, E_n) = \mathbf{P}(C) \prod_{i=1}^n \mathbf{P}(E_i \mid C)$$

Proof.

- **Chain rule:** $\mathbf{P}(E_1, \dots, E_n, C) = \mathbf{P}(E_1 \mid E_2, \dots, E_n, C) \cdots \mathbf{P}(E_n \mid C) \cdot \mathbf{P}(C)$
- All $E_1, \dots, E_n$ are **conditionally independent** given C
- $\mathbf{P}(E_1 \mid E_2, \dots, E_n, C) = \mathbf{P}(E_1 \mid C)$ and $\mathbf{P}(E_2 \mid E_3, \dots, E_n, C) = \mathbf{P}(E_2 \mid C)$ and so on $\cdots$
- Hence, $\mathbf{P}(C, E_1, \dots, E_n) = \mathbf{P}(C) \prod_{i=1}^n \mathbf{P}(E_i \mid C)$

□

Def 2.23. Marginalization Generalized: Given **random variables** $X_1, \dots, X_n$ and $Y_1, \dots, Y_m$, we have

$$\mathbf{P}(X_1, \dots, X_n) = \sum_{y_1 \in \text{dom}(Y_1), \dots, y_m \in \text{dom}(Y_m)} \mathbf{P}(X_1, \dots, X_n, Y_1 = y_1, \dots, Y_m = y_m)$$

Def 2.24. Given a **Naive Bayes Model** with a set of effects $E = \{E_1, \dots, E_n\}$, we refer to the **observable** effect as $O = \{O_1, \dots, O_{no}\}$ and the set of **unobservable** (unknown) effects as $U = \{U_1, \dots, U_{nu}\}$ where $E = O \cup U$

Notation Understanding

Assume we have two unknown effects ($nu = 2$): *Headache* (U_1) where $D_H = \{\text{yes}, \text{no}\}$ and *Fever* (U_2) where $D_F = \{\text{high}, \text{low}\}$

We define $D_u = \text{dom}(U_1) \times \dots \times \text{dom}(U_{nu}) = \{(y, h), (y, l), (n, h), (n, l)\}$

We can sum all [joint probabilities](#) of the unknown effects as :

$$\sum_{u \in D_u} P(U_1 = u_1, \dots, U_{nu} = u_{nu}) = 1$$

- For $u = (y, h)$, we have $P(U_1 = y, U_2 = h)$
- For $u = (y, l)$, we have $P(U_1 = y, U_2 = l)$
- For $u = (n, h)$, we have $P(U_1 = n, U_2 = h)$
- For $u = (n, l)$, we have $P(U_1 = n, U_2 = l)$

For a specific tuple $u = \langle u_1, u_2 \rangle \in D_u$, assuming *Headache* and *Fever* are [independent](#), the [joint probability](#) of $P(H = u_1, F = u_2)$ can be also computed from the product:

$$\prod_{j=1}^{nu} P(U_j = u_{u_j}) \quad \text{where } u_{u_j} \text{ means the value little } u \text{ at index } j \text{ given the tuple } u$$

Let's take the tuple $\langle n, h \rangle$ and we are iterating over $U = \{H(U_1), F(U_2)\}$

$$\prod_{j=1}^{nu} P(U_j = u_{u_j}) = P(H = n) \cdot P(F = h)$$

If we want to sum all the [joint probabilities](#) for all values, we write:

$$\sum_{u \in D_u} \prod_{j=1}^{nu} P(U_j = u_{u_j})$$

Which will evaluate to 1 too (assuming [independence](#)).

Let $D_u = \text{dom}(U_1) \times \dots \times \text{dom}(U_{nu})$, then:

$$\mathbf{P}(C \mid O_1, \dots, O_{no}) = \frac{\mathbf{P}(C, O_1, \dots, O_{no})}{\mathbf{P}(O_1, \dots, O_{no})}$$

We [marginalize](#) over the unknowns to compute the numerator:

$$\mathbf{P}(C, O_1, \dots, O_{no}) = \sum_{u \in D_u} \mathbf{P}(C, O_1, \dots, O_{no}, U_1 = u_1, \dots, U_{nu} = u_{nu})$$

$$\mathbf{P}(C, O_1, \dots, O_{no}) = \sum_{u \in D_u} \mathbf{P}(C) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C) \prod_{j=1}^{nu} \mathbf{P}(U_j = u_{u_j} \mid C)$$

$$\mathbf{P}(C, O_1, \dots, O_{no}) = \mathbf{P}(C) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C) \sum_{u \in D_u} \prod_{j=1}^{nu} \mathbf{P}(U_j = u_{u_j} \mid C)$$

Also, for the denominator, we have:

$$\mathbf{P}(O_1, \dots, O_{no}) = \sum_{c \in \text{dom}(C)} \sum_{u \in D_u} \mathbf{P}(O_1, \dots, O_{no}, C = c, U_1 = u_1, \dots, U_{nu} = u_{nu})$$

$$\mathbf{P}(O_1, \dots, O_{no}) = \sum_{c \in \text{dom}(C)} \sum_{u \in D_u} P(C = c) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C = c) \prod_{j=1}^{nu} P(U_j = u_{u_j} \mid C = c)$$

$$\mathbf{P}(O_1, \dots, O_{no}) = \sum_{c \in \text{dom}(C)} P(C = c) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C = c) \sum_{u \in D_u} \prod_{j=1}^{nu} P(U_j = u_{u_j} \mid C = c)$$

Therefore,

$$\mathbf{P}(C \mid O_1, \dots, O_{no}) = \frac{\mathbf{P}(C) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C) \sum_{u \in D_u} \prod_{j=1}^{nu} \mathbf{P}(U_j = u_{u_j} \mid C)}{\sum_{c \in \text{dom}(C)} P(C = c) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C = c) \sum_{u \in D_u} \prod_{j=1}^{nu} P(U_j = u_{u_j} \mid C = c)}$$

Note

$$\sum_{u \in D_u} \prod_{j=1}^{nu} P(U_j = u_{u_j} \mid C = c) = 1$$

Hence,

$$\mathbf{P}(C \mid O_1, \dots, O_{no}) = \frac{\mathbf{P}(C) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C)}{\sum_{c \in \text{dom}(C)} P(C = c) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C = c)}$$

Note

The denominator is

- the same for any given observations $O_1, \dots, O_{no}$ independent of the value of C
- the sum over all numerators

The denominator acts as a **normalizing constant**. It scales the product in the numerator so that the resulting values for $\mathbf{P}(C \mid O_1, \dots, O_{no})$ sum to 1. This process is called **Normalization**.

Def 2.25. Normalization: Given a vector $w := \langle w_1, \dots, w_k \rangle$ where each entry $w_i \in [0, 1]$ and $\sum_{i=1}^k w_i \leq 1$, we define a **function** $\alpha(w)$ that normalizes the vector by dividing each entry by the sum of all its components.

$$(\alpha(w))_i = \frac{w_i}{\sum_{j=1}^k w_j}$$

By construction, $\sum_{i=1}^k (\alpha(w))_i = 1$, which ensures that $\alpha(w)$ is a valid **probability distribution**.

Inference Theorem in a Naive Bayes Model

$$\mathbf{P}(C \mid O_1, \dots, O_{no}) = \alpha \left(\mathbf{P}(C) \prod_{i=1}^{no} \mathbf{P}(O_i \mid C) \right)$$

Example 2.16. Consider we are modeling *cavity* (C) and *toothache* (T) with Boolean domains $\{T, F\}$ and we know the following probabilities:

- $P(C = T) = 0.1 \quad P(T = T \mid C = T) = 0.8 \quad P(T = T \mid C = F) = 0.05$

What is the probability distribution for a patient having *cavity* given they have *toothache*?

$$\begin{aligned} \mathbf{P}(C \mid t = T) &:= \alpha(\mathbf{P}(C) \cdot \mathbf{P}(T = T \mid C)) \\ \mathbf{P}(C) &:= \langle P(C = T), P(C = F) \rangle \\ \mathbf{P}(T = T \mid C) &:= \langle P(T = T \mid C = T), P(T = T \mid C = F) \rangle \\ \mathbf{P}(C \mid t = T) &:= \alpha(\langle P(C = T), P(C = F) \rangle \cdot \langle P(T = T \mid C = T), P(T = T \mid C = F) \rangle) \\ &:= \alpha(\langle P(C = T) \cdot P(T = T \mid C = T), P(C = F) \cdot P(T = T \mid C = F) \rangle) \\ &:= \alpha(\langle 0.1 \cdot 0.8, 1 - P(C = T) \cdot 0.05 \rangle) \\ &:= \alpha(\langle 0.1 \cdot 0.8, 0.9 \cdot 0.05 \rangle) \\ &:= \alpha(\langle 0.08, 0.045 \rangle) \\ &:= \langle 0.64, 0.36 \rangle \end{aligned}$$

Theorem 2.8. Let $Q, E_1, \dots, E_{ne}, U_1, \dots, U_{nu}$ be **random variables** and $D_u = \text{dom}(U_1) \times \dots \times \text{dom}(U_{nu})$. Then:

$$\mathbf{P}(Q \mid E_1 = e_1, \dots, E_{ne} = e_{ne}) = \alpha \left(\sum_{u \in D_u} \mathbf{P}(Q, E_1 = e_1, \dots, E_{ne} = e_{ne}, U_1 = u_1, \dots, U_{nu} = u_{nu}) \right)$$

We call Q the **query variable**, $E_1, \dots, E_{ne}$ the **evidence**, and $U_1, \dots, U_{nu}$ the **unknown** (or **hidden**) **variables**, and computing a **conditional probability** this way **enumeration**

Note

General Law of Inference by Enumeration: "To find out what you care about (Q), you must look at every possible world (U) that could exist alongside what you already saw (E).” concretely:

- sum over all relevant entries of the **JPD** of the variables, and
- normalize the results to yield a **probability distribution**

3 Bayesian Networks

Def 3.1. A **Bayesian Network** consists of

1. a **DAG** $\langle \mathcal{X}, E \rangle$ of **random variables** $\mathcal{X} = \{X_1, \dots, X_n\}$, and
2. a **CPD** $\mathbf{P}(X_i \mid \text{Parents}(X_i))$ for every $X_i \in \mathcal{X}$,


such that every X_i is **conditionally independent** of any conjunctions of **non-descendants** of X_i given $\text{Parents}(X_i)$

Note

Roots of the **DAG** are **conditionally independent** given the empty set $\emptyset$; i.e. they are **independent**. This follows directly from the definition:

- By definition, a root node X_i has no parents, meaning $\text{Parents}(X_i) = \emptyset$.
- Any other root node X_j is a **non-descendant** of X_i (since there are no directed paths between roots).
- By the definition of the **Bayesian Network**, X_i is conditionally independent of its non-descendant X_j given its parents ($\emptyset$).
- Conditional independence given the empty set is exactly equivalent to marginal **independence**.

Example 3.1. Consider a simple **Bayesian Network** modeling a security system with variables: B (Burglar), E (Earthquake), A (Alarm rings), and C (Call Police).



- Variables $\mathcal{X} = \{B, E, A, C\}$
- Edges $E = \{(B, A), (E, A), (A, C)\}$
- **CPDs**: $\mathbf{P}(B)$, $\mathbf{P}(E)$, $\mathbf{P}(A \mid B, E)$, and $\mathbf{P}(C \mid A)$

By definition, each node is **conditionally independent** of its **non-descendants** given its parents. Let's analyze each node:

- **Node B**: $\text{Parents}(B) = \emptyset$, $\text{NonDesc}(B) = \{E\}$.
Property: B is independent of E given $\emptyset \implies \mathbf{P}(B \mid E) = \mathbf{P}(B)$.
- **Node E**: $\text{Parents}(E) = \emptyset$, $\text{NonDesc}(E) = \{B\}$.
Property: E is independent of B given $\emptyset \implies \mathbf{P}(E \mid B) = \mathbf{P}(E)$.
- **Node A**: $\text{Parents}(A) = \{B, E\}$, $\text{NonDesc}(A) = \emptyset$.
Property: Vacuously true (no non-descendants exists to be independent from).
- **Node C**: $\text{Parents}(C) = \{A\}$, $\text{NonDesc}(C) = \{B, E\}$.
Property: C is independent of $\{B, E\}$ given $A \implies \mathbf{P}(C \mid A, B, E) = \mathbf{P}(C \mid A)$.
(Intuition: Once we know the alarm rang, knowing exactly **why** it rang doesn't give us any extra information about whether the police will be called).

Theorem 3.1. The **JPD** of a **Bayesian Network** $\langle \mathcal{X}, E \rangle$ is given by:

$$\mathbf{P}(X_1, \dots, X_n) = \prod_{X_i \in \mathcal{X}} \mathbf{P}(X_i \mid \text{Parents}(X_i))$$

Proof.

1. Let $X_1, X_2, \dots, X_n$ be a **topological sort** (an ordering) of the variables in the **DAG**. This ordering guarantees that if $X_i \rightarrow X_j$ is an edge (meaning X_i is a parent of X_j), then $i < j$.
2. Because joint probability is commutative, we can examine the variables in the **JPD** as $\mathbf{P}(X_n, \dots, X_1)$. Applying the **Chain Rule**, we get:

$$\begin{aligned} \mathbf{P}(X_n, \dots, X_1) &= \mathbf{P}(X_n \mid X_{n-1}, \dots, X_1) \cdot \mathbf{P}(X_{n-1} \mid X_{n-2}, \dots, X_1) \cdots \mathbf{P}(X_2 \mid X_1) \cdot \mathbf{P}(X_1) \\ &= \prod_{i=1}^n \mathbf{P}(X_i \mid X_{i-1}, \dots, X_1) \end{aligned}$$

3. Because we used a topological sort, all variables $X_{i-1}, \dots, X_1$ appear before X_i in the ordering. This guarantees they must be **non-descendants** of X_i (since descendants can only appear *after* X_i).
4. Thus, the conditioning set $\{X_{i-1}, \dots, X_1\}$ consists exactly of $\text{Parents}(X_i)$ plus some other non-descendants of X_i .
5. By the definition of **Bayesian Networks**, X_i is **conditionally independent** of its **non-descendants** given its parents. Therefore, the extra **non-descendants** drop out of the probability:

$$\mathbf{P}(X_i \mid X_{i-1}, \dots, X_1) = \mathbf{P}(X_i \mid \text{Parents}(X_i))$$

6. Substituting this back into the **Chain Rule** from Step 2 yields:

$$\mathbf{P}(X_n, \dots, X_1) = \prod_{i=1}^n \mathbf{P}(X_i \mid \text{Parents}(X_i))$$

By commutativity of the joint probability, this is equivalent to $\mathbf{P}(X_1, \dots, X_n)$. Validating the theorem. □

Def 3.2. Given **random variables** $X_1, \dots, X_n$ with **finite domains** $D_1, \dots, D_n$, the **size** of the **Bayesian Network** $\mathcal{B} = \langle \{X_1, \dots, X_n\}, E \rangle$ is defined as:

$$\text{size}(\mathcal{B}) := \sum_{i=1}^n |D_i| \cdot \left(\prod_{X_j \in \text{Parents}(X_i)} |D_j| \right)$$

Note

- $\text{size}(\mathcal{B}) \triangleq$ The total number of entries in the **CPD**.
- Smaller **Bayesian Networks** need to estimate less probabilities $\rightsquigarrow$ more efficient inference.
- Explicit Full **JPD** has size $\prod_{i=1}^n |D_i|$
- If $|\text{Parents}(X_i)| \leq k$ for every X_i , and $D_{\max}$ is the largest **random variable domain**, then $\text{size}(\mathcal{B}) \leq n |D_{\max}|^{k+1}$
- In the worst case, $\text{size}(\mathcal{B}) = \sum_{i=1}^n \prod_{j=1}^i |D_j|$, namely if every variable depends on all its predecessors in the chosen variable ordering.
- **Bayesian Networks** are compact if each variable is directly influenced only by few of its predecessor variables.

Example 3.2. Consider a **Bayesian Network** $\mathcal{B}$ with 4 boolean **random variables** (meaning their domain sizes are $|D_i| = 2$ for all i) arranged in a simple chain: $X_1 \rightarrow X_2 \rightarrow X_3 \rightarrow X_4$.

We execute the formula term by term:

- **Node** X_1 : No parents ($\text{Parents}(X_1) = \emptyset \implies$ independent product is 1).
Term expands to: $|D_1| \cdot 1 = 2 \cdot 1 = 2$
- **Node** X_2 : One parent ($\text{Parents}(X_2) = \{X_1\}$).
Term expands to: $|D_2| \cdot |D_1| = 2 \cdot 2 = 4$
- **Node** X_3 : One parent ($\text{Parents}(X_3) = \{X_2\}$).
Term expands to: $|D_3| \cdot |D_2| = 2 \cdot 2 = 4$
- **Node** X_4 : One parent ($\text{Parents}(X_4) = \{X_3\}$).
Term expands to: $|D_4| \cdot |D_3| = 2 \cdot 2 = 4$

Summing these up gives the total size of the probability tables to be stored:

$$\text{size}(\mathcal{B}) = 2 + 4 + 4 + 4 = 14 \text{ probability entries}$$

We can verify the efficiency by comparing it to an explicit Full **JPD**:

$$\text{Full JPD Size} = \prod_{i=1}^4 |D_i| = 2 \cdot 2 \cdot 2 \cdot 2 = 16 \text{ entries}$$

Even on just 4 variables, the conditional independencies in the **Bayesian Network** make it tangibly more compact.

Upper Bound Analysis

In our chain, no node has more than 1 parent, so the maximum in-degree is $k = 1$. The largest domain is $D_{\max} = 2$ (because all variables are boolean). For our $n = 4$ nodes, the sparsity upper bound formula yields:

$$\text{size}(\mathcal{B}) \leq n \cdot |D_{\max}|^{k+1} = 4 \cdot 2^{1+1} = 4 \cdot 4 = 16$$

Our actual calculated size of 14 is indeed strictly bounded by 16.

To see the real power of this bound, imagine extending this exact chain to $n = 100$ boolean variables. The Full JPD would require 2^{100} entries (a computational impossibility). In contrast, our [Bayesian Network](#) upper bound guarantees the memory footprint is simply $\leq 100 \cdot 2^2 = 400$ entries. This demonstrates why bounding the maximum number of parents k is critical for scalability!

Question

How to keep our [Bayesian Network](#) small?

Answer

- Reduce the number of edges $\rightsquigarrow$ Order the variables to allow for exploiting [conditional independence](#) (causes before effects), or
- Represent the [CPD](#) efficiently:
 - For a Boolean [random variable](#) X , we only need to store $P(X = T \mid \text{Parents}(X))$
 - Introduce different kinds of nodes exploiting domain knowledge.

Def 3.3. BN Construction Algorithm:

1. Initialize $BN = \langle \{X_1, \dots, X_n\}, E \rangle$ where $E = \emptyset$
2. Fix any variable ordering, $X_1, \dots, X_n$
3. **for** $i := 1, \dots, n$ **do**:
 - (a) Choose a minimal set $\text{Parents}(X_i) \subseteq \{X_1, \dots, X_{i-1}\}$ such that:

$$\mathbf{P}(X_i \mid X_{i-1}, \dots, X_1) = \mathbf{P}(X_i \mid \text{Parents}(X_i))$$

- (b) For each $X_j \in \text{Parents}(X_i)$, insert $\langle X_j, X_i \rangle$ into E
- (c) Associate X_i with $\mathbf{P}(X_i \mid \text{Parents}(X_i))$

Note

The size of the resulting [Bayesian Network](#) depends on the chosen **variable ordering**. The size of a [Bayesian Network](#) is not a fixed property (it depends on the skill of the designer)

Rule of Thumb

Order *causes* before *effects*

Def 3.4. A node X in a [Bayesian Network](#) is called **deterministic**, if its value is completely determined by the values of $\text{Parents}(X)$


Example 3.3. The sum of two dice throws S is entirely determined by the values of the two dice **First** and **Second**

Note

- [Deterministic](#) nodes model direct, causal relationships.
- If X is [deterministic](#), then $P(X \mid \text{Parents}(X)) \in \{0, 1\}$
- We can replace the [CPD](#) $\mathbf{P}(X \mid \text{Parents}(X))$ by a boolean function

Noisy Nodes

Sometimes, values of nodes are *almost deterministic*. For example, assume we have a network that model all causes of fever:



If there is a fever, then one of them (at least) must be the cause, but none of them necessarily cause a fever. We say that the causal relation between parent and child is **inhibited**. We can model **inhibition** by individual inhibition factors q_d

Def 3.5. The **CPD** of a **noisy disjunction node** X with $\text{Parents}(X) = \{X_1, \dots, X_n\}$ in a **Bayesian Network** is given by:

$$P(X \mid X_1, \dots, X_n) = 1 - \left(\prod_{\{j \mid X_j = T\}} q_j \right)$$

where the q_i are the **inhibition factors** of $X_i \in \text{Parents}(X)$, defined as:

$$q_i := P(\neg X \mid \neg X_1, \dots, \neg X_{i-1}, X_i, \neg X_{i+1}, \dots, \neg X_n)$$

Note

Instead of a distribution with 2^k parameters, we only need k parameters.

Example 3.4. Assume the following **inhibition factors**:

$$\begin{aligned} q_{\text{cold}} &= P(\neg \text{fever} \mid \text{cold}, \neg \text{flu}, \neg \text{malaria}) = 0.6 \\ q_{\text{flu}} &= P(\neg \text{fever} \mid \neg \text{cold}, \text{flu}, \neg \text{malaria}) = 0.2 \\ q_{\text{malaria}} &= P(\neg \text{fever} \mid \neg \text{cold}, \neg \text{flu}, \text{malaria}) = 0.1 \end{aligned}$$

If we model *fever* as a **noisy disjunction node**, then the general rule for the **CPD**

$$P(X_i \mid \text{Parents}(X_i)) = \prod_{\{j \mid X_j = T\}} q_j$$

gives the following table:

Cold	Flu	Malaria	$P(\text{Fever})$	$P(\neg \text{Fever})$
F	F	F	0.0	1.0
F	F	T	0.9	0.1
F	T	F	0.8	0.2
F	T	T	0.98	$0.02 = 0.2 \cdot 0.1$
T	F	F	0.4	0.6
T	F	T	0.94	$0.06 = 0.6 \cdot 0.1$
T	T	F	0.88	$0.12 = 0.6 \cdot 0.2$
T	T	T	0.988	$0.012 = 0.6 \cdot 0.2 \cdot 0.1$

Def 3.6. John, Mary and the Alarm: This is the well-known **Bayesian Network** example found in textbooks. It goes as follows:

- I got very valuable stuff at home. So I bought an alarm. Unfortunately, the alarm just rings at home, doesn't call me on my mobile.
- I've got two neighbors, Mary and John, who'll call me if they hear the alarm.
- The problem is that, sometimes, the alarm is caused by an earthquake.
- Also, John might confuse the alarm with his telephone, and Mary might miss the alarm altogether because she typically listens to loud music.

↪ **Random variables:** B, E, A, J, M for Burglary, Earthquake, Alarm, John, and Mary, respectively.

Given that both John and Mary call me, what is the probability of burglary?




Figure 1: Burglary-Earthquake Example

Question

How to calculate $\mathbf{P}(B \mid J, M)$?

(How inference in [Bayesian Networks](#) works?)

Let's use Enumeration ([Theorem 2.8](#)):

$$\mathbf{P}(Q \mid E_1 = e_1, \dots, E_{ne} = e_{ne}) = \alpha \left(\sum_{u \in D_u} \mathbf{P}(Q, E_1 = e_1, \dots, E_{ne} = e_{ne}, U_1 = u_1, \dots, U_{nu} = u_{nu}) \right)$$

with $Q = B, E = \{J, M\}, U = \{A, E\}$

$$\mathbf{P}(B \mid J, M) = \alpha \left(\sum_{b_a, b_e \in \{T, F\}} \mathbf{P}(B, J, M, A = b_a, E = b_e) \right)$$

Let's exploit [Theorem 3.1](#)

$$\mathbf{P}(X_1, \dots, X_n) = \prod_{X_i \in \mathcal{X}} \mathbf{P}(X_i \mid \text{Parents}(X_i))$$

$$\mathbf{P}(B \mid J, M) = \alpha \left(\sum_{b_a, b_e \in \{T, F\}} P(J \mid A = b_a) \cdot P(M \mid A = b_a) \cdot \mathbf{P}(A = b_a \mid E = b_e, B) \cdot P(E = b_e) \cdot \mathbf{P}(B) \right)$$

These are 5 factors in 4 summands over two cases for Burglary ↪ 38 arithmetic operations +3 for α ↪ General worst case: $\mathcal{O}(n \cdot 2^n)$

Question

Can we do better?

Answer

We can optimize a little bit:

$$\mathbf{P}(B \mid J, M) = \alpha \left(\mathbf{P}(B) \cdot \left(\sum_{b_a, b_e \in \{T, F\}} P(J \mid A = b_a) \cdot P(M \mid A = b_a) \cdot \mathbf{P}(A = b_a \mid E = b_e, B) \cdot P(E = b_e) \right) \right)$$

$$\mathbf{P}(B \mid J, M) = \alpha \left(\mathbf{P}(B) \cdot \left(\sum_{b_e \in \{T, F\}} P(E = b_e) \cdot \sum_{b_a \in \{T, F\}} \mathbf{P}(A = b_a \mid E = b_e, B) \cdot P(J \mid A = b_a) \cdot P(M \mid A = b_a) \right) \right)$$

This way, we have 3 factors in 2 summands + two factors in the outer product, over two cases $\rightsquigarrow$ 28 arithmetic operations +3 for α


Question

Can we do better?

Answer

Yes! $\rightsquigarrow$ **Variable Elimination**

Notice that enumeration corresponds to a depth-first traversal of an arithmetic expression tree, here is an example with some example probabilities:



Notice the redundancies!

- $P(e)$ and $P(\neg e)$ (No need for both!)
- $P(j \mid \neg a), P(j \mid a), P(m \mid \neg a), P(m \mid a)$ each is being computed twice! That is because $P(J \mid A = b_a)$ and $P(M \mid A = b_a)$ only depends on A , but they are trapped behind the summation over E , hence computed twice.

We can precompute these and store them (trading time complexity by space complexity) $\rightsquigarrow$ Variable Elimination.

The idea is, instead of left-to-right (top-down DFS), we operate right-to-left (bottom-up) and store intermediate *factors (functions)* along with their *dependencies*

Lemma 3.2. Given a query $\mathbf{P}(Q_1, \dots, Q_{n_Q} \mid E_1 = e_1, \dots, E_{n_E} = e_{n_E})$, we can ignore and remove all hidden variables that are not ancestors of any of the $Q_1, \dots, Q_{n_Q}$ or $E_1, \dots, E_{n_E}$

4 Decision Networks

We now know how to update our world model, represented as a **set** of **random variables**, given observations. Now we need to talk about **taking actions**.

Question

Given a world model and a **set** of actions, what are the likely consequences of action? How good are these consequences?

Idea

- (1) Represent actions as *special random variables*.
 - Given disjoint actions $a_1, \dots, a_n$, we introduce a **random variable** A with domain $\{a_1, \dots, a_n\}$. Then we can query $\mathbf{P}(X \mid A = a_i)$
- (2) Assign numerical values to the outcomes of actions, i.e. a **function** $u : \text{dom}(X) \rightarrow \mathbb{R}$
- (3) Choose the action that **maximizes** the **expected value** if u

Def 4.1. Decision Theory investigates **decision problems**, i.e. how a **utility-based agent** a deals with choosing among actions based on the desirability of their outcomes given by a real-valued **utility function** U on states $s \in S$, i.e. $U : S \rightarrow \mathbb{R}$

Note

If our states are **random variables**, then we obtain a **random variable** for the **utility function**.
 Let $X_i : \Omega \rightarrow D_i$ **random variables** on a **probability space** $\langle \Omega, \mathcal{F}, P \rangle$ and $f : D_1 \times \dots \times D_n \rightarrow \mathbb{R}$. Then $F(x) := f(X_0(x), \dots, X_n(x))$ is a **random variable** $\Omega \rightarrow \mathbb{R}$. Meaning, any **function** on the domains of **random variables** is itself a **random variable**.

Example 4.1. Consider commuting to work as an example experiment where we have $\Omega = \{\text{accident on the road, minor crash on the road, smooth commute}\}$. We define a **random variable** X representing travel time in minutes:

- $X(\text{accident on the road}) = 90$
- $X(\text{minor crash on the road}) = 45$
- $X(\text{smooth commute}) = 20$

These are the states our world can be in. We want to define a utility function F representing our happiness. Because X is a **random**, F becomes **random** too, we do not get one happiness level, we get a distribution of possible happiness levels based on the underlying reality Ω .

Say that $f(t) = 100 - t$ (the longer the trip, the less happy I am). Then we have:

Reality (ω)	Travel Time ($X(\omega)$)	Happiness ($f(X(\omega))$)
accident on the road	90	10
minor crash on the road	45	55
smooth commute	20	80

Def 4.2. Given a **probability space** $\langle \Omega, \mathcal{F}, P \rangle$ and a **random variable** $X : \Omega \rightarrow D$ with $D \subseteq \mathbb{R}$, then

$$\mathbb{E}(X) := \sum_{x \in D} P(X = x) \cdot x$$

is called the **Expected Value** of X . Analogously, let $e_1, \dots, e_n$ be a sequence of **events**. Then the **Expected Value** of X given $e_1, \dots, e_n$ is defined as

$$\mathbb{E}(X \mid e_1, \dots, e_n) := \sum_{x \in D} P(X = x \mid e_1, \dots, e_n) \cdot x$$

Example 4.2. Continuing with the Commute Time example, assume we have the following probabilities:

- $P(X = 20) = 0.7$
- $P(X = 45) = 0.2$
- $P(X = 90) = 0.10$

$$\begin{aligned}
\mathbb{E}(X) &:= \sum_{x \in D} P(X = x) \cdot x \\
&:= P(X = 20) \cdot 20 + P(X = 45) \cdot 45 + P(X = 90) \cdot 90 \\
&:= 0.7 \cdot 20 + 0.2 \cdot 45 + 0.10 \cdot 90 \\
\mathbb{E}(X) &:= 32 \text{ minutes}
\end{aligned}$$

This means, (while I might never actually spend exactly 32 minutes on any single day), if I commute for, e.g. a year, my average travel time will be 32 minutes.

Def 4.3. Given:

- $D = \{a_1, \dots, a_n\}$: a set of possible action, where we have a special **random variable** $A : \Omega \rightarrow D$ representing the action.
- $X_i : \Omega \rightarrow D_i$: **random variables** (the states)
- $U : D \times D_i \rightarrow \mathbb{R}$: a **utility function**

Then, the **Expected Utility** of the action $a \in D$ is the **expected value** of U (interpreted as a **random variable**) given $A = a$:

$$\text{EU}(a) := \sum_{\langle x_1, \dots, x_n \rangle \in D_1 \times \dots \times D_n} P(X_1 = x_1, \dots, X_n = x_n \mid A = a) \cdot U(x_1, \dots, x_n, a)$$

Example 4.3.

- $D = \{\text{take}, \text{leave}\}$ for taking or leaving the umbrella.
- X : random variable for the state of the weather outside with $D_x = \{\text{rain}, \text{sun}\}$ with **priors**: $P(X = \text{rain}) = 0.3, P(X = \text{sun}) = 0.7$
- $U : D \times D_x \rightarrow \mathbb{R}$ the utility function defined as:
 - $U(\text{sun}, \text{leave}) = 100$
 - $U(\text{sun}, \text{take}) = 80$
 - $U(\text{rain}, \text{leave}) = 0$
 - $U(\text{rain}, \text{take}) = 60$

$$\begin{aligned}
\text{EU}(\text{take}) &:= \sum_{x_i \in \{\text{rain}, \text{sun}\}} P(X = x_i \mid A = \text{take}) \cdot U(x_i, \text{take}) \\
&:= P(X = \text{rain} \mid A = \text{take}) \cdot U(\text{rain}, \text{take}) + P(X = \text{sun} \mid A = \text{take}) \cdot U(\text{sun}, \text{take}) \\
&:= P(X = \text{rain}) \cdot U(\text{rain}, \text{take}) + P(X = \text{sun}) \cdot U(\text{sun}, \text{take}) \\
&:= 0.3 \cdot 60 + 0.7 \cdot 80 = 74
\end{aligned}$$

$$\begin{aligned}
\text{EU}(\text{leave}) &:= \sum_{x_i \in \{\text{rain}, \text{sun}\}} P(X = x_i \mid A = \text{leave}) \cdot U(x_i, \text{leave}) \\
&:= P(X = \text{rain} \mid A = \text{leave}) \cdot U(\text{rain}, \text{leave}) + P(X = \text{sun} \mid A = \text{leave}) \cdot U(\text{sun}, \text{leave}) \\
&:= P(X = \text{rain}) \cdot U(\text{rain}, \text{leave}) + P(X = \text{sun}) \cdot U(\text{sun}, \text{leave}) \\
&:= 0.3 \cdot 0 + 0.7 \cdot 100 = 70
\end{aligned}$$

Def 4.4. MEU Principle for Rationality: We call an action **rational** if it **Maximizes the Expected Utility (MEU)**. A **utility-based agent** is called **rational**, iff it always chooses a **rational** action.

Def 4.5. Decision Network: A **decision network** is a **Bayesian Network** with two additional kinds of nodes:


1. **Action Nodes**: representing a **set** of possible actions

(square)

2. A single **Utility Node** (also called value node)

(diamond)

Simple Example (no dependencies):



General Algorithm

Given **evidence** $E_j = e_j$ and **action nodes** $A_1, \dots, A_k$, compute the **expected utility** of each action, given **evidence**, for each action we have:

$$EU(a \mid E = e_j) := \sum_{\langle x_1, \dots, x_n \rangle \in D_1 \times \dots \times D_n} P(X_1 = x_1, \dots, X_n = x_n \mid A = a, E_j = e_j) \cdot U(x_1, \dots, x_n, a)$$

And we are interested in the single action that **MEU**:

$$A_{\text{MEU}} = \arg \max_{a_1, \dots, a_k} EU(a_i \mid E = e_j)$$

And we say $\text{MEU} = EU(A_{\text{MEU}} \mid E = e_j)$ (maximum expected utility is the expected utility we get if we would choose the action that maximizes the expected utility)

Note

Decision Theorists are mostly *economists*, not statisticians. Economists usually talk about *prizes* (which we call states).

Def 4.6. Given prizes A and B , an agent can express preferences of the form:

- $A \succ B$: A is preferred over B
- $A \sim B$: indifference between A and B
- $A \succcurlyeq B$: B is not preferred over A (the agent wants A at least as much as $B \rightsquigarrow$ *weak preference*)

Note


In computer science, we define **relations** $\succ, \sim$ on our **set** of states $\mathcal{S}$

But that works only for **deterministic environments**. In a **non-deterministic** one, the agent simply does not know the consequences of the actions. Now we look at what **Decision Theorists** call *lotteries*

Def 4.7. Let $\mathcal{S}$ be a **set** of states. We call a **random variable** X with domain $\{A_1, \dots, A_n\} \subseteq \mathcal{S}$ a **lottery** and write $[p_1, A_1; p_2, A_2; \dots; p_n, A_n]$ where $p_i = P(X = A_i)$


Idea

A **lottery** represents the result of a **nondeterministic** action that can have outcomes A_i with **prior** p_i . For the binary case, we use $[p, A; 1 - p, B]$. We can then extend **preferences** to include **lotteries**, as a measure of how *strongly* we prefer one prize over another.



Note

We assume the **set** of states $\mathcal{S}$ to be closed under **lotteries**, i.e. **lotteries** are also states. That allows us to consider **lotteries** such as $[p, A; 1 - p, [q, B; 1 - q, C]]$



Preference Relations

- $X \succsim Y \wedge Y \succsim X \Rightarrow X \sim Y$
- $X \succsim Y \wedge \neg(Y \succsim X) \Rightarrow X \succ Y$

Def 4.8. We call a set $\succ$ of **preferences rational**, iff the following constraints hold:


- **Orderability:** $A \succ B \vee B \succ A \vee A \sim B$
- **Transitivity:** $A \succ B \wedge B \succ C \Rightarrow A \succ C$
- **Continuity:** $A \succ B \succ C \Rightarrow (\exists p. [p, A; 1 - p, C] \sim B)$
- **Substitutability:** $A \sim B \Rightarrow [p, A; 1 - p, C] \sim [p, B; 1 - p, C]$
- **Monotonicity:** $A \succ B \Rightarrow ((p > q) \Leftrightarrow [p, A; 1 - p, B] \succ [q, A; 1 - q, B])$
- **Decomposability:** $[p, A; 1 - p, [q, B; 1 - q, C]] \sim [p, A; ((1 - p)q), B; ((1 - p)(1 - q)), C]$

From a **set** of **rational preferences**, we can obtain a meaningful utility function.

In plain English, this definition says when a person's preferences are considered *rational* in decision theory.

- **Orderability:** for any two options, you must be able to compare them. You either prefer one, prefer the other, or are indifferent.
- **Transitivity:** if you prefer A to B , and B to C , then you should also prefer A to C .
- **Continuity:** if A is better than B , and B is better than C , then there is some probability p such that you are indifferent between B for sure and the lottery $[p, A; 1 - p, C]$.
- **Substitutability:** if you are indifferent between A and B , then replacing A by B inside any lottery should not change your preference.
- **Monotonicity:** if A is better than B , then a lottery with a higher probability of A should be preferred to one with a lower probability of A .

- **Decomposability:** a nested lottery should be equivalent to a single flat lottery with the same overall probabilities.



Theorem 4.1. Ramsey's Theorem: Given **rational preferences**, there exists a real valued function U such that $U(A) \geq U(B)$, iff $A \succsim B$ and $U([p, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$

Note

This theorem states that if an agent's / person's **preferences** are **rational**, then we can assign numbers to outcomes in a way that matches those **preferences**. i.e. there exists a utility **function** U that gives each outcome a number. If $A \succsim B$ then $U(A) \geq U(B)$. The higher the utility number, the more preferred the outcome. For a **lottery** $[p, S_1; \dots; p_n, S_n]$, its utility is just the probability-weighted average of the utilities of its outcomes.

Def 4.9. We call a **linear (total) ordering** on states a **value function** or **ordinal utility function**. With **deterministic** prizes only (no **lottery** choices), only a **total ordering** on prizes can be determined.

Note


If we don't need to care about *relative* utilities of states, e.g. to compute non-trivial **expected utilities**, that's all we need anyway! i.e. when there are no probabilities to compare, "utility" is just a way to sort outcomes from best to worst, not a precise measurement.

Def 4.10. Standard approach to assessment of human utilities : Compare a given state A to a **standard lottery** L_p that has:

- best possible prize $u_{\top}$ with probability p
- worst possible catastrophe $u_{\perp}$ with probability $1 - p$

adjust **lottery** probability p until $A \sim L_p$. Then $U(A) = p$

Example 4.4. Choose $u_{\top} \hat{=}$ current state, $u_{\perp} \hat{=}$ instant death



Popular Utility Functions

Def 4.11. Normalized utilities: $u_{\top} = 1, u_{\perp} = 0$ (not meaningful)

Def 4.12. Money: very intuitive, often easy to determine, but actually not well-suited as utility function

Def 4.13. Micromorts: one millionth chance of instant death. (useful for paying to reduce product risk, etc.). But not necessarily a good measure of risk, if the risk is only severe injury or illness

Def 4.14. QALYs: quality adjusted life years $\rightsquigarrow$ very better, more informative, useful for medical decisions involving substantial risk

Def 4.15. Value of Perfect Information (VPI): Let A be the **set** of available actions and F is a **random variable**. Given

evidence $E_i = e_i$, let α be the action that MEU a priori, and α_f the action that MEU given $F = f$, i.e.

$$\alpha = \arg \max_{a \in A} \text{EU}(a \mid E_i = e_i)$$


$$\alpha_f = \arg \max_{a \in A} \text{EU}(a \mid E_i = e_i, F = f)$$

The **value of perfect information (VPI)** on F given evidence $E_i = e_i$ is defined as:

$$\text{VPI}_{E_i=e_i}(F) := \left(\sum_{f \in \text{dom}(F)} P(F = f \mid E_i = e_i) \cdot \text{EU}(\alpha_f \mid E_i = e_i, F = f) \right) - \text{EU}(\alpha \mid E_i = e_i)$$

Example 4.5. consider two blocks A, B and exactly one has oil which worth $\$K$. You can drill in one location. You know the priors: $P(\text{oil at A}) = P(\text{oil at B}) = 0.5$.

We can model this a **Decision Network** with one chance node (oil location), one **action node** (drilling location), and one **utility node**:



Question

Without any information, what is the action that MEU?

$$\begin{aligned} \text{EU}(D = a) &= P(O = a) \cdot U(D = a, O = a) + P(O = b) \cdot U(D = a, O = b) \\ &= 0.5 \cdot K + 0.5 \cdot 0 = \frac{K}{2} \\ \text{EU}(D = b) &= P(O = a) \cdot U(D = b, O = a) + P(O = b) \cdot U(D = b, O = b) \\ &= 0.5 \cdot 0 + 0.5 \cdot K = \frac{K}{2} \end{aligned}$$

Answer

Which means, without any additional information, both actions have the same expected utility. We write $\text{MEU}(\emptyset)$ for the maximum expected utility given no extra information:

$$\text{MEU}(\emptyset) = \arg \max_a \text{EU}(a) = \frac{K}{2}$$

Question

What if someone offered to give us information about where the oil is. What is the value of that information? i.e. How much we are willing to pay to get such information?

Answer

Assume we have information that says *the oil location is b* (our **evidence**), then we get EU of 0 for $D = a$ and we get EU of K for $D = b$ and we will choose to dig in b immediately.

The value of this perfect information is simply what we get with information (K) minus what we would get without it ($\frac{K}{2}$).

$$\text{VPI}(E) = \text{MEU}(A \mid E) - \text{MEU}(\emptyset)$$

5 Hidden Markov Models

Motivation

- The world changes in *stochastically predictable ways*
- For example, the weather changes, but the weather tomorrow is somewhat predictable given today's weather and other factors.
- The stock market changes, but the stock price tomorrow is probably related to today's price
- A patient's blood sugar changes, but their blood sugar is related to their blood sugar 10 minutes ago

Question

How can we model this?

Def 5.1. Given a **probability space** $\langle \Omega, \mathcal{F}, P \rangle$, and a time structure S with an order $\langle S, \preceq \rangle$. Then a **sequence** of **random variables** $(X_t)_{t \in S}$ with $\text{dom}(X_t) = D$ is called **stochastic process** over the **time structure** S

Note

X_t models the outcome of the **random variable** X at a given time step t . The sample space Ω corresponds to the **set** of all possible **sequences** of outcomes.

We will almost exclusively use $\langle S, \preceq \rangle = \langle \mathbb{N}, \leq \rangle$

Notation

Given a **stochastic process** X_t over S and $a, b \in S$ with $a \preceq b$, we write $X_{a:b}$ for the **sequence** $X_a, X_{a+1}, \dots, X_{b-1}, X_b$ and $X_{a:b}^x$ for $X_a = x_a, \dots, X_b = x_b$

Idea


Construct a **Bayesian Network** from these variables

Def 5.2. Let $(X_t)_{t \in S}$ be a **stochastic process**. We say that X has **n-th order Markov property** iff X_t only depends on a bounded subset of $X_{0:t-1}$, i.e. for all $t \in S$ we have $\mathbf{P}(X_t \mid X_0, \dots, X_{t-1}) = \mathbf{P}(X_t \mid X_{t-n}, \dots, X_{t-1})$ for some $n \in S$.

A **stochastic process** with the **Markov property** for some n is called an **n-th order Markov process**


Def 5.3. First-Order Markov Property (Markov Chain):

$$\mathbf{P}(X_t \mid X_{0:t-1}) = \mathbf{P}(X_t \mid X_{t-1})$$



Def 5.4. Second-Order Markov Property:

$$\mathbf{P}(X_t \mid X_{0:t-1}) = \mathbf{P}(X_t \mid X_{t-2}, X_{t-1})$$



Note

Given a **Markov chain**, and we are now at time t , then the future is **conditionally independent** of the past given the present:

$$\mathbf{P}(X_{t+1} \mid X_{0:t}) = \mathbf{P}(X_{t+1} \mid X_t)$$

Def 5.5. Let $(X_t)_{t \in S}$ be a **stochastic process** where X has the **Markov property** (for any n). We call the **probability distribution** $\mathbf{P}(X_t | X_{0:t-1})$ the **transition model** of the **stochastic process** $(X_t)_{t \in S}$

Def 5.6. Stationary Assumption: A **Markov chain** is called **stationary** if the **transition model** is **independent** of time, i.e. $\mathbf{P}(X_t | X_{t-1})$ is the same for all t

Example 5.1. Consider we represent the weather as a **Markov chain**, we introduce a variable X with domain $\{\text{sun}, \text{rain}\}$ and we fix the initial distribution $\mathbf{P}(X) = \langle 1, 0 \rangle$ (sunny), and we have a **stationary transition model** that might be:


X_{t-1}	X_t	$P(X_t X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

Different ways of representing transition models

If a **Markov chain** X is **stationary** and with **finite domain**, we can represent the **transition model** as a matrix T where $T_{ij} := P(X_{t=j} | X_{t-1=i})$

$$T = \begin{pmatrix} P(\text{sun} | \text{sun}) & P(\text{rain} | \text{sun}) \\ P(\text{sun} | \text{rain}) & P(\text{rain} | \text{rain}) \end{pmatrix} = \begin{pmatrix} 0.9 & 0.1 \\ 0.3 & 0.7 \end{pmatrix}.$$

We can also represent it to look like a finite state machine :



Question

Given our weather example (example 5.1), what is the probability distribution of X after one time step (knowing it is initially sunny)?

Answer

- Trivially from the transition model $P(X_2 = \text{sun} | X_1 = \text{sun}) = 0.9$. Hence, $\mathbf{P}(X_2) = \langle 0.9, 0.1 \rangle$ But let's try to reason about it.
- We know that for $t = 1$, $P(\text{sun}) = 1$ and $P(\text{rain}) = 0$, i.e. the initial distribution is $\mathbf{P}(X_1) = \langle 1, 0 \rangle$. We also have a **stationary transition model**. We are interested in knowing $\mathbf{P}(X_2)$.

We compute it for *sun* by **marginalization** over the different values of $X_{t=1}$:

$$\begin{aligned}
 P(X_2 = \text{sun}) &= \sum_{x \in \{\text{sun}, \text{rain}\}} P(X_2 = \text{sun}, X_1 = x) \\
 &= \sum_{x \in \{\text{sun}, \text{rain}\}} P(X_2 = \text{sun} | X_1 = x) \cdot P(X_1 = x) \\
 P(X_2 = \text{sun}) &= P(X_2 = \text{sun} | X_1 = \text{sun}) \cdot P(X_1 = \text{sun}) + P(X_2 = \text{sun} | X_1 = \text{rain}) \cdot P(X_1 = \text{rain}) \\
 &= 0.9 \cdot 1.0 + 0.1 \cdot 0 = 0.9
 \end{aligned}$$

What we did is the simplest algorithm for **Markov chains**, it is called **mini forward**

Def 5.7. Mini Forward Algorithm: Given a **stationary Markov chain** X with initial distribution $\mathbf{P}(X_1)$, we the distribution at any time step t by recursively applying:

$$\mathbf{P}(X_n) = \sum_{x_{n-1}} P(X_n | X_{n-1} = x_{n-1}) \cdot P(X_{n-1} = x_{n-1})$$

for $n = 2, 3, \dots, t$. At each step, we **marginalize** out the previous time step by weighting each possible transition by its prior probability.

Conversion

If we were to do this for infinitely many time steps, the **probability distribution** $\mathbf{P}(X)$ of the **Markov chain** X converges.

$$\mathbf{P}(X_1) = \langle 1.0, 0.0 \rangle \rightsquigarrow \mathbf{P}(X_2) = \langle 0.9, 0.1 \rangle \rightsquigarrow \mathbf{P}(X_3) = \langle 0.84, 0.16 \rangle \rightsquigarrow$$

$$\mathbf{P}(X_4) = \langle 0.804, 0.196 \rangle \rightsquigarrow \dots \rightsquigarrow \mathbf{P}(X_\infty) = \langle 0.75, 0.25 \rangle$$

This applies to any initial distribution $\mathbf{P}(X) = \langle p, 1 - p \rangle$!

Def 5.8. As $t \rightarrow \infty$, $\mathbf{P}(X)$ converges towards a fixed point called the **stationary distribution** of the **Markov chain** X

In other words, the **Markov chain** essentially “forgets” its initial state; the long-term probabilities become completely independent of the starting conditions.

Computing the Stationary Distribution

To compute the **stationary distribution** S , we can use the matrix representation of the **transition model** T . Because the stationary distribution is a fixed point, passing it through the transition model returns the same distribution. If we treat the distribution S as a column vector and T is the transition matrix (as defined earlier, where rows sum to 1), then taking the next step involves the transpose of T . The stationary distribution satisfies:

$$S = T^T \cdot S$$

Note: If you were to represent S as a row vector instead, the equation would simply be $S = S \cdot T$.

To compute S (as a column vector), we solve the system of linear equations:

$$(T^T - I) \cdot S = 0$$

along with the necessary constraint that all probabilities must sum to 1 ($\sum_i S_i = 1$). This means S is an eigenvector of T^T corresponding to the eigenvalue $\lambda = 1$.

Note

Markov chains are not so useful for most agents, usually we need another ingredient, observations.

We begin to speak of **hidden Markov models** (HMMs) when we have an underlying **Markov chain** X that is **hidden** to us and we **observe** some **evidence** E that is affected by that **Markov chain**.

Def 5.9. We call a **stochastic process** of *hidden* variables a **state variable**


Example 5.2. You are a security guard in a secret underground facility, want to know it if is raining outside. Your only source of information is whether the director comes in with an umbrella.

- we have a **stochastic process** $\text{Rain}_0, \text{Rain}_1, \text{Rain}_2, \dots$ of **hidden** variables, and
- a related **stochastic process** $\text{Umbrella}_0, \text{Umbrella}_1, \text{Umbrella}_2, \dots$ of **evidence** variables.

And a combined **stochastic process** $\langle \text{Rain}_0, \text{Umbrella}_0 \rangle, \langle \text{Rain}_1, \text{Umbrella}_1 \rangle, \dots$

Note: Umbrella_t only depends on Rain_t , not on e.g. Umbrella_{t-1}

We model this as a **Bayesian Network**:



Def 5.10. We call the **conditional probability** distribution $\mathbf{P}(E_t \mid X_{0:t}, E_{1:t-1})$ a **sensor model**


Def 5.11. We say that a **sensor model** has the **sensor Markov property**, iff $\mathbf{P}(E_t \mid X_{0:t}, E_{1:t-1}) = \mathbf{P}(E_t \mid X_t)$, i.e. the **sensor model** depends only the current state.

Observation Matrix

If a **sensor model** has the **sensor Markov property**, we can represent each observation $E_t = e_t$ at time t as the *diagonal* matrix O_t with $O_{t_{ii}} := P(E_t = e_t \mid X_t = i)$

Def 5.12. A pair $\langle X, E \rangle$ where X is a **stationary Markov chain** (hidden state variable), and E has the **sensor markov property** (**evidence**) is called a **(stationary) Hidden Markov Model (HMM)**

Example 5.3. Our umbrella example (**example 5.2**) is a **HMM** with the following example **transition model** and **sensor model**:



Umbrella Matrices

$$O_T = \begin{pmatrix} 0.9 & 0 \\ 0 & 0.2 \end{pmatrix} \quad O_F = \begin{pmatrix} 0.1 & 0 \\ 0 & 0.8 \end{pmatrix} \quad T = \begin{pmatrix} 0.7 & 0.3 \\ 0.3 & 0.7 \end{pmatrix}$$

Def 5.13. Markov Inference Tasks: Given a **Markov chain** with state variables X_t and **evidence** variables E_t , we are interested in the following **Markov inference tasks**:

- **Filtering (Monitoring):** The task of tracking / updating the **probability distribution** $B_t(X) = \mathbf{P}(X_t \mid E_{1:t}^=)$ (**belief state**) over time. i.e. given sequence $e_1, \dots, e_t$ of observations up until time t , we compute the likely state of the world at current time t .
- **Prediction (State Estimation):** Given sequence $e_1, \dots, e_t$ of observations up until time t , compute the likely *future* state of the world at time $t+k$ for $k > 0$. i.e. calculating the **probability distribution** $\mathbf{P}(X_{t+k} \mid E_{1:t}^=)$.
- **Smoothing (Hindsight):** Given sequence $e_1, \dots, e_t$ of observations up until time t , compute the likely *past* state of the world at time $t-k$ for $0 < k < t$. i.e. calculate the **probability distribution** $\mathbf{P}(X_{t-k} \mid E_{1:t}^=)$

- **Most Likely Explanation (MLE)**: Given sequence $e_1, \dots, e_t$ of observations up until time t , compute the most likely sequence of states that led to these observations, i.e. calculate:

$$\arg \max_x P(X_{1:t}^x \mid E_{1:t}^e)$$

Note


The most likely sequence of states is *not* (necessarily) the sequence of most likely states

Def 5.14. Given a **Markov chain** X and **evidence** variables E , let $\mathbf{P}(X_0)$ be the initial **prior** (initial belief), then we can compute the full **JPD** as:

$$\mathbf{P}(X_{0:t}, E_{1:t}) = \mathbf{P}(X_0) \cdot \left(\prod_{i=1}^t \mathbf{P}(X_i \mid X_{i-1}) \cdot \mathbf{P}(E_i \mid X_i) \right)$$

Explanation

This is a **Bayesian Network** (recall [theorem 3.1](#)):



Filtering Base Case

We want to compute the belief state at $t = 1$, which is $B(X_1) = \mathbf{P}(X_1 \mid E_1)$. We can derive this in two steps as follows:

By Bayes' theorem, we incorporate the new evidence to update our belief:

$$\begin{aligned} \mathbf{P}(X_1 \mid E_1) &= \frac{\mathbf{P}(E_1 \mid X_1) \cdot \mathbf{P}(X_1)}{\mathbf{P}(E_1)} \\ &= \alpha(\mathbf{P}(E_1 \mid X_1) \cdot \mathbf{P}(X_1)) \end{aligned}$$

Here $\mathbf{P}(E_1 \mid X_1)$ is our **sensor model**.

Next, we expand the prior prediction $\mathbf{P}(X_1)$ by **marginalizing** over the initial state X_0 :

$$\begin{aligned} \mathbf{P}(X_1) &= \sum_{x_0} \mathbf{P}(X_1, X_0 = x_0) \\ &= \sum_{x_0} \mathbf{P}(X_1 \mid X_0 = x_0) \cdot \mathbf{P}(X_0 = x_0) \end{aligned}$$

Here, $\mathbf{P}(X_1 \mid X_0)$ is the **transition model**, and $\mathbf{P}(X_0)$ is the initial belief prior.

Combining these gives the complete filtering formula for $t = 1$:

$$\mathbf{P}(X_1 \mid E_1) = \alpha \left(\mathbf{P}(E_1 \mid X_1) \sum_{x_0} \mathbf{P}(X_1 \mid x_0) \cdot \mathbf{P}(x_0) \right)$$

Prediction Base Case

If we know the initial belief distribution $\mathbf{P}(X_0)$ and want to compute the distribution of the next state $\mathbf{P}(X_1)$ prior to any observations, we perform a 1-step prediction.

We can compute this by [marginalizing](#) over all possible values of the initial state X_0 :

$$\mathbf{P}(X_1) = \sum_{x_0 \in \text{dom}(X_0)} \mathbf{P}(X_1, X_0 = x_0)$$

Using the chain rule of probability, we split the joint probability into the [transition model](#) and the initial [prior](#):

$$\mathbf{P}(X_1) = \sum_{x_0} \mathbf{P}(X_1 \mid X_0 = x_0) \cdot \mathbf{P}(X_0 = x_0)$$

This equation tells us that the probability of being in a state at $t = 1$ is the sum of the probabilities of transitioning to it from every possible previous state, weighted by how likely those previous states were.

Filtering Derivation

To compute the belief state at time t given all evidence up to time t , we want to find $\mathbf{P}(X_t \mid E_{1:t})$. We split the evidence into the current observation E_t and all past observations $E_{1:t-1}$.

Step 1: Update (Bayes' Rule / Product Rule)

We want to find $\mathbf{P}(X_t \mid E_t, E_{1:t-1})$. By definition of conditional probability:

$$\begin{aligned} \mathbf{P}(X_t \mid E_{1:t}) &= \mathbf{P}(X_t \mid E_t, E_{1:t-1}) \\ &= \frac{\mathbf{P}(X_t, E_t, E_{1:t-1})}{\mathbf{P}(E_t, E_{1:t-1})} \end{aligned}$$

We apply the product rule to the numerator ($\mathbf{P}(A, B, C) = \mathbf{P}(A \mid B, C) \cdot \mathbf{P}(B, C)$):

$$\mathbf{P}(X_t, E_t, E_{1:t-1}) = \mathbf{P}(E_t \mid X_t, E_{1:t-1}) \cdot \mathbf{P}(X_t, E_{1:t-1})$$

And apply the product rule once more to $\mathbf{P}(X_t, E_{1:t-1})$:

$$\mathbf{P}(X_t, E_t, E_{1:t-1}) = \mathbf{P}(E_t \mid X_t, E_{1:t-1}) \cdot \mathbf{P}(X_t \mid E_{1:t-1}) \cdot \mathbf{P}(E_{1:t-1})$$

We also apply the product rule to the denominator:

$$\mathbf{P}(E_t, E_{1:t-1}) = \mathbf{P}(E_t \mid E_{1:t-1}) \cdot \mathbf{P}(E_{1:t-1})$$

Now substitute both back into the fraction:

$$\mathbf{P}(X_t \mid E_{1:t}) = \frac{\mathbf{P}(E_t \mid X_t, E_{1:t-1}) \cdot \mathbf{P}(X_t \mid E_{1:t-1}) \cdot \mathbf{P}(E_{1:t-1})}{\mathbf{P}(E_t \mid E_{1:t-1}) \cdot \mathbf{P}(E_{1:t-1})}$$

The term $\mathbf{P}(E_{1:t-1})$ cancels out:

$$\mathbf{P}(X_t \mid E_{1:t}) = \frac{\mathbf{P}(E_t \mid X_t, E_{1:t-1}) \cdot \mathbf{P}(X_t \mid E_{1:t-1})}{\mathbf{P}(E_t \mid E_{1:t-1})}$$

Since $\frac{1}{\mathbf{P}(E_t \mid E_{1:t-1})}$ is just a normalizing constant, we can simplify our equation by replacing it with α :

$$\mathbf{P}(X_t \mid E_{1:t}) = \alpha \left(\mathbf{P}(E_t \mid X_t, E_{1:t-1}) \cdot \mathbf{P}(X_t \mid E_{1:t-1}) \right)$$

By the [Sensor Markov property](#), E_t only depends on X_t , simplifying the first term (the [sensor model](#)):

$$\mathbf{P}(X_t \mid E_{1:t}) = \alpha \left(\mathbf{P}(E_t \mid X_t) \cdot \mathbf{P}(X_t \mid E_{1:t-1}) \right)$$

Step 2: Prediction (Law of Total Probability)

Now we need to compute the 1-step prediction $\mathbf{P}(X_t | E_{1:t-1})$. We do this by [marginalizing](#) over the previous state X_{t-1} :

$$\mathbf{P}(X_t | E_{1:t-1}) = \sum_x \mathbf{P}(X_t, X_{t-1} = x | E_{1:t-1})$$

To show how the terms split, we expand the conditional joint probability using the definition of conditional probability:

$$\mathbf{P}(X_t, X_{t-1} | E_{1:t-1}) = \frac{\mathbf{P}(X_t, X_{t-1}, E_{1:t-1})}{\mathbf{P}(E_{1:t-1})}$$

Applying the product rule to the numerator gives $\mathbf{P}(X_t, X_{t-1}, E_{1:t-1}) = \mathbf{P}(X_t | X_{t-1}, E_{1:t-1}) \cdot \mathbf{P}(X_{t-1}, E_{1:t-1})$:

$$\mathbf{P}(X_t, X_{t-1} | E_{1:t-1}) = \frac{\mathbf{P}(X_t | X_{t-1}, E_{1:t-1}) \cdot \mathbf{P}(X_{t-1}, E_{1:t-1})}{\mathbf{P}(E_{1:t-1})}$$

Notice that $\frac{\mathbf{P}(X_{t-1}, E_{1:t-1})}{\mathbf{P}(E_{1:t-1})}$ is simply the definition of conditional probability for $\mathbf{P}(X_{t-1} | E_{1:t-1})$. Therefore:

$$\mathbf{P}(X_t, X_{t-1} | E_{1:t-1}) = \mathbf{P}(X_t | X_{t-1}, E_{1:t-1}) \cdot \mathbf{P}(X_{t-1} | E_{1:t-1})$$

Substituting this conditionalized product rule back into our original marginalization sum gives:

$$\mathbf{P}(X_t | E_{1:t-1}) = \sum_x \mathbf{P}(X_t | X_{t-1} = x, E_{1:t-1}) \cdot \mathbf{P}(X_{t-1} = x | E_{1:t-1})$$

By the First-Order [Markov property](#), the transition to X_t only depends on X_{t-1} , simplifying the first term (the [transition model](#)):

$$\mathbf{P}(X_t | E_{1:t-1}) = \sum_x \mathbf{P}(X_t | X_{t-1} = x) \cdot \mathbf{P}(X_{t-1} = x | E_{1:t-1})$$

Step 3: Combine

Substituting the prediction back into the update equation gives the recursive filtering formula.

$$\mathbf{P}(X_t | E_{1:t}^{\bar{e}}) = \alpha \left(\mathbf{P}(E_t = e_t | X_t) \cdot \left(\sum_x \mathbf{P}(X_t | X_{t-1} = x) \cdot \mathbf{P}(X_{t-1} = x | E_{1:t-1}^{\bar{e}}) \right) \right)$$

Note

Notice that $\mathbf{P}(X_{t-1} = x | E_{1:t-1}^{\bar{e}})$ is simply the filtering result from the *previous* time step. This means we can compute the belief state recursively!

Def 5.15. We call the inner part of the above expression the **forward** algorithm:

$$\mathbf{P}(X_t | E_{1:t}^{\bar{e}}) = \alpha (\text{FORWARD}(e_t, \mathbf{P}(X_{t-1} | E_{1:t-1}^{\bar{e}}))) = f_{1:t}$$

Filtering in Matrix Notation

$$\mathbf{P}(X_t | E_{1:t}^{\bar{e}}) = \alpha (O_t \cdot T^T \cdot \mathbf{P}(X_{t-1} | E_{1:t-1}^{\bar{e}})) = f_{1:t}$$

Filtering the Umbrellas

Let's assume we were given the following information:

- $\mathbf{P}(R_0) = \langle 0.5, 0.5 \rangle$
- $\mathbf{P}(R_t | R_{t-1}) = 0.6, \mathbf{P}(\neg R_t | \neg R_{t-1}) = 0.8$

- $P(U_t | R_t) = 0.9, P(\neg U_t | \neg R_t) = 0.85$

From that, we can infer:

- $P(\neg R_t | R_{t-1}) = 0.4$
- $P(R_t | \neg R_{t-1}) = 0.2$
- $P(\neg U_t | R_t) = 0.1$
- $P(U_t | \neg R_t) = 0.15$

Hence, we construct the matrices:

$$O_T = \begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \quad O_F = \begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix} \quad T = \begin{pmatrix} 0.6 & 0.4 \\ 0.2 & 0.8 \end{pmatrix}$$

Assume that the director carries an umbrella on days 1 and 2, and not on day 3, hence:

$$O_1 = O_2 = O_T = \begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \quad O_3 = O_F = \begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix}$$

Remember:

$$\mathbf{P}(X_t | E_{1:t}^{\neg e}) = \alpha(O_t \cdot T^T \cdot \mathbf{P}(X_{t-1} | E_{1:t-1}^{\neg e})) = f_{1:t}$$

Note

Throughout these calculations, our state vector order is fixed: the top value represents the probability of Rain (R_t) and the bottom value represents the probability of No Rain ($\neg R_t$). i.e. $\begin{pmatrix} P(R_t | E_{1:t}) \\ P(\neg R_t | E_{1:t}) \end{pmatrix}$. This is because of how we initially constructed the transition and observation matrices.

For $t = 1$, we have:

$$\begin{aligned} f_{1:1} &= \alpha(O_1 \cdot T^T \cdot \mathbf{P}(R_0)) \\ &= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} 0.6 & 0.2 \\ 0.4 & 0.8 \end{pmatrix} \begin{pmatrix} 0.5 \\ 0.5 \end{pmatrix}\right) \\ &= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} (0.6 \cdot 0.5) + (0.2 \cdot 0.5) \\ (0.4 \cdot 0.5) + (0.8 \cdot 0.5) \end{pmatrix}\right) \\ &= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} 0.4 \\ 0.6 \end{pmatrix}\right) \\ &= \alpha\begin{pmatrix} 0.9 \cdot 0.4 \\ 0.15 \cdot 0.6 \end{pmatrix} \\ &= \alpha\begin{pmatrix} 0.36 \\ 0.09 \end{pmatrix} \\ &= \frac{1}{0.36 + 0.09} \begin{pmatrix} 0.36 \\ 0.09 \end{pmatrix} = \begin{pmatrix} \frac{0.36}{0.45} \\ \frac{0.09}{0.45} \end{pmatrix} \\ &= \begin{pmatrix} 0.8 \\ 0.2 \end{pmatrix} \end{aligned}$$

For $t = 2$, we have:

$$\begin{aligned}
f_{1:2} &= \alpha(O_2 \cdot T^T \cdot f_{1:1}) \\
&= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} 0.6 & 0.2 \\ 0.4 & 0.8 \end{pmatrix} \begin{pmatrix} 0.8 \\ 0.2 \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} (0.6 \cdot 0.8) + (0.2 \cdot 0.2) \\ (0.4 \cdot 0.8) + (0.8 \cdot 0.2) \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} 0.52 \\ 0.48 \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.9 \cdot 0.52 \\ 0.15 \cdot 0.48 \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.468 \\ 0.072 \end{pmatrix}\right) \\
&= \frac{1}{0.468 + 0.072} \begin{pmatrix} 0.468 \\ 0.072 \end{pmatrix} = \begin{pmatrix} \frac{0.468}{0.54} \\ \frac{0.072}{0.54} \end{pmatrix} \\
&\approx \begin{pmatrix} 0.867 \\ 0.133 \end{pmatrix}
\end{aligned}$$

For $t = 3$, the director did *not* bring an umbrella, so we use $O_3 = O_F$:

$$\begin{aligned}
f_{1:3} &= \alpha(O_3 \cdot T^T \cdot f_{1:2}) \\
&= \alpha\left(\begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix} \begin{pmatrix} 0.6 & 0.2 \\ 0.4 & 0.8 \end{pmatrix} \begin{pmatrix} 0.867 \\ 0.133 \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix} \begin{pmatrix} (0.6 \cdot 0.867) + (0.2 \cdot 0.133) \\ (0.4 \cdot 0.867) + (0.8 \cdot 0.133) \end{pmatrix}\right) \\
&\approx \alpha\left(\begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix} \begin{pmatrix} 0.547 \\ 0.453 \end{pmatrix}\right) \\
&= \alpha\left(\begin{pmatrix} 0.1 \cdot 0.547 \\ 0.85 \cdot 0.453 \end{pmatrix}\right) \\
&\approx \alpha\left(\begin{pmatrix} 0.055 \\ 0.385 \end{pmatrix}\right) \\
&= \frac{1}{0.055 + 0.385} \begin{pmatrix} 0.055 \\ 0.385 \end{pmatrix} = \begin{pmatrix} \frac{0.055}{0.440} \\ \frac{0.385}{0.440} \end{pmatrix} \\
&\approx \begin{pmatrix} 0.125 \\ 0.875 \end{pmatrix}
\end{aligned}$$

At day 3, there is a 12.5% chance it is raining, and an 87.5% it is not raining.

Prediction

Idea

Prediction is **filtering** without new **evidence**, i.e. we can use **filtering** until t , and then continue as follows:

$$\mathbf{P}(X_{t+k} \mid E_{1:t}^{\neg e}) = \sum_x \mathbf{P}(X_{t+k} \mid X_{t+k-1} = x) \cdot P(X_{t+k-1} = x \mid E_{1:t}^{\neg e})$$

Note

By recursively applying this formula k times, starting from our current filtered belief state $\mathbf{P}(X_t \mid E_{1:t}^{\neg e})$, we can predict k steps into the future.

$$\mathbf{P}(X_{t+k} \mid E_{1:t}^{\neg e}) = T^T \cdot \mathbf{P}(X_{t+k-1} \mid E_{1:t}^{\neg e})$$

Smoothing Derivation

Idea

Smoothing computes $\mathbf{P}(X_{t-k} \mid E_{1:t}^{\neg e})$ for $k > 0$. The intuition is to split the evidence into two independent parts: the past/current evidence up to the target state ($E_{1:t-k}$) and the future evidence after the target state ($E_{t-k+1:t}$). We then compute a forward filtering pass up to time $t - k$, and a backward pass from time t down to $t - k$.

Step 1: Divide the Evidence (Product Rule)

We want to find our belief about state X_{t-k} given all evidence $E_{1:t}$. Let's divide the evidence at the target state: $E_{1:t} = E_{1:t-k}, E_{t-k+1:t}$

$$\mathbf{P}(X_{t-k} \mid E_{1:t-k}, E_{t-k+1:t}) = \frac{\mathbf{P}(X_{t-k}, E_{t-k+1:t}, E_{1:t-k})}{\mathbf{P}(E_{t-k+1:t}, E_{1:t-k})}$$

We apply the product rule to the numerator ($\mathbf{P}(A, B, C) = \mathbf{P}(A \mid B, C) \cdot \mathbf{P}(B \mid C) \cdot \mathbf{P}(C)$) and denominator ($\mathbf{P}(A, B) = \mathbf{P}(A \mid B) \cdot \mathbf{P}(B)$), keeping $E_{1:t-k}$ as the conditioning context:

$$\mathbf{P}(X_{t-k} \mid E_{1:t}) = \frac{\mathbf{P}(E_{t-k+1:t} \mid X_{t-k}, E_{1:t-k}) \cdot \mathbf{P}(X_{t-k} \mid E_{1:t-k}) \cdot \mathbf{P}(E_{1:t-k})}{\mathbf{P}(E_{t-k+1:t} \mid E_{1:t-k}) \cdot \mathbf{P}(E_{1:t-k})}$$

we cancel out $\mathbf{P}(E_{1:t-k})$ and replace the normalizing denominator with α :

$$\mathbf{P}(X_{t-k} \mid E_{1:t}) = \alpha \left(\mathbf{P}(E_{t-k+1:t} \mid X_{t-k}, E_{1:t-k}) \cdot \mathbf{P}(X_{t-k} \mid E_{1:t-k}) \right)$$

Step 2: Apply Conditional Independence

By the **Markov property**, if we lock in our knowledge of the target state X_{t-k} , then the future evidence $E_{t-k+1:t}$ becomes independent of the past evidence $E_{1:t-k}$. (We say the state X_{t-k} "d-separates" the past and future). Thus, we drop the past evidence from the conditional:

$$\mathbf{P}(E_{t-k+1:t} \mid X_{t-k}, E_{1:t-k}) = \mathbf{P}(E_{t-k+1:t} \mid X_{t-k})$$

Substituting this back gives us the foundational equation for smoothing:

$$\mathbf{P}(X_{t-k} \mid E_{1:t}) = \alpha \left(\underbrace{\mathbf{P}(E_{t-k+1:t} \mid X_{t-k})}_{\text{backward message } b_{t-k+1:t}} \cdot \underbrace{\mathbf{P}(X_{t-k} \mid E_{1:t-k})}_{\text{forward message } f_{1:t-k}} \right)$$

Notice that the second term is exactly our **filtering** formula result! We call this the **forward message** $f_{1:t-k}$. The first term is a new concept, denoting the probability of observing all upcoming future evidence given the target state. We formally define this as the **backward message** $b_{t-k+1:t}$.

Step 3: Deriving the Backward Message Recursion

Let $b_{m:t} = \mathbf{P}(E_{m:t} \mid X_{m-1})$ be a generic backward message representing future evidence $E_{m:t}$ given the state immediately preceding it (X_{m-1}). We want to recursively express this in terms of the next state down the line, X_m . We do this by **marginalizing** over all possible values of X_m :

$$\mathbf{P}(E_{m:t} \mid X_{m-1}) = \sum_{x \in \text{dom}(X)} \mathbf{P}(E_{m:t}, X_m = x \mid X_{m-1})$$

To split the inner joint probability, we apply the definition of conditional probability ($\mathbf{P}(A, B \mid C) = \frac{\mathbf{P}(A, B, C)}{\mathbf{P}(C)}$) and then use the product rule on the numerator ($\mathbf{P}(A, B, C) = \mathbf{P}(A \mid B, C) \cdot \mathbf{P}(B, C)$):

$$\begin{aligned} \mathbf{P}(E_{m:t}, X_m = x \mid X_{m-1}) &= \frac{\mathbf{P}(E_{m:t}, X_m = x, X_{m-1})}{\mathbf{P}(X_{m-1})} \\ &= \frac{\mathbf{P}(E_{m:t} \mid X_m = x, X_{m-1}) \cdot \mathbf{P}(X_m = x, X_{m-1})}{\mathbf{P}(X_{m-1})} \end{aligned}$$

Since $\frac{\mathbf{P}(X_m=x, X_{m-1})}{\mathbf{P}(X_{m-1})} = \mathbf{P}(X_m = x \mid X_{m-1})$ (by definition of conditional probability again), substituting this back gives us the conditionalized product rule:

$$\mathbf{P}(E_{m:t} \mid X_{m-1}) = \sum_{x \in \text{dom}(X)} \mathbf{P}(E_{m:t} \mid X_m = x, X_{m-1}) \cdot \mathbf{P}(X_m = x \mid X_{m-1})$$

Next, we split the future evidence $E_{m:t}$ into purely the current observation E_m and the remaining future observations $E_{m+1:t}$:

$$= \sum_{x \in \text{dom}(X)} \mathbf{P}(E_m, E_{m+1:t} \mid X_m = x, X_{m-1}) \cdot \mathbf{P}(X_m = x \mid X_{m-1})$$

First, by the [Markov property](#), future observations ($E_{m:t}$) are conditionally independent of the past state X_{m-1} given the current state X_m . In other words, knowing the current state completely shields the future from the past. We can thus safely drop X_{m-1} from the conditioning:

$$= \sum_{x \in \text{dom}(X)} \mathbf{P}(E_m, E_{m+1:t} \mid X_m = x) \cdot \mathbf{P}(X_m = x \mid X_{m-1})$$

Next, we must separate the joint probability $\mathbf{P}(E_m, E_{m+1:t} \mid X_m = x)$. Unpacking this completely using the definition of conditional probability ($\mathbf{P}(A, B \mid C) = \frac{\mathbf{P}(A, B, C)}{\mathbf{P}(C)}$) gives:

$$\mathbf{P}(E_m, E_{m+1:t} \mid X_m = x) = \frac{\mathbf{P}(E_{m+1:t}, E_m, X_m = x)}{\mathbf{P}(X_m = x)}$$

We apply the standard product rule to the numerator ($\mathbf{P}(A, B, C) = \mathbf{P}(A \mid B, C) \cdot \mathbf{P}(B, C)$):

$$= \frac{\mathbf{P}(E_{m+1:t} \mid E_m, X_m = x) \cdot \mathbf{P}(E_m, X_m = x)}{\mathbf{P}(X_m = x)}$$

And applying the definition of conditional probability once more on the remaining fraction ($\frac{\mathbf{P}(A, B)}{\mathbf{P}(B)} = \mathbf{P}(A \mid B)$):

$$= \mathbf{P}(E_{m+1:t} \mid E_m, X_m = x) \cdot \mathbf{P}(E_m \mid X_m = x)$$

Finally, we recognize that the current evidence E_m and future evidence $E_{m+1:t}$ are conditionally independent given the current state X_m (because the sensor model depends solely on the immediate state). Due to this conditional independence, knowing E_m provides no extra information about $E_{m+1:t}$ if we already know X_m , so we can drop E_m from the second term:

$$\mathbf{P}(E_m, E_{m+1:t} \mid X_m = x) = \mathbf{P}(E_m \mid X_m = x) \cdot \mathbf{P}(E_{m+1:t} \mid X_m = x)$$

Substituting this fully separated expression back into our main summation yields:

$$= \sum_{x \in \text{dom}(X)} \mathbf{P}(E_m \mid X_m = x) \cdot \mathbf{P}(E_{m+1:t} \mid X_m = x) \cdot \mathbf{P}(X_m = x \mid X_{m-1})$$

Rearranging these terms, we see the recursive relationship emerges beautifully:

$$b_{m:t} = \sum_{x \in \text{dom}(X)} \underbrace{\mathbf{P}(E_m \mid X_m = x)}_{\text{sensor model}} \cdot \underbrace{\mathbf{P}(E_{m+1:t} \mid X_m = x)}_{\text{next backward msg } b_{m+1:t}} \cdot \underbrace{\mathbf{P}(X_m = x \mid X_{m-1})}_{\text{transition model}}$$

If we substitute our target $m = t - k + 1$, we get the recursive definition:

$$b_{t-k+1:t} = \sum_{x \in \text{dom}(X)} \mathbf{P}(E_{t-k+1} = e_{t-k+1} \mid X_{t-k+1} = x) \cdot b_{t-k+2:t} \cdot \mathbf{P}(X_{t-k+1} = x \mid X_{t-k})$$

Note

As a starting point for the backward recursion at time t , we let $b_{t+1:t}$ be the explicit uniform vector with 1 in every component, since there is no evidence after t .

Backward Algorithm Matrix Notation

The generic backward message recursion algorithm represented as matrix multiplication is:

$$b_{t-k+1:t} = T \cdot O_{t-k+1} \cdot b_{t-k+2:t}$$

Notice that we evaluate using T , not T^T like we do in forward filtering.

Smoothing Final Algorithm Matrix Notation

Combining the forward and backward passes, the full **smoothing algorithm** becomes:

$$\mathbf{P}(X_{t-k} \mid E_{1:t}^{\neq e}) = \alpha(f_{1:t-k} \times b_{t-k+1:t})$$

Where $\times$ denotes component-wise (Hadamard) vector multiplication.

Smoothing Umbrellas

Using [filtering](#) we already calculated

$$f_{1:1} = \begin{pmatrix} 0.8 \\ 0.2 \end{pmatrix}, f_{1:2} = \begin{pmatrix} 0.867 \\ 0.133 \end{pmatrix}, f_{1:3} = \begin{pmatrix} 0.125 \\ 0.875 \end{pmatrix}$$

Let's compute our smoothed belief for day 1 ($k = 2$):

$$\mathbf{P}(R_1 \mid E_{1:3}) = \mathbf{P}(R_1 \mid U_1 = T, U_2 = T, U_3 = F) = \alpha(f_{1:1} \times b_{2:3})$$

For that, we need to compute $b_{2:3}$ and $b_{3:3}$ using the backward algorithm: $b_{m:t} = T \cdot O_m \cdot b_{m+1:t}$. The base case at time $t = 3$ is $b_{4:3} = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$.

1. Calculate backward message $b_{3:3}$

$$\begin{aligned} b_{3:3} &= T \cdot O_3 \cdot b_{4:3} \\ &= \begin{pmatrix} 0.6 & 0.4 \\ 0.2 & 0.8 \end{pmatrix} \begin{pmatrix} 0.1 & 0 \\ 0 & 0.85 \end{pmatrix} \begin{pmatrix} 1 \\ 1 \end{pmatrix} \\ &= \begin{pmatrix} 0.6 & 0.4 \\ 0.2 & 0.8 \end{pmatrix} \begin{pmatrix} 0.1 \\ 0.85 \end{pmatrix} \\ &= \begin{pmatrix} (0.6 \cdot 0.1) + (0.4 \cdot 0.85) \\ (0.2 \cdot 0.1) + (0.8 \cdot 0.85) \end{pmatrix} = \begin{pmatrix} 0.06 + 0.34 \\ 0.02 + 0.68 \end{pmatrix} = \begin{pmatrix} 0.4 \\ 0.7 \end{pmatrix} \end{aligned}$$

2. Calculate backward message $b_{2:3}$

$$\begin{aligned} b_{2:3} &= T \cdot O_2 \cdot b_{3:3} \\ &= \begin{pmatrix} 0.6 & 0.4 \\ 0.2 & 0.8 \end{pmatrix} \begin{pmatrix} 0.9 & 0 \\ 0 & 0.15 \end{pmatrix} \begin{pmatrix} 0.4 \\ 0.7 \end{pmatrix} \\ &= \begin{pmatrix} 0.6 & 0.4 \\ 0.2 & 0.8 \end{pmatrix} \begin{pmatrix} 0.36 \\ 0.105 \end{pmatrix} \\ &= \begin{pmatrix} (0.6 \cdot 0.36) + (0.4 \cdot 0.105) \\ (0.2 \cdot 0.36) + (0.8 \cdot 0.105) \end{pmatrix} = \begin{pmatrix} 0.216 + 0.042 \\ 0.072 + 0.084 \end{pmatrix} = \begin{pmatrix} 0.258 \\ 0.156 \end{pmatrix} \end{aligned}$$

3. Combine Forward and Backward messages

Now we take our forward filtered probability for day 1 ($f_{1:1}$) and element-wise multiply ($\times$) it with the backward message

from the future ($b_{2:3}$):

$$\begin{aligned}
 \mathbf{P}(R_1 \mid E_{1:3}) &= \alpha(f_{1:1} \times b_{2:3}) \\
 &= \alpha\left(\begin{pmatrix} 0.8 \\ 0.2 \end{pmatrix} \times \begin{pmatrix} 0.258 \\ 0.156 \end{pmatrix}\right) \\
 &= \alpha\begin{pmatrix} 0.8 \cdot 0.258 \\ 0.2 \cdot 0.156 \end{pmatrix} \\
 &= \alpha\begin{pmatrix} 0.2064 \\ 0.0312 \end{pmatrix}
 \end{aligned}$$

Calculating the normalization constant $\alpha = \frac{1}{0.2064+0.0312} = \frac{1}{0.2376} \approx 4.2088$:

$$\begin{aligned}
 \mathbf{P}(R_1 \mid E_{1:3}) &\approx 4.2088 \cdot \begin{pmatrix} 0.2064 \\ 0.0312 \end{pmatrix} \\
 &\approx \begin{pmatrix} 0.8687 \\ 0.1313 \end{pmatrix}
 \end{aligned}$$

Note

Compare this with the original filtered probability for day 1 ($f_{1:1} = \langle 0.8, 0.2 \rangle$). Given that the director brought an umbrella on day 2, it made sense in hindsight that it was even more likely to be raining on day 1 (86.87% vs 80%).

Def 5.16. Dynamic Bayesian Networks (DBN): A **Bayesian Network** $\mathcal{D}$ is called **dynamic** iff its **random variables** are **indexed** by a time structure. We assume $\mathcal{D}$ is:

- **time sliced**, i.e. that the **time slices** $\mathcal{D}_t$ - the subgraphs of t -indexed **random variables** and the edges between them - are **isomorphic**.
- a **stationary Markov chain**, i.e. that variables X_t can only have parents in $\mathcal{D}_t$ and $\mathcal{D}_{t-1}$

Observation

- Every **HMM** is a single-variable **DBN** *(trivially)*
- Every **DBN** can be turned into an **HMM** *(combine variables into tuple $\rightsquigarrow$ lose information about dependencies)*
- **DBNs** have sparse dependencies $\rightsquigarrow$ exponentially fewer parameters

6 Markov Decision Processes


Motivation

In **deterministic** environments, the next state is completely determined by the current state and the chosen action. For these problems, we define a corresponding **deterministic** search problem where the objective is to find a solution consisting of a specific **sequence** of actions—an optimal plan—leading from the initial state to a goal state.

However, in a **stochastic** environment, the next state is **not** uniquely determined by the current state and action. In these **stochastic** search problems, the **agent** controls the action but lacks direct control over the environment's transitions. While the exact outcome of an action is unknown, the **agent** knows the probabilities of transitioning to various states given a specific action.

Consequently, a fixed **sequence** of actions is no longer sufficient. Instead, the goal is to find an **optimal policy**, π^* , which is a **function** $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ that dictates the best action for any given state. By following this policy, the **agent** navigates the environment's uncertainty to reach a goal state efficiently.


Example 6.1. Consider a grid world where the agent can go west, east, north, or south. Notice the difference between **deterministic** and **stochastic** environments:



Note

We are still assuming **fully observable** environments

Def 6.1. MDP Running Example: We will consider the following **fully observable** 4×3 grid world with **stochastic** actions:



Def 6.2. MDP: A Markov Decision Process (MDP) is the **structure** $\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{I}, R \rangle$ where:

- $\mathcal{S}$: the **set** of states
- $\mathcal{I}$: the **set** of initial states. Since we are in a **fully observable** environment, $\mathcal{I} = \{s_0\}$.
- $\mathcal{A}$: the **set** of actions, with $\mathcal{A}_s$ denoting the **set** of action that can be taken from state s .
- $\mathcal{T}$: a transition model $\mathcal{T}(s, a) = \mathbf{P}(\mathcal{S} \mid s, a)$
- R : a **reward function** $R : \mathcal{S} \rightarrow \mathbb{R}$ (we call $R(s)$ a **reward**)

Idea

We use the **rewards** as a utility function: the goal is to choose actions such that the expected cumulative **rewards** for the foreseeable future is maximized.

Note

In **MDPs**, the aim is to find an **optimal policy** $\pi(s)$, which tells us the best action for every possible state s . (because we can't predict where we might end up, we need to consider all states)

Def 6.3. Policy: A **policy** π for an **MDP** is a **function** mapping each state $s \in \mathcal{S}$ to an action $a \in \mathcal{A}_s$. An **optimal policy** is a policy that maximizes the expected total **rewards**.

Example 6.2.




Figure 2: Different optimal policies for different reward functions

For $R(s) = -0.01$, we have a small negative reward for each state. Hence, each action will cost a little bit and not too much. Therefore, in state $(3, 2)$, the optimal action is to go west where there is a 80% chance we stay where we are. We cannot go north since there is a 10% of ending up in the pit (-1) . Because the action costs only 0.01, the agent prefers to bang his head to the wall until by chance goes north or south. Similarly, in $(4, 1)$.

For $R(s) = -0.03$, the action cost is a bit high. The agent prefer to take the chance of going north (in $(3, 2)$) and west in $(4, 1)$ regardless of the 10% probability of falling in the pit.

For $R(s) = -0.4$, the cost now is much higher, the agent will try to finish as soon as possible.

For $R(s) = -2$, life is very painful, the agent might commit suicide to alleviate the pain.

Question

Assume we have two sequences of three actions, the first one will collect the rewards 1, 2, 2, and second one will collect 2, 3, 4. Which sequence should the agent choose?

Answer

Well, since we agreed that we want to maximize the total sum of rewards, obviously we should prefer 2, 3, 4 because they sum to $9 > 5$.

Question

What about 0, 0, 1 and 1, 0, 0?

Observation

They have the same total sum of rewards.

Idea

- It is reasonable to maximize the sum of rewards
- But it is also reasonable to prefer rewards **now** over rewards **later**

Answer

Then the agent should prefer 1, 0, 0 over 0, 0, 1

Question

But how can we model this into the agent's brain?

Answer

We should let the values of the rewards decay exponentially over time

Def 6.4. We call **preferences** on **reward sequences** **stationary** iff

$$[r, r_0, r_1, r_2, \dots] \succ [r, r'_0, r'_1, r'_2, \dots] \Leftrightarrow [r_0, r_1, r_2, \dots] \succ [r'_0, r'_1, r'_2, \dots]$$

Theorem 6.1. For **stationary preferences**, there are only two ways to combine **rewards** over time:

- **additive rewards:** $U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$
- **discounted rewards:** $U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$ where $0 \leq \gamma \leq 1$ is called **discount factor**

Question


Given **discounted rewards** with $\gamma = 0.5$, what is the utility of the **rewards sequences** [1, 2, 3], and [3, 2, 1]?

Answer

- $U([1, 2, 3]) = 1 + 0.5 \cdot 2 + 0.25 \cdot 3 = 2.75$
 - $U([3, 2, 1]) = 3 + 0.5 \cdot 2 + 0.25 \cdot 1 = 4.25$
- $U([3, 2, 1]) > U([1, 2, 3])$ (the sooner the better)

Discounting Exercise


Consider the following world:



- $\mathcal{S} = \{a, b, c, d, e\}$
- $\mathcal{A}_s = \{E, W\} \quad \forall s \in \{b, c, d\}$ (East and West)
- $\mathcal{T}$ is **deterministic**
- If we are in a , we get **reward** 10.
- If we are in e , we get **reward** 1.
- $R(s) = 0 \quad \forall s \in \{b, c, d\}$


Question

For $\gamma = 1$ (no **discounted rewards**), what is the optimal policy ?

Answer**Question**

For $\gamma = 0.1$, what is the optimal policy ?

Answer



Question

For which γ are West and East equally good when agent in state d ?

Answer

- Going East: $0 + \gamma \cdot 1$
- Going West: $0 + \gamma \cdot 0 + \gamma^2 \cdot 0 + \gamma^3 \cdot 10$
- $\gamma = 10 \cdot \gamma^3$
- $1 = 10 \cdot \gamma^2$
- $\gamma^2 = \frac{1}{10} \rightsquigarrow \gamma = \frac{1}{\sqrt{10}}$

What if we do not have terminal states ?

Observation

If we are using **additive rewards** and we do not have terminal states, we can get infinite reward!

Idea

We have two solutions to this problem:

- Finite Horizon: we terminate after a fixed T steps. $\rightsquigarrow$ Gives **non-stationary** policies (π depends on the time left)
- Absorbing States: for any policy π , the agent eventually dies with probability = 1. $\rightsquigarrow$ **Expected utility** of every state is finite

Observation

A better solution is to use **discounted rewards** with $0 \leq \gamma < 1$.

Proof.

- Let $R_{\max}$ be the maximum reward the agent can get.
- $U([r_0, r_1, r_2, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t$
- $\sum_{t=0}^{\infty} \gamma^t r_t \leq \sum_{t=0}^{\infty} \gamma^t R_{\max}$
- $\sum_{t=0}^{\infty} \gamma^t R_{\max}$ is a convergent geometric series (since $\gamma < 1$)
- $\sum_{t=0}^{\infty} \gamma^t R_{\max} = \frac{R_{\max}}{1-\gamma}$
- Therefore, $U([r_0, r_1, r_2, \dots, r_\infty]) \leq \frac{R_{\max}}{1-\gamma}$

□

Remember - The Geometric Series

Def 6.5. A **geometric series** is the sum of an infinite sequence where each term is found by multiplying the previous one by a fixed, non-zero constant r , called the **common ratio**. Given an initial term $a \neq 0$, the series is expressed as:

$$\sum_{n=0}^{\infty} ar^n = a + ar + ar^2 + ar^3 + \dots$$

To determine if the infinite sum exists, we define the **k -th partial sum**, S_k , as the sum of the first k terms:

$$S_k = \sum_{n=0}^{k-1} ar^n = a + ar + ar^2 + \dots + ar^{k-1}$$

Through algebraic manipulation, this can be written in closed form as:

$$S_k = a \left(\frac{1 - r^k}{1 - r} \right) \quad \text{for } r \neq 1$$

The infinite series converges if the limit of the partial sums exists and is finite:

$$S = \lim_{k \rightarrow \infty} S_k$$

If the magnitude of the ratio is less than one ($|r| < 1$), then the term r^k vanishes as k approaches infinity:

$$\lim_{k \rightarrow \infty} r^k = 0$$

Substituting this into the partial sum formula yields the **sum of the infinite geometric series**:

$$S = \lim_{k \rightarrow \infty} a \left(\frac{1 - r^k}{1 - r} \right) = a \left(\frac{1 - 0}{1 - r} \right) = \frac{a}{1 - r}$$

- If $|r| < 1$, the series **converges** to $\frac{a}{1-r}$.
- If $|r| \geq 1$, the series **diverges** (the sum does not approach a finite limit).

Def 6.6. Given a **sequence** of state $S = s_0, s_1, s_2, \dots$ and a **discount factor** $0 \leq \gamma < 1$, the **utility** of the **sequence** is :

$$U(S) = \sum_{t=0}^{\infty} \gamma^t R(s_t)$$

Def 6.7. Given a **policy** π and a starting state s_0 . We define a **random variable** $S_{s_0}^\pi$ which represent the **sequence** of states resulting from executing π at every state starting from s_0 (because the environment is **stochastic**, we do not know the exact sequence, hence the **random variable** treatment).

Def 6.8. The **expected utility** obtained by executing the policy π starting in s_0 is given by:

$$U^\pi(s_0) := \text{EU}(S_{s_0}^\pi)$$

Def 6.9. The **optimal policy** π^* is a policy that maximizes the **expected utility** for the agent:

$$\pi_{s_0}^* := \arg \max_{\pi} U^\pi(s_0)$$

Observation

While we initially define it relative to a starting state s_0 . It is a fundamental property of MDPs that there exists a single policy π^* that is optimal for **all** starting states simultaneously.

Def 6.10. The **true utility** $U(s)$ of a state is the **expected utility** obtained by following the **optimal policy** starting in that state:

$$U(s) = U^{\pi^*}(s)$$

Def 6.11. Once these **utilities** are known, the **optimal policy** can be extracted locally using the **MEU** principle. Choosing the best action is just maximizing the **expected utility** of the immediate successor states:

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}_s} \left(\sum_{s'} P(s' | s, a) \cdot U(s') \right)$$

Given the **true utilities**, we can compute the **optimal policy** and vice versa.

Example 6.3. Assume we know the **utilities** of our grid world:

3	0.812	0.868	0.912	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388
	1	2	3	4

What is the **optimal action** for state (3,4)?

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}_s} \left(\sum_{s'} P(s' | s, a) \cdot U(s') \right)$$

$$\pi^*((3,4)) = \arg \max_{a \in \{N, S, E, W\}} \left(\sum_{s'} P(s' | s, a) \cdot U(s') \right)$$

For $a = N$, we have $s' \in \{(3,2), (4,1), (2,1)\}$

$$\sum_{s'} P(s' | s, a) \cdot U(s') = 0.8 \cdot 0.660 + 0.1 \cdot 0.388 + 0.1 \cdot 0.655 = 0.6326$$

For $a = S$

$$\sum_{s'} P(s' | s, a) \cdot U(s') = 0.8 \cdot 0.611 + 0.1 \cdot 0.388 + 0.1 \cdot 0.655 = 0.5931$$

For $a = E$

$$\sum_{s'} P(s' | s, a) \cdot U(s') = 0.8 \cdot 0.388 + 0.1 \cdot 0.660 + 0.1 \cdot 0.611 = 0.4375$$

For $a = W$

$$\sum_{s'} P(s' | s, a) \cdot U(s') = 0.8 \cdot 0.655 + 0.1 \cdot 0.611 + 0.1 \cdot 0.660 = 0.6511$$

Therefore:

$$\pi^*((3,4)) = \arg \max_{a \in \{N, S, E, W\}} \left(\sum_{s'} P(s' | s, a) \cdot U(s') \right) = W$$

Question

How to compute these **utilities** without knowing the **optimal policy**?

Observation

The expected sum of **rewards** = current **reward** + $\gamma \cdot$ expected **reward** sum after taking the **optimal action**

Bellman Equation

Theorem 6.2.

$$U(s) := R(s) + \gamma \cdot \max_{a \in \mathcal{A}_s} \sum_{s'} U(s') \cdot P(s' | s, a)$$

Example 6.4. Assume our grid world with $R(s) = -0.04$. Then the utility for, e.g. $(1, 1)$ is:

$$\begin{aligned}
 U(s) &:= R(s) + \gamma \cdot \max_{a \in \mathcal{A}_s} \sum_{s'} U(s') \cdot P(s' \mid s, a) \\
 U((1, 1)) &:= R((1, 1)) + \gamma \cdot \max_{a \in \{N, S, E, W\}} \sum_{s'} U(s') \cdot P(s' \mid (1, 1), a) \\
 &:= -0.04 + \gamma \cdot \max_{a \in \{N, S, E, W\}} \sum_{s'} U(s') \cdot P(s' \mid (1, 1), a) \\
 U(1, 1) &= -0.04 + \gamma \cdot \max\{ \\
 &\quad 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), & \text{North} \\
 &\quad 0.9U(1, 1) + 0.1U(1, 2), & \text{West} \\
 &\quad 0.9U(1, 1) + 0.1U(2, 1), & \text{South} \\
 &\quad 0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1)\} & \text{East}
 \end{aligned}$$

The Computational Challenge: The Bellman equations are inherently **nonlinear** due to the max operator. This coupling, combined with a high-dimensional state space, renders direct analytical solutions via standard Linear Algebra **intractable**.

Idea

Solve iteratively using e.g., dynamic programming:

- Start with arbitrary **utility** values
- At every iteration, update the **utility** values for each state using the values from the previous iteration
- Keep iterating until convergence $\rightsquigarrow$ every s , eventually will have the true $U(s)$

Def 6.12.

Value Iteration

$$\begin{aligned}
 U_{i+1}(s) &\leftarrow R(s) + \gamma \cdot \max_{a \in \mathcal{A}_s} \sum_{s'} U_i(s') \cdot P(s' \mid s, a) \\
 \text{Terminate when} \quad &\max_{s \in S} |U_i(s) - U_{i+1}(s)| \quad \text{is small enough}
 \end{aligned}$$

Def 6.13. The **maximum norm** is defined as $\|U\| = \max_s |U(s)|$. Consequently, $\|U - V\|$ represents the maximum difference between U and V .

Example 6.5. Consider a state space $S = \{s_1, s_2, s_3\}$ with the following utility values:

$$U = \begin{bmatrix} 2 \\ -5 \\ 4 \end{bmatrix}$$

To find the maximum norm $\|U\|$, we calculate the absolute value for each state and select the largest:

$$\begin{aligned}
 \|U\| &= \max_{s \in S} |U(s)| \\
 &= \max(|2|, |-5|, |4|) \\
 &= \max(2, 5, 4) \\
 &= 5
 \end{aligned}$$

The maximum norm is 5 because -5 is the value furthest from zero in this vector.

Example 6.6. Consider two utility vectors U and V for the state space $S = \{s_1, s_2, s_3\}$:

$$U = \begin{bmatrix} 10 \\ 2 \\ 5 \end{bmatrix}, \quad V = \begin{bmatrix} 7 \\ 2 \\ 8 \end{bmatrix}$$

To find $\|U - V\|$, we find the maximum absolute difference across all states:

$$\begin{aligned}
 \|U - V\| &= \max_{s \in S} |U(s) - V(s)| \\
 &= \max(|10 - 7|, |2 - 2|, |5 - 8|) \\
 &= \max(|3|, |0|, |-3|) \\
 &= \max(3, 0, 3) \\
 &= 3
 \end{aligned}$$

This value represents the largest "error" or shift between the two approximations.

Theorem 6.3. Let U^t and U^{t+1} be successive approximations to the **true utility** vector U during **value iteration**. Then for any two approximations U^t, V^t we have:

$$\|U^{t+1} - V^{t+1}\| \leq \gamma \|U^t - V^t\|$$

where:

- $U, V \in \mathbb{R}^{|S|}$ are vectors where each component $U(s)$ represents the expected long-term reward from state s .
- $\|U - V\|$ is the "distance" between these two utility estimates.

$\rightsquigarrow$ The update rule ensures that U and V move closer together at a rate defined by the discount factor γ .

Note

- That means, any distinct approximations get closer to each other over time.
- In our case, they get closer to the **true utility** vector U
- Consequently, **value iteration** converges to a unique, stable, optimal solution

Theorem 6.4. If $\|U^{t+1} - U^t\| < \epsilon$, then:

$$\|U^{t+1} - U\| < \frac{2\epsilon\gamma}{1-\gamma}$$

Note

- In practice, we cannot check the distance to the **true utility** vector U because it is unknown. We can only check the change between steps: $\|U^{t+1} - U^t\|$.
- **Theorem 6.3** guarantees that if the change in the last step is less than ϵ , then the absolute error from the optimal solution is no larger than $\frac{2\epsilon\gamma}{1-\gamma}$.
- As the discount factor γ approaches 1, the denominator $(1 - \gamma)$ becomes very small, which makes the error bound much larger. This reflects how "far-sighted" problems are harder to converge.
- This allows me to pre-calculate a stopping threshold ϵ to ensure the final result is within a desired accuracy range.

Idea

To bridge the gap between theory and code, we use a specific termination condition in the **value iteration** algorithm:

- We want the final error to be small, specifically $\|U^{t+1} - U\| < \epsilon$.
- During the loop, we track $\delta = \|U^{t+1} - U^t\|$, which is the change between the current and previous guess.
- Based on **Theorem 6.4**, if we want the true error to be less than ϵ , we must iterate until our observed change δ satisfies:

$$\delta < \epsilon \frac{(1-\gamma)}{\gamma}$$

- As the discount factor γ gets closer to 1, the value $(1 - \gamma)$ becomes very small. This forces the algorithm to wait until δ is extremely tiny, ensuring that even small "leftover" changes don't accumulate into a large error over an infinite time horizon.

Def 6.14. The **Value Iteration Algorithm** is given by **Algorithm 1**

Algorithm 1 Value Iteration

```
1: function VALUE-ITERATION( $mdp, \epsilon$ )
2:   inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ , rewards  $R(s)$ , discount  $\gamma$ 
3:   inputs:  $\epsilon$ , the maximum error allowed in the utility of any state
4:   local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
5:   local variables:  $\delta$ , the maximum change in the utility of any state in an iteration
6:   repeat
7:      $U \leftarrow U'$ 
8:      $\delta \leftarrow 0$ 
9:     for each state  $s$  in  $S$  do
10:       $U'[s] \leftarrow R(s) + \gamma \cdot \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
11:      if  $|U'[s] - U[s]| > \delta$  then
12:         $\delta \leftarrow |U'[s] - U[s]|$ 
13:      end if
14:    end for
15:  until  $\delta < \epsilon(1 - \gamma)/\gamma$ 
16:  return  $U$ 
17: end function
```

Example 6.7. Assume again our grid world. We will use $\gamma = 1$ (additive rewards) for simplicity and $R(s) = -0.04$. We start with $U_0(s) = 0$, we want to iterate twice to calculate $U_1((3,2))$, $U_1((3,3))$, $U_2((3,2))$ and $U_2((3,3))$:




Figure 3: U_0

$$U_{i+1}(s) \leftarrow R(s) + \gamma \cdot \max_a \sum_{s'} U_i(s') \cdot P(s' | s, a)$$

$$U_1(s_{3,2}) \leftarrow R(s_{3,2}) + 1 \cdot \max_a \sum_{s'} U_0(s') \cdot P(s' | s_{3,2}, a)$$

$$U_1(s_{3,2}) \leftarrow -0.04 + 1 \cdot \max_a \sum_{s'} U_0(s') \cdot P(s' | s_{3,2}, a)$$

$$\begin{aligned} &\leftarrow -0.04 + \max\{ \\ &0 \cdot 0.8 + -1 \cdot 0.1 + 0 \cdot 0.1, \\ &-1 \cdot 0.8 + 0 \cdot 0.1 + 0 \cdot 0.1, \\ &0 \cdot 0.8 + 0 \cdot 0.1 + 0 \cdot 0.1, \\ &0 \cdot 0.8 + -1 \cdot 0.1 + 0 \cdot 0.1\} \end{aligned}$$

North
East
West
South

$$U_1(s_{3,2}) \leftarrow -0.04 + \max\{-0.1, -0.8, 0, -0.1\}$$

$$U_1(s_{3,2}) \leftarrow -0.04 + 0$$

$$U_1(s_{3,2}) \leftarrow -0.04$$

$$\begin{aligned}
U_1(s_{3,3}) &\leftarrow -0.04 + 1 \cdot \max_a \sum_{s'} U_0(s') \cdot P(s' \mid s_{3,3}, a) \\
&\leftarrow -0.04 + \max\{ \\
&\quad 0 \cdot 0.8 + 1 \cdot 0.1 + 0 \cdot 0.1, & \text{North} \\
&\quad 1 \cdot 0.8 + 0 \cdot 0.1 + 0 \cdot 0.1, & \text{East} \\
&\quad 0 \cdot 0.8 + 0 \cdot 0.1 + 0 \cdot 0.1, & \text{West} \\
&\quad 0 \cdot 0.8 + 1 \cdot 0.1 + 0 \cdot 0.1\} & \text{South} \\
U_1(s_{3,3}) &\leftarrow -0.04 + \max\{0.1, 0.8, 0, 0.1\} \\
U_1(s_{3,3}) &\leftarrow -0.04 + 0.8 \\
U_1(s_{3,3}) &\leftarrow 0.76
\end{aligned}$$

Note

For the first iteration, the only state that will have a positive **utility** is the one that is adjacent to the state where $R(s) = +1$. For every other state, one step will not take us to that state, hence, we will end up having $U(s) = R(s) = -0.04$ for every other state.




Figure 4: U_1

To generalize, If we are calculating the **utility** estimate for iteration i , only the states that can reach the "+1" state in i steps or less will have a positive estimate of **utility**.

$$\begin{aligned}
U_{i+1}(s) &\leftarrow R(s) + \gamma \cdot \max_a \sum_{s'} U_i(s') \cdot P(s' \mid s, a) \\
U_2(s_{3,2}) &\leftarrow R(s_{3,2}) + 1 \cdot \max_a \sum_{s'} U_1(s') \cdot P(s' \mid s_{3,2}, a) \\
U_2(s_{3,2}) &\leftarrow -0.04 + 1 \cdot \max_a \sum_{s'} U_1(s') \cdot P(s' \mid s_{3,2}, a) \\
&\leftarrow -0.04 + \max\{ \\
&\quad 0.76 \cdot 0.8 + -1 \cdot 0.1 + -0.04 \cdot 0.1, & \text{North} \\
&\quad -1 \cdot 0.8 + 0.76 \cdot 0.1 + -0.04 \cdot 0.1, & \text{East} \\
&\quad -0.04 \cdot 0.8 + 0.76 \cdot 0.1 + -0.04 \cdot 0.1, & \text{West} \\
&\quad -0.04 \cdot 0.8 + -1 \cdot 0.1 + -0.04 \cdot 0.1\} & \text{South} \\
U_2(s_{3,2}) &\leftarrow -0.04 + \max\{0.504, -0.728, 0.04, -0.136\} \\
U_2(s_{3,2}) &\leftarrow -0.04 + 0.504 \\
U_2(s_{3,2}) &\leftarrow 0.464
\end{aligned}$$

$$\begin{aligned}
U_{i+1}(s) &\leftarrow R(s) + \gamma \cdot \max_a \sum_{s'} U_i(s') \cdot P(s' | s, a) \\
U_2(s_{3,3}) &\leftarrow R(s_{3,3}) + 1 \cdot \max_a \sum_{s'} U_1(s') \cdot P(s' | s_{3,3}, a) \\
U_2(s_{3,3}) &\leftarrow -0.04 + 1 \cdot \max_a \sum_{s'} U_1(s') \cdot P(s' | s_{3,3}, a) \\
&\leftarrow -0.04 + \max\{ \\
&\quad 0.76 \cdot 0.8 + -0.04 \cdot 0.1 + 1 \cdot 0.1, & \text{North} \\
&\quad 1 \cdot 0.8 + 0.76 \cdot 0.1 + -0.04 \cdot 0.1, & \text{East} \\
&\quad -0.04 \cdot 0.8 + 0.76 \cdot 0.1 + -0.04 \cdot 0.1, & \text{West} \\
&\quad -0.04 \cdot 0.8 + 1 \cdot 0.1 + -0.04 \cdot 0.1\} & \text{South} \\
U_2(s_{3,3}) &\leftarrow -0.04 + \max\{0.704, 0.872, 0.04, 0.064\} \\
U_2(s_{3,3}) &\leftarrow -0.04 + 0.872 \\
U_2(s_{3,3}) &\leftarrow 0.832
\end{aligned}$$

$U_2(s_{2,3})$ will be 0.56, All other states will have **utility** = -0.08. Therefore, second iteration **utilities**:

3	-0.08	0.56	0.832	+1
2	-0.08		0.464	-1
1	-0.08	-0.08	-0.08	-0.08
	1	2	3	4

Figure 5: U_2

Again, the states that have positive **utilities** in the second iteration are those which can reach the "+1" state in 2 steps or less.

Observation

While executing the **value iteration algorithm**, you may observe that despite the utility values haven't converged yet, the **optimal policy** already stopped changing.

Deriving the **optimal policy** doesn't necessarily require us to have accurate estimates of the **true utilities**.

Idea

Search for **optimal policy** and **utility values** simultaneously $\rightsquigarrow$ **Policy Iteration**

Def 6.15. Policy Iteration alternates between two steps:

- **Policy Evaluation:** Given a **policy** π_i , calculate $U_i^\pi(s)$ which is the **utility** of each state s if π_i were to be executed.
- **Policy Improvement:** Calculate a new **policy** π_{i+1} using U_i^π

Terminate when there is no change in **policy**.

Value vs Policy Iteration

Value Iteration:

$$U_0(s) \rightarrow U_1(s) \rightarrow U_2(s) \rightarrow \dots \rightarrow U^*(s)$$

Then calculate **optimal action** for each using:

$$\pi^*(s) = \arg \max_{a \in \mathcal{A}_s} \left(\sum_{s'} P(s' | s, a) \cdot U(s') \right)$$

Policy Iteration:

$$\pi_0(s) \xrightarrow{\text{eval}} U_0(s) \xrightarrow{\text{impr}} \pi_1(s) \xrightarrow{\text{eval}} U_1(s) \dots \pi^*(s) \xrightarrow{\text{eval}} U^*(s) \xrightarrow{\text{impr}} \pi^*(s)$$

Policy Improvement

$$\pi_{i+1}(s) := \arg \max_a \sum_{s'} P(s' | a, s) \cdot U_i(s')$$

Policy Evaluation

$$U_i(s) := R(s) + \gamma \cdot \sum_{s'} P(s' | s, \pi_i(s)) \cdot U_i(s')$$

Note

- Notice that we do not have the max operator anymore since we are already using a policy $\pi_i(s)$.
- This yields a linear system of equations where we know everything except the **utilities** (what we solve for)
- Solving this system is easier than solving the **Bellman Equation** where we have a nonlinearity because we do not know the actions (max operator)
- For n states, this will take $\mathcal{O}(n^3)$ time.

Policy Iteration Mini Exercise

Consider the following grid world with $R(s) = -0.04$ and $\gamma = 1$. Assume the initial policy $\pi_0(s_{1,1}) = \pi_0(s_{1,2}) = E$ (east) and the same transition probabilities as before.

2		+1
1		-1
	1	2

We evaluate the initial policy:

$$U_i(s) := R(s) + \gamma \cdot \sum_{s'} P(s' | s, \pi_i(s)) \cdot U_i(s')$$

$$U_0(s) := R(s) + \gamma \cdot \sum_{s'} P(s' | s, \pi_0(s)) \cdot U_0(s')$$

$$U_0(s_{1,1}) := -0.04 + \sum_{s'} P(s' | s_{1,1}, E) \cdot U_0(s')$$

$$U_0(s_{1,1}) := -0.04 + 0.8 \cdot U_0(s_{2,1}) + 0.1 \cdot U_0(s_{1,2}) + 0.1 \cdot U_0(s_{1,1})$$

$$U_0(s_{1,1}) := -0.04 + 0.8 \cdot -1 + 0.1 \cdot U_0(s_{1,2}) + 0.1 \cdot U_0(s_{1,1})$$

$$U_0(s_{1,2}) := -0.04 + \sum_{s'} P(s' | s_{1,2}, E) \cdot U_0(s')$$

$$U_0(s_{1,2}) := -0.04 + 0.8 + 0.1 \cdot U_0(s_{1,2}) + 0.1 \cdot U_0(s_{1,1})$$

We have the following system (two equations and two variables):

$$U_0(s_{1,1}) = -0.84 + 0.1U_0(s_{1,2}) + 0.1U_0(s_{1,1}) \quad (1)$$

$$U_0(s_{1,2}) = 0.76 + 0.1U_0(s_{1,2}) + 0.1U_0(s_{1,1}) \quad (2)$$

Rearranging terms to the left-hand side:

$$0.9U_0(s_{1,1}) - 0.1U_0(s_{1,2}) = -0.84 \quad (1')$$

$$-0.1U_0(s_{1,1}) + 0.9U_0(s_{1,2}) = 0.76 \quad (2')$$

Adding (1') and (2') yields:

$$0.8U_0(s_{1,1}) + 0.8U_0(s_{1,2}) = -0.08 \implies U_0(s_{1,1}) + U_0(s_{1,2}) = -0.1$$

Solving for $U_0(s_{1,1}) = -0.1 - U_0(s_{1,2})$ and substituting into (2'):

$$-0.1(-0.1 - U_0(s_{1,2})) + 0.9U_0(s_{1,2}) = 0.76$$

$$0.01 + 0.1U_0(s_{1,2}) + 0.9U_0(s_{1,2}) = 0.76$$

$$U_0(s_{1,2}) = 0.75$$

Substituting back to find $U_0(s_{1,1})$:

$$U_0(s_{1,1}) = -0.1 - 0.75 = -0.85$$

The final utilities are:

$$U_0(s_{1,1}) = -0.85, \quad U_0(s_{1,2}) = 0.75$$

Now we improve:

$$\pi_{i+1}(s) := \arg \max_a \sum_{s'} P(s' | a, s) \cdot U_i(s')$$

$$\pi_1(s) := \arg \max_a \sum_{s'} P(s' | a, s) \cdot U_0(s')$$

$$\pi_1(s_{1,1}) := \arg \max_a \sum_{s'} P(s' | a, s_{1,1}) \cdot U_0(s')$$

$$\pi_1(s_{1,1}) := \arg \max_a \{$$

$$0.8 \cdot 0.75 + 0.1 \cdot -1 + 0.1 \cdot -0.85,$$

North

$$0.8 \cdot -0.85 + 0.1 \cdot -0.85 + 0.1 \cdot -1,$$

South

$$0.8 \cdot -1 + 0.1 \cdot 0.75 + 0.1 \cdot -0.85,$$

East

$$0.8 \cdot -0.85 + 0.1 \cdot 0.75 + 0.1 \cdot -0.85\}$$

West

$$\pi_1(s_{1,1}) := \arg \max_a \{$$

$$0.415,$$

North

$$-0.865,$$

South

$$-0.81,$$

East

$$-0.69\}$$

West

$$\pi_1(s_{1,1}) := N$$

Similarly, for $s_{1,2}$:

$$\begin{aligned}
\pi_1(s_{1,2}) &:= \arg \max_a \sum_{s'} P(s' \mid a, s_{1,2}) \cdot U_0(s') \\
\pi_1(s_{1,2}) &:= \arg \max_a \{ \\
&\quad 0.8 \cdot 0.75 + 0.1 \cdot 0.75 + 0.1 \cdot 1, && \text{North} \\
&\quad 0.8 \cdot -0.85 + 0.1 \cdot 0.75 + 0.1 \cdot 1, && \text{South} \\
&\quad 0.8 \cdot 1 + 0.1 \cdot 0.75 + 0.1 \cdot -0.85, && \text{East} \\
&\quad 0.8 \cdot 0.75 + 0.1 \cdot 0.75 + 0.1 \cdot -0.85 \} && \text{West} \\
\pi_1(s_{1,2}) &:= \arg \max_a \{ \\
&\quad 0.775, && \text{North} \\
&\quad -0.505, && \text{South} \\
&\quad 0.79, && \text{East} \\
&\quad 0.59 \} && \text{West} \\
\pi_1(s_{1,2}) &:= E && (\text{didn't change})
\end{aligned}$$

Now we evaluate π_1 :

$$\begin{aligned}
U_i(s) &:= R(s) + \gamma \cdot \sum_{s'} P(s' \mid s, \pi_i(s)) \cdot U_i(s') \\
U_1(s) &:= R(s) + \gamma \cdot \sum_{s'} P(s' \mid s, \pi_1(s)) \cdot U_1(s') \\
U_1(s_{1,1}) &:= -0.04 + \sum_{s'} P(s' \mid s, \pi_1(s_{1,1})) \cdot U_1(s') \\
U_1(s_{1,1}) &:= -0.04 + 0.8 \cdot U_1(s_{1,2}) + 0.1 \cdot -1 + 0.1 \cdot U_1(s_{1,1}) \\
U_1(s_{1,2}) &:= -0.04 + \sum_{s'} P(s' \mid s, \pi_1(s_{1,2})) \cdot U_1(s') \\
U_1(s_{1,2}) &:= -0.04 + 0.8 \cdot 1 + 0.1 \cdot U_1(s_{1,2}) + 0.1 \cdot U_1(s_{1,1})
\end{aligned}$$

We have the following system (two equations and two variables):

$$U_1(s_{1,1}) := -0.04 + 0.8 \cdot U_1(s_{1,2}) - 0.1 + 0.1 \cdot U_1(s_{1,1}) \quad (3)$$

$$U_1(s_{1,2}) := -0.04 + 0.8 + 0.1 \cdot U_1(s_{1,2}) + 0.1 \cdot U_1(s_{1,1}) \quad (4)$$

Rearranging:

$$0.9U_1(s_{1,1}) - 0.8U_1(s_{1,2}) = -0.14 \quad (3')$$

$$-0.1U_1(s_{1,1}) + 0.9U_1(s_{1,2}) = 0.76 \quad (4')$$

From (4'), we express $U_1(s_{1,1})$ in terms of $U_1(s_{1,2})$:

$$0.1U_1(s_{1,1}) = 0.9U_1(s_{1,2}) - 0.76 \implies U_1(s_{1,1}) = 9U_1(s_{1,2}) - 7.6$$

Substituting this into (3'):

$$0.9(9U_1(s_{1,2}) - 7.6) - 0.8U_1(s_{1,2}) = -0.14$$

$$8.1U_1(s_{1,2}) - 6.84 - 0.8U_1(s_{1,2}) = -0.14$$

$$7.3U_1(s_{1,2}) = 6.7$$

$$U_1(s_{1,2}) = \frac{67}{73} \approx 0.918$$

Substituting back to find $U_1(s_{1,1})$:

$$U_1(s_{1,1}) = 9 \left(\frac{67}{73} \right) - 7.6 = \frac{603 - 554.8}{73} = \frac{48.2}{73} \approx 0.660$$

The final utilities for π_1 are:

$$U_1(s_{1,1}) \approx 0.660, \quad U_1(s_{1,2}) \approx 0.918$$

Now we improve:

$$\begin{aligned} \pi_{i+1}(s) &:= \arg \max_a \sum_{s'} P(s' \mid a, s) \cdot U_i(s') \\ \pi_2(s) &:= \arg \max_a \sum_{s'} P(s' \mid a, s) \cdot U_1(s') \\ \pi_2(s_{1,1}) &:= \arg \max_a \sum_{s'} P(s' \mid a, s_{1,1}) \cdot U_1(s') \\ \pi_2(s_{1,1}) &:= \arg \max_a \{ \\ &\quad 0.8 \cdot 0.918 + 0.1 \cdot -1 + 0.1 \cdot 0.660, & \text{North} \\ &\quad 0.8 \cdot 0.660 + 0.1 \cdot 0.660 + 0.1 \cdot -1, & \text{South} \\ &\quad 0.8 \cdot -1 + 0.1 \cdot 0.918 + 0.1 \cdot 0.660, & \text{East} \\ &\quad 0.8 \cdot 0.660 + 0.1 \cdot 0.918 + 0.1 \cdot 0.660 \} & \text{West} \\ \pi_2(s_{1,1}) &:= \arg \max_a \{ \\ &\quad 0.7004, & \text{North} \\ &\quad 0.494, & \text{South} \\ &\quad -0.6422, & \text{East} \\ &\quad 0.6858 \} & \text{West} \\ \pi_2(s_{1,1}) &:= N & \text{(didn't change)} \end{aligned}$$

Similarly, for $s_{1,2}$:

$$\begin{aligned} \pi_2(s_{1,2}) &:= \arg \max_a \sum_{s'} P(s' \mid a, s_{1,2}) \cdot U_1(s') \\ \pi_2(s_{1,2}) &:= \arg \max_a \{ \\ &\quad 0.8 \cdot 0.918 + 0.1 \cdot 0.918 + 0.1 \cdot 1, & \text{North} \\ &\quad 0.8 \cdot 0.660 + 0.1 \cdot 0.918 + 0.1 \cdot 1, & \text{South} \\ &\quad 0.8 \cdot 1 + 0.1 \cdot 0.918 + 0.1 \cdot 0.660, & \text{East} \\ &\quad 0.8 \cdot 0.918 + 0.1 \cdot 0.918 + 0.1 \cdot 0.660 \} & \text{West} \\ \pi_2(s_{1,2}) &:= \arg \max_a \{ \\ &\quad 0.9262, & \text{North} \\ &\quad 0.7198, & \text{South} \\ &\quad 0.9578, & \text{East} \\ &\quad 0.8922 \} & \text{West} \\ \pi_2(s_{1,2}) &:= E & \text{(didn't change)} \end{aligned}$$

Since $\pi_1(s) = \pi_2(s) \quad \forall s$, we terminate.

2	0.918 →	+1
	0.660 ↑	-1
1	1	2

7 Introduction to Machine Learning

Motivation

- You have data, and you think there is some *model* that can describe how your data would be generated (if it were)
- You do not know that *model*, but you are trying to find something that is as close as possible to that *model* (from some *set* of *models* (hypotheses) you are willing to consider)
- Simplest Example: Learning a *function*
 - There is a **target function** g (We do not know it)
 - We have input-output pairs $X_i = (x, g(x))$
 - The *set* of hypotheses you are willing to consider are in a space $\mathcal{H}$
 - Given a *set* (called training set) of examples X_i , we are trying to find a hypothesis $h \in \mathcal{H}$ such that $h \approx g$

Def 7.1. A **Supervised (Inductive) Learning Problem** is the triple $\langle \mathcal{H}, T, f \rangle$ where:

- $\mathcal{H}$: The *set* of **hypotheses** consisting of *functions* $h : A \rightarrow B$
- $T \subseteq A \times B$: A *set* of **examples** called the **Training Set**, such that for every $a \in A$, there is at most one $b \in B$ with $\langle a, b \rangle \in T$


We assume there is an *unknown function* $f : A \rightarrow B$ called the **Target Function** with $T \subseteq f$

Def 7.2. **Inductive Learning Algorithms** solve *supervised learning problems* by finding a *hypothesis* $h \in \mathcal{H}$ such that $h \sim f$ (for some notion of similarity)


Def 7.3. We call a *supervised learning problems* with *target function* $f : A \rightarrow B$ a **classification problem** if B is *finite*, and call the *members* of B **classes/classifications**. We call it a **regression problem** if $B = \mathbb{R}$

Def 7.4. We call a *hypothesis* $h \in \mathcal{H}$ **consistent with** the unknown target *function* f (on a *set* $T \subseteq \text{dom}(f)$), if it agrees with f (on all *examples* in T)

Example 7.1. Assume we have the following data:




We can choose our *hypotheses space* $\mathcal{H}$ to be the linear *functions*, then the best linear *hypothesis* we can find might look like:




Which is *partially, approximately consistent*.


If $\mathcal{H}$ were the quadratic *functions*, we can find another *partially consistent hypothesis* that might look like:



If we consider Degree-4 polynomials, we get a [consistent hypothesis](#):



If we consider High-degree polynomials we get more (noisy) [consistent hypotheses](#), e.g.



Observation

Whether we can find a [consistent hypothesis](#) for a given [training set](#) depends on the chosen [hypotheses space](#)

Recall

Ockham's-razor: Maximize a combination of [consistency](#) and simplicity.

Def 7.5. We say that a [supervised learning problem](#) is **realizable**, iff there is a [hypothesis](#) $h \in \mathcal{H}$ that is [consistent](#) on the [training set](#) T

Problem

We do not always know whether a given learning problem is [realizable](#), unless we have prior knowledge. A naive solution would be to make $\mathcal{H}$ as large as possible, e.g. all Turing Machines. However, the complexity of the [supervised learning problem](#) is tied to the size of the [hypotheses space](#) and [consistency](#) is not even decidable for general Turing Machines.

Much of the research in machine learning has concentrated on simple [hypotheses spaces](#)

Consistency vs. Simplicity

Following **Ockham's Razor**, among competing **hypotheses**, we prefer those that achieve a good balance between:

- **Consistency** with the observed **training set**
- **Simplicity** of the chosen **hypothesis**

This tradeoff is fundamental in machine learning:

- If the **hypothesis** is *too simple*, it may fail to capture important regularities in the data.
- If the **hypothesis** is *too complex*, it may adapt excessively to noise or accidental patterns in the **training set**.


Def 7.6. We say that a **hypothesis** **underfits** the data if it is too simple to adequately represent the underlying regularities of the unknown **target function**. Such **hypotheses** usually fail to achieve good **consistency** even on the **training set**

Def 7.7. We say that a **hypothesis** **overfits** the data if it becomes excessively adapted to the specific **training set**, including noise or accidental irregularities. Such **hypotheses** may achieve very high **consistency** on the **training data** while generalizing poorly to unseen **examples**

Bias vs. Variance

The tradeoff between simplicity and **consistency** is often called the **Bias-Variance Tradeoff**:

- **High Bias** corresponds to overly simple models that tend to **underfit**
- **High Variance** corresponds to overly complex models that tend to **overfit**



Why is it called Variance?

Imagine repeatedly drawing different random **training sets** from the same underlying process.

For each **training set**, we compute the “best” **hypothesis** $h \in \mathcal{H}$.

This generates a sequence of learned **hypotheses**:

$$h_1, h_2, h_3, \dots$$

If small changes in the **training set** produce large changes in the resulting **hypothesis**, then the learning method has **high variance**.

Thus, high variance means the learned model is highly sensitive to the particular sampled data. This sensitivity often causes **overfitting**, since the model begins adapting to noise and accidental fluctuations in the training data rather than the underlying structure.

Why is it called Bias?

Bias measures the systematic error introduced by restricting ourselves to an overly simple **hypotheses space**.

If the considered **hypotheses** are all “far” from the unknown **target function**, then even the best possible learned model will make systematic errors.

In this case, the learning method has **high bias**.

High bias typically causes **underfitting**, because the model is too simple to capture important regularities in the data.

Observation

In practice, good machine learning models attempt to balance:

simplicity and *consistency*

or equivalently:

bias and variance

Too much simplicity leads to **underfitting**, while too much complexity leads to **overfitting**.

Problem

We want to learn a **hypothesis** that will best predict future data.

Note

This only works if the **training set** is “*representative*” of the underlying process.

Idea

We think of **examples** (seen and unseen) as a **sequence**, and formalize “*representativeness*” via a *stationary assumption* on the **probability distribution**.

Each **example**, before observation, is a **random variable** E_j . The observed value $e_j = \langle x_j, y_j \rangle$ is a **sample** from its **distribution**.

Def 7.8. IID: A **sequence** $E_1, E_2, \dots, E_n$ of **random variables** is **independent and identically distributed (IID)** iff they are:

- **independent**, i.e. $\mathbf{P}(E_j \mid E_{j-1}, E_{j-2}, \dots) = \mathbf{P}(E_j)$ for every j , and
- **identically distributed**, i.e. $\mathbf{P}(E_i) = \mathbf{P}(E_j)$ for every pair of indices i, j .

Def 7.9. Stationary Assumption: We assume the sequence $\varepsilon = (E_t)$ of **examples** is **stationary** iff for every $k \geq 0$ and every pair of indices i, j :

$$\mathbf{P}(E_i, E_{i+1}, \dots, E_{i+k}) = \mathbf{P}(E_j, E_{j+1}, \dots, E_{j+k}).$$

Note

Stationarity does not imply **independence**; **IID** implies **stationarity**.

Def 7.10. In **attribute-based representations**, **examples** are described by

- **attributes**: **functions** on input samples,
- their **values**, and
- **classifications**

Example 7.2. Situations where I will/won’t wait for a table in a Restaurant

Table 3: The Restaurant Dataset (Russell & Norvig)

Example	Attributes										Target
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	WillWait
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0–10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30–60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0–10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10–30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0–10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0–10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0–10	T
X_9	F	T	F	T	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10–30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0–10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30–60	T

The **attributes** are:

- **Alt:** Whether a suitable alternative restaurant is available nearby.
- **Bar:** Whether the restaurant has a comfortable bar area to wait in.
- **Fri:** True if the current day is Friday or Saturday.
- **Hun:** Whether the customer is genuinely hungry.
- **Pat:** The number of patrons currently in the restaurant (*None*, *Some*, or *Full*).
- **Price:** The restaurant's price tier (\$, \$\$, or \$\$\$).
- **Rain:** Whether it is actively raining outside.
- **Res:** Whether the customer made a reservation in advance.
- **Type:** The culinary style of the restaurant (*French*, *Thai*, *Burger*, or *Italian*).
- **Est:** The estimated waiting time in minutes (*0–10*, *10–30*, *30–60*, or *>60*).

Def 7.11. For a Boolean **classification**, we say that an **examples** is **positive** (*T*) or **negative** (*F*) depending on its **class**

Def 7.12. Decision Trees: A **Decision Tree (DT)** for a given **attribute-based representation** is a **tree**, where the **non-leaf nodes** are **labeled** by **attributes**, their **outgoing edges** by **disjoint sets** of **attribute values**, and the **leaf nodes** are **labeled** by the **classifications**.

Decision Trees are one possible representation for **hypotheses**

Example 7.3. Here is the *true* DT for [example 7.2](#):




Figure 6: The true decision tree for the restaurant dataset.


Def 7.13. We call an **attribute** together with a **set** of **attribute values** (an inner node with outgoing edge labels) an **attribute test**.

Note

The **target function** is a **function** $A_1 \times \dots \times A_n \rightarrow C$, where A_i are the domains of the **attributes** and C is the **set** of **classifications**.

Decision Trees can express any **function** of the input **attributes** $\rightsquigarrow \mathcal{H} = A_1 \times \dots \times A_n$

Example 7.4. For Boolean **functions**, a **path** from the **root** to a **leaf** corresponds to a row in a truth table (and the **DT** corresponds to a truth table):



Problem

Trivially, for any **training set** there is a **consistent hypothesis** with one **path** to a **leaf** for each **example**, but it probably won't generalize to new **examples**.

Solution

We should find more *compact* **Decision Trees**.

Idea

A *good* **attribute** splits the **examples** into **subsets** that are *ideally* all positive or all negative.

A Mathematical Vernacular and Induction

A.1 Mathematical Language

Def A.1. Natural Language: A natural language is any form of spoken or signed means of communication that has evolved naturally in humans through use and repetition without conscious planning or premeditation.

Def A.2. Domain of Discourse: The [set](#) of entities, objects, or concepts that a language, discussion, or formal system refers to. It specifies the scope within which terms, statements, and reasoning apply.

Def A.3. Jargon (Terminology): A jargon (or terminology) is a [set](#) of specialized words or phrases (called technical terms or just terms) relating to concepts from a particular [domain of discourse](#).

Def A.4. Technical Language: A technical language is a [natural language](#) extended by a [terminology](#) and (possibly) special idioms, discourse markers, and notations.

Def A.5. Stylized Language: Mathematicians use a stylized language that uses formulae to represent mathematical objects, uses math idioms for special situations (e.g. "iff", "hence", "let... be...", then..."), and classifies statements by role (e.g., definition, Lemma, Theorem, Proof, Example).

Def A.6. Mathematical Vernacular: A [technical language](#) that you need to master to be successful when moving to a new environment ([symbolic AI](#)).

Def A.7. Formulae: A mathematical formula can be a **mathematical statement** (clause) which can be true or false (e.g. $x > 5, 3 + 5 = 7$), or a **mathematical object** (e.g. $3, n, x^2 + y^2 + z^2, \int_1^0 x^{3/2} dx$)

Observation: [mathematical vernacular](#) loves to name [objects/statements](#) for precise references. For example: Let $p = 3x^2 + 7x$, then $p^3 + 17p + 1$ is ...

Moreover, [mathematical vernacular](#) aggregates [objects/statements](#) for cognitive efficiency, examples:

- $a \in S, b \in S$, and $c \in S \rightsquigarrow a, b, c \in S$ ([object](#) aggregation)
- $i \geq 0$ and $i \leq n \rightsquigarrow 0 \leq i \leq n$ ([statement](#) aggregation)

Def A.8. Sequence: [mathematical vernacular](#) uses the concept of sequences instead of lists. Sequences are usually finite, but can be infinite as well.

Def A.9. Ellipses: [mathematical vernacular](#) uses ellipses ($\dots$) as a constructor for [sequences](#). The meaning of an ellipsis is usually considered "obvious" and left for interpretation by the reader.

[ellipses](#) allow to write down large [objects](#) easily, examples:

- $1, \dots, n \rightsquigarrow$ the [sequence](#) of natural numbers between 1 and n in order.
- $1, 4, 9, 16, \dots \rightsquigarrow$ the [sequence](#) of squares in orders.
- $e_1, \dots, e_n \rightsquigarrow$ a [sequence](#) of [objects](#) e_i for $1 \leq i \leq n$

[Mathematical statements](#) come in different epistemic varieties:

- **Definitions:** [statements](#) that introduce new global identifier for important [objects](#)
- **Assertions:** [statements](#) that state properties of [mathematical objects](#)
- **Examples:** [statements](#) that exhibit a witness for some property
- **Axioms:** [statements](#) that characterize the [objects](#) of a certain [domain of discourse](#) or theory. An axiom (postulate) is a statement about [mathematical objects](#) that we *assume* to be true.

[Mathematical assertions](#) are pragmatically classified into categories:

- **proofs** are arguments that justify the truth of [statements](#) beyond any doubt. [Proofs](#) are not really [statements](#), but we sometimes treat them together.
- **proposition** is a [statement](#) which is interesting in its own right.
- **lemma** is an easily proved [statement](#) which is helpful for proving other [propositions](#) and [theorems](#), but is usually not particularly interesting in its own right. A **lemma** is an intermediate [theorem](#) that serves as part of a [proof](#) of a larger [theorem](#).

- **theorem** is a more important [statement](#) than a [proposition](#) which says something definitive on the subject, and often takes more effort to prove than a [proposition](#) or [lemma](#). A [theorem](#) is a [statement](#) about [mathematical object](#) that we *know* to be true.
- **corollary** is a quick consequence of a [proposition](#) or [theorem](#) that was proven recently. It is a [theorem](#) that follows directly from another [theorem](#).
- **conjecture** is a [statement](#) that is thought to be provable, but has not been yet. A [conjecture](#) that is proven is a [theorem](#)

[Mathematical statement](#) use abbreviations, examples:

- $\wedge$ and $\vee$ for "and" and "or",
- $\neg$ for "not"
- $\forall x.P$ ($\forall x \in S.P$) for "condition P holds for all x (in S)"
- $\exists x.P$ ($\exists x \in S.P$) for "there exists an x (in S) such that the condition P holds"
- $\nexists x.P$ ($\nexists x \in S.P$) for "there exists no x (in S) such that the condition P holds"
- $\exists^1 x.P$ ($\exists^1 x \in S.P$) for "there exists one and only x (in S) such that the condition P holds"
- "*iff*" for "if and only if", symbolized by $\Leftrightarrow$
- $\Rightarrow$ for "implies"; we can read $A \Rightarrow B$ as "if A then B "

Def A.10. Declaration: In [mathematical vernacular](#) we call a clause that introduce new identifiers together with some properties a declaration. For example Let $\epsilon, \delta > 0 \dots$

Def A.11. Scope: The scope of an identifier is the part of an expression where the reference is valid; that is, where the identifier can be used to refer. In the other parts, the identifier may refer to a different entity, or to nothing at all (unbound)

Observations: [definition](#) have *global scope*, [declarations](#) have *local scope*.

Def A.12. Definiendum: The definiendum is the symbol, expression, or term being introduced by a [definition](#). It is the newly named [object](#) or concept.

Def A.13. Definiens: The definiens is the symbol, expression, or construction that specifies the meaning of the [defineiendum](#). It explains what the [defineiendum](#) denotes.

Mathematics has developed various forms of [definition](#) (definition schemata). A **simple definition** introduces a name (the [defineiendum](#)) for a compound [object](#) or concept (the [definiens](#)). The [defineiendum](#) must be new, i.e. may not have been used for anything else, in particular, the [defineiendum](#) may not occur in the [definiens](#). We use the symbols $:=$ (and the inverse $=:$) to write [simple definitions](#). For example, we can give the natural number 2 the name ϕ . In a [formula](#) we write this as $\phi := 2$ or $2 =: \phi$.

On the other hand, in a **pattern definition** the [defineiendum](#) is the operator or [relation](#) applied to n distinct variables -called pattern variables- $v_1, \dots, v_n$ as arguments, and the [definiens](#) is an expression in these variables. When the new operator is applied to arguments $a_1, \dots, a_n$, then its value is the [definiens](#) expression where the v_i are replaced by the a_i . We use the symbol $:=$ for operator definitions and $:\Leftrightarrow$ for [relation](#) definitions. For example, the following is a [pattern definition](#) for the [set intersection operator](#) $\cap$:

$$A \cap B := \{x \mid x \in A \wedge x \in B\}$$

The pattern variables are A and B , and with this definition we have e.g. $\Phi \cap \Phi = \{x \mid x \in \Phi \wedge x \in \Phi\}$.

An **implicit definition** (definition by description) is a clause or expression A , such that we can prove "there is exactly one n such that A " ($\exists^1 n.A$), where n is a new name - the [defineiendum](#). If such a unique existence proof exists, we call n *well-defined*. For example, $\forall x.x \notin \Phi$ is an [implicit definition](#) for the empty set Φ (instead of explicitly saying $\Phi := \{\}$). We can prove unique existence of Φ by just exhibiting $\{\}$ and showing that any other [set](#) S with $\forall x.x \notin S$ we have $S \equiv \Phi$. S cannot have elements, so it has the same elements as Φ , and thus $S \equiv \Phi$.

Here is another [implicit definition](#): The exponential function is that function $f : \mathbb{R} \rightarrow \mathbb{R}$ with $f' = f$ and $f(0) = 1$. Here A is the clause " $f' = f$ and $f(0) = 1$ ". Well-definedness is mathematically non-trivial.

Mathematics uses examples and counterexamples to support understanding a property P . Examples give us a sense of the extent of P and counterexamples help delineate the border of P .

Def A.14. Example: An example E is a [mathematical statement](#) that consists of:

- symbol p for the **exemplandum** (plural: exemplanda): the property to be exemplified,
- the **exemplans** (plural: exemplantia): an expression A denoting a **mathematical object** that acts as witness object for the property p , and
- (optionally) a **justification** of E , i.e. a **proof** π that $p(A)$ holds in the current context.

Def A.15. Counter Example: An **example** for the complement of the property p where π is a **proof** that $p(A)$ does not hold.

Def A.16. Running Example: An **example** that is recalled time and again during an argument, applying the newly discovered knowledge to it, presumably to show how things work.

For instance, the following **mathematical statement** is a **mathematical example**:

The identity function on $\mathbb{R}$ is *continuous*.

- the **exemplandum** is "*continuous*",
- the **exemplans** A is "The identity function on $\mathbb{R}$ ", and
- the **justification** π , a **proof** of continuity $\text{Id}_{\mathbb{R}}$: Let $\epsilon > 0$, then we choose $\delta := \epsilon \cdots$ is omitted.

We can use **examples** for $\forall$ -statements as well. Let $B = \forall x.A$ be a **mathematical statement**, then we call an **example/counter example** for B whose **exemplandum** is the property " A holds for x " an **example/counter example** for B . For example, 3 is an **example** for the **statement** $B = \forall n.\text{odd}(n)$.

Note

An **example** (or even several) for a **mathematical statement** B do not constitute a **proof** for B . On the other hand, one **counter example** for a **mathematical statement** B does show that $\neg B$ is true ($\exists x.\neg A \equiv \neg(\forall x.A)$).

Def A.17. Syllogism: A syllogism (**proof** method) is an argument that applies **deductive reasoning** to arrive at a **conclusion** based on two (or more) **propositions** that are **asserted** or assumed to be true. A **syllogistic system** is simply a **set** of **syllogisms**

Note

Today we usually reserve the term **syllogism** to informal arguments, **formal** arguments are called **inference** rules.

In a **proof by contradiction**, we make an assumption ($\neg A$) to the contrary of what we want to prove (A), and then we show that the assumption leads to a contradiction and therefore must be false. That lets us to **conclude** that $(\neg\neg A)$ must be true, and thus (A) .

If we want to prove a **statement** of the form "*If A , then B* ", then we do that by:

- assuming A and proving B from that.
- after this subproof, we may not use A anymore.

We call this **proof** method a **proof by local hypothesis**

If we know a general rule of the form "*If A then B* ", and we also know a specific instance A' , which arises from A by replacing its variables with concrete values. Then by **chaining** we can **conclude** the corresponding instance B' , obtained from B by the same variable replacement.

Example.

1. We know: "Socrates is a human." This is A' .
2. We also know a general rule:

$$\forall x (x \text{ is human} \Rightarrow x \text{ is mortal}).$$

This is the "If A then B " statement, where

$$A(x) = (x \text{ is human}), \quad B(x) = (x \text{ is mortal}).$$

3. Substitute $x = \text{Socrates}$:

$$A(x) \rightarrow \text{Socrates is human}, \quad B(x) \rightarrow \text{Socrates is mortal}.$$

4. Since “Socrates is human” is true, [chaining](#) gives us:

Socrates is mortal.

If we want to prove “ A for all x ”, then we prove A without regard for x , then argue something like “as x was chosen arbitrarily when we proved A , we know A for every x ”

If we know that one of the cases $A_1, \dots, A_k$ must always hold, then we use **proof by cases** to show that “if A_1 then C ” and “if A_2 then C ” and so on $\forall i : 1 \leq i \leq k$, hence we can [conclude](#) that C holds. The following is an example:

Lemma A.1. For all rational numbers a and b , if $ab = 0$, then $a = 0$ or $b = 0$.

Proof.

1. Let $a, b \in \mathbb{Q}$ and assume $ab = 0$.
2. Obviously, either $a = 0$ or $a \neq 0$.
3. We prove that $a = 0$ or $b = 0$ by the two cases induced.
 4. Case $a = 0$:
 - 4.1 Then the conclusion of the lemma is trivially true, and there is nothing to prove.
 5. Case $a \neq 0$:
 - 5.1 We can multiply both sides of the equation $ab = 0$ by $\frac{1}{a}$ and obtain

$$\frac{1}{a} \cdot ab = \frac{1}{a} \cdot 0.$$
 - 5.2 Reducing the fractions gives $b = 0$.
6. In both cases, either $a = 0$ or $b = 0$, so the statement holds. □

Def A.18. Without Loss of Generality (WLOG): An [inference](#) principle used in [proofs](#) when a [statement](#) $Q(x)$ must be shown for all cases of x , but the different cases are equivalent by symmetry or trivial reduction. [Formally](#), if the domain of x can be partitioned into cases $A_1, \dots, A_k$ and it can be shown that:

- proving $Q(x)$ under one representative case A_i suffices, because
- for every other case A_j there is either a trivial [proof](#) of $Q(x)$ or a direct reduction to the representative case,

then we may assume A_i holds “without loss of generality”. This allows us to simplify the argument by only proving the representative case.

For example, if we want to prove the property $Q(p)$ for all primes p . We know that every prime is either 2 (the only even prime) or an odd prime. If the case $p = 2$ is easy, and the general odd prime case is the only real work, then we can say:

[WLOG](#), assume p is odd

This does not exclude the even case; it means we have already checked that the even case adds no new difficulty.

In [mathematical vernacular](#), different alphabets are used to convey precise meanings. [Table 4](#) presents the Greek letters that you are expected to read, recognize, and write. In addition, [table 5](#) shows the curly Roman and Fraktur letters, which are also standard in mathematical writing. Becoming fluent with these alphabets is an essential part of mathematical literacy.

α	A	alpha	β	B	beta	γ	Γ	gamma
δ	Δ	delta	ϵ	E	epsilon	ζ	Z	zeta
η	H	eta	θ, ϑ	Θ	theta	ι	I	iota
κ	K	kappa	λ	Λ	lambda	μ	M	mu
ν	N	nu	ξ	Ξ	xi	$\omicron$	O	omicron
π, ϖ	Π	pi	ρ	P	rho	σ	Σ	sigma
τ	T	tau	υ	Υ	upsilon	φ, ϕ	Φ	phi
χ	X	chi	ψ	Ψ	psi	ω	Ω	omega

Table 4: Greek Alphabet

Roman	$\mathcal{A}, \mathcal{B}, \mathcal{C}, \mathcal{D}, \mathcal{E}, \mathcal{F}, \mathcal{G}, \mathcal{H}, \mathcal{I}, \mathcal{J}, \mathcal{K}, \mathcal{L}, \mathcal{M}, \mathcal{N}, \mathcal{O}, \mathcal{P}, \mathcal{Q}, \mathcal{R}, \mathcal{S}, \mathcal{T}, \mathcal{U}, \mathcal{V}, \mathcal{W}, \mathcal{X}, \mathcal{Y}, \mathcal{Z}$
Fraktur	$\mathfrak{A}, \mathfrak{B}, \mathfrak{C}, \mathfrak{D}, \mathfrak{E}, \mathfrak{F}, \mathfrak{G}, \mathfrak{H}, \mathfrak{I}, \mathfrak{J}, \mathfrak{K}, \mathfrak{L}, \mathfrak{M}, \mathfrak{N}, \mathfrak{O}, \mathfrak{P}, \mathfrak{Q}, \mathfrak{R}, \mathfrak{S}, \mathfrak{T}, \mathfrak{U}, \mathfrak{V}, \mathfrak{W}, \mathfrak{X}, \mathfrak{Y}, \mathfrak{Z}$

Table 5: Roman and Fraktur Letters

A.2 Unary Natural Numbers

Def A.19. Unary Natural Numbers: Terms generated by the following: $u ::= 0 \mid s(u)$ where 0 denotes zero and $s(u)$ denotes the successor of u . i.e. $\{0, s(0), s(s(0)), s(s(s(0))), \dots\}$.

Def A.20. Piano Axioms: The following set of **axioms** are called the Piano axioms. They can be used to *reason* about natural numbers:

- **P1.** 0 is a **unary natural number**.
- **P2.** Every **unary natural number** has a successor that is **unary natural number** and that is different from it.
- **P3.** 0 is not a successor of any **unary natural number**.
- **P4.** Different **unary natural numbers** have different successors.
- **P5. Induction Axiom.** Every **unary natural number** possess a property P if
 - 0 has the property P , and
 - the successor of every **unary natural number** that has the property P also possesses property P

More formally, we call the **set** of **unary natural numbers** $\mathbb{N}_1$, and we say $P(n)$ if a **unary natural number** $n \in \mathbb{N}_1$ has a property P . Therefore, we express the axioms as follows:

- $0 \in \mathbb{N}_1$
- $\forall n \in \mathbb{N}_1. s(n) \in \mathbb{N}_1 \wedge n \neq s(n)$
- $\neg(\exists n \in \mathbb{N}_1. 0 = s(n))$
- $\forall n \in \mathbb{N}_1. \forall m \in \mathbb{N}_1. n \neq m \Rightarrow s(n) \neq s(m)$
- $\forall P. P(0) \wedge (\forall n \in \mathbb{N}_1. P(n) \Rightarrow P(s(n))) \Rightarrow (\forall m \in \mathbb{N}_1. P(m))$

Theorem A.2. Every **unary natural number** is either zero (0) or the successor of a **unary natural number**.

Proof. We make use of the induction axiom P5: We use the property P of "being zero or a successor" and prove the statement by convincing ourselves of the prerequisites of:

1. 0 is zero so 0 is "zero or a successor".
2. Let n be arbitrary **unary natural number** that is "zero or a successor".
3. Then its successor is a "successor", so the successor of n is also "zero or a successor".
4. Since we have taken n arbitrary we have shown that for any n , its successor has property P .
5. Property P holds for all **unary natural number** by P5, so we have proven the assertion.

□

Def A.21. Mathematical Induction: A method for proving a **statement** $P(n)$ is true for every natural number n . A **proof** by mathematical induction consists of two steps:

- The **base case**: Prove that $P(0)$ is true.
- The **induction step**: Prove that for every n , if $P(n)$ is true then $P(n+1)$ is also true.

The hypothesis in the **induction step**, that $P(n)$ is true for a particular n , is called the **induction hypothesis**

Def A.22. The addition "function": We define the addition operation $+$ procedurally:

- Adding 0 to a number does not change it, expressed as: $n + 0 = n$.
- Adding m to the successor of n yields the successor of $m + n$, expressed as: $m + s(n) = s(m + n)$.

Theorem A.3. For all **unary natural number** n, m, l , we have $(n + m) + l = n + (m + l)$.

Proof. We prove this by induction on l .

Let $P(l)$ be the property:

$$(n + m) + l = n + (m + l).$$

- **base case:**

$$(n + m) + 0 = (n + m) = n + (m + 0)$$

- **step case:** Assume $(n + m) + l = n + (m + l)$ for an arbitrary l . We must show

$$(n + m) + s(l) = n + (m + s(l)).$$

$$\begin{aligned} (n + m) + s(l) &= s((n + m) + l) \\ &\stackrel{\text{IH}}{=} s(n + (m + l)) \\ &= n + s(m + l) \\ &= n + (m + s(l)). \end{aligned}$$

Hence the property holds for $s(l)$.

□

Def A.23. Multiplication: The **unary multiplication** operation can be defined by the equations:

- $n \cdot 0 = 0$, and
- $n \cdot s(m) = n + (n \cdot m)$

Def A.24. Exponentiation: The **unary exponentiation** operation can be defined by the equations:

- $\text{exp}(n, 0) = s(0)$, and
- $\text{exp}(n, s(m)) = n \cdot \text{exp}(n, m)$

Def A.25. Summation: The **unary summation** operation can be defined by the equations:

- $\bigoplus_{i=0}^0 (n_i) = 0$, and
- $\bigoplus_{i=0}^{s(m)} (n_i) = n_{s(m)} \bigoplus \bigoplus_{i=0}^m (n_i)$

Example:

$$\begin{aligned} \bigoplus_{i=0}^{s(s(0))} (n_i) &= n_{s(s(0))} \bigoplus \bigoplus_{i=0}^{s(0)} (n_i) \\ &= n_{s(s(0))} \bigoplus n_{s(0)} \bigoplus \bigoplus_{i=0}^0 (n_i) \\ &= n_{s(s(0))} \bigoplus n_{s(0)} \bigoplus 0 \end{aligned}$$

Def A.26. Product: The **unary product** operation can be defined by the equations:

- $\bigodot_{i=0}^0 (n_i) = s(0)$, and
- $\bigodot_{i=0}^{s(m)} (n_i) = n_{s(m)} \bigodot \bigodot_{i=0}^m (n_i)$

B Set Theory and Foundations of Mathematics

B.1 Sets

Def B.1. Set: A collection of different things (called **elements** or **members** of the set).

We can represent **sets** by listing the **elements** within curly brackets, e.g. $\{a, b, c\}$ or by describing the **elements** via a property, e.g. $\{x \mid x \bmod 2 = 0\}$ (the **set** of even numbers). We can state **element-hood** ($a \in S$) or not ($b \notin S$)

Note

Learn to distinguish between **objects** and their representations. ($\{a, b, c\}$ and $\{b, a, a, c\}$ are different representations of the same **set**)

Def B.2. Set Equality: $A \equiv B := \forall x. x \in A \Leftrightarrow x \in B$

Def B.3. Subset: $A \subseteq B := \forall x. x \in A \Rightarrow x \in B$

Def B.4. Proper Subset: $A \subset B := (A \subseteq B) \wedge (A \neq B)$

Def B.5. Super Set: $A \supseteq B := \forall x. x \in B \Rightarrow x \in A$

Def B.6. Proper Superset: $A \supset B := (A \supseteq B) \wedge (A \neq B)$

Def B.7. Set Union: $A \cup B := \{x \mid x \in A \vee x \in B\}$

Def B.8. Set Intersection: $A \cap B := \{x \mid x \in A \wedge x \in B\}$

Def B.9. Disjoint: Two **sets** A, B are **disjoint** iff $A \cap B = \emptyset$

Def B.10. Set Difference: $A \setminus B := \{x \mid x \in A \wedge x \notin B\}$

Def B.11. Power Set: $\mathcal{P}(A) := \{S \mid S \subseteq A\}$ (the set of all subsets)

Note

Set operators such as **union** ($\cup$), **intersection** ($\cap$), and **set difference** ($\setminus$) always return a **set**. For example, consider

$$\{a, b\} \cap \{b, c\} = \{b\}.$$

The result $\{b\}$ is a **set**.

Now observe:

- $\{b\} \notin \{a, b\}$, since the **elements** of $\{a, b\}$ are a and b , not the **set** $\{b\}$.
- $\{b\} \notin \{b, c\}$, since the **elements** of $\{b, c\}$ are b and c , not the **set** $\{b\}$.
- $\{b\} \subseteq \{a, b\}$, because every **element** of $\{b\}$ (namely b) is also in $\{a, b\}$.
- $\{b\} \subseteq \{b, c\}$, for the same reason.

Thus, the outcome of a **set** operation should always be compared to other **sets** using $\subseteq$ (**subset relation**), not $\in$ (**element relation**).

Def B.12. Cartesian Product: $A \times B := \{(a, b) \mid a \in A \wedge b \in B\}$, (a, b) is a *pair*

Def B.13. Family: A **set** of **sets**

Def B.14. Topology: A **family** of open **sets**

Def B.15. Indexing Function: Let I and X be **sets**. A function $f : I \rightarrow X$ is called an indexing function. The set I is called the index set. For each $i \in I$, we define $x_i := f(i)$, here i is called the index (or parameter) of x_i .

Def B.16. Indexed Family: Given an **indexing function** $f : I \rightarrow X$, the **set** of values $f(I) = \{f(i) : i \in I\}$ is called an indexed **family**. It is often written as $(x_i)_{i \in I}$, $\langle x_i \rangle_{i \in I}$, or $\{x_i\}_{i \in I}$.

Example: Consider **index set** $I := \{1, 2, 3\}$. Consider **indexing function** f defined as:

$$f(1) = \{a, b\} \quad f(2) = \{c\} \quad f(3) = \{d, f\}$$

the **set** of values $f(I) = \{f(1), f(2), f(3)\}$ is the **indexed family**:

$$(X_i)_{i \in I} = f(I) = \{\{a, b\}, \{c\}, \{d, f\}\}$$

where $x_i = f(i)$, i.e. $x_1 = f(1) = \{a, b\}, \dots$

Def B.17. Union over a Family: Let $(X_i)_{i \in I}$ be an **indexed family**, then $\bigcup_{i \in I} X_i := \{x \mid \exists i \in I. x \in X_i\}$

Def B.18. Intersection over a Family: Let $(X_i)_{i \in I}$ be an **indexed family**, then $\bigcap_{i \in I} X_i := \{x \mid \forall i \in I. x \in X_i\}$

Def B.19. n -fold Cartesian Product: $\prod_{i=1}^n X_i := X_1 \times \dots \times X_n := \{\langle x_1, \dots, x_n \rangle \mid \forall i. 1 \leq i \leq n \Rightarrow x_i \in X_i\}$ where $\langle x_1, \dots, x_n \rangle$ is called an n -tuple.

Example: consider the **indexed family** $(X_i)_{i \in \{1,2,3\}} = \{\{a, b\}, \{c\}, \{d, f\}\}$. $X_1 \times X_2 \times X_3 = \{\langle a, c, d \rangle, \langle a, c, f \rangle, \langle b, c, d \rangle, \langle b, c, f \rangle\}$.

Def B.20. n -dim Cartesian Space: $X^n := \{\langle x_1, \dots, x_n \rangle \mid 1 \leq i \leq n \Rightarrow x_i \in X\}$ where $\langle x_1, \dots, x_n \rangle$ is called a vector.

Example: Let $X = \{1, 2\}$, then $X^2 = \{\langle 1, 1 \rangle, \langle 1, 2 \rangle, \langle 2, 1 \rangle, \langle 2, 2 \rangle\}$. Another example is the usual 2D plane: $\mathbb{R}^2 := \{\langle x, y \rangle \mid x \in \mathbb{R}, y \in \mathbb{R}\}$

Def B.21. Size of a Set: The size $\Gamma(A)$ of a **set** A is the number of **elements** in A .

Note

- $\Gamma(A \cup B) \leq \Gamma(A) + \Gamma(B)$
- $\Gamma(A \cap B) \leq \min(\Gamma(A), \Gamma(B))$
- $\Gamma(A \times B) = \Gamma(A) \cdot \Gamma(B)$

Def B.22. Pairwise Disjoint: A **family** of **sets** is called **pairwise disjoint** or **mutually disjoint**, if any two of those **sets** are **disjoint**.

Def B.23. Multiset: A **multiset** (or **bag**) is a generalization of a **set** where **elements** can appear more than once. Formally, a multiset $\mathcal{M}$ over a domain D is a pair (D, m) , where $m : D \rightarrow \mathbb{N}$ is a function mapping each **element** $x \in D$ to its **multiplicity** (the number of occurrences).

Def B.24. Multiset Membership: $x \in \mathcal{M} \equiv m(x) > 0$

Def B.25. Multiset Size: The size of a multiset $\mathcal{M} = (D, m)$, denoted $\Gamma(\mathcal{M})$, is the sum of the multiplicities of all its **elements**: $\Gamma(\mathcal{M}) = \sum_{x \in D} m(x)$.

Def B.26. Multiset Sum (Union): The sum of two multisets $\mathcal{M}_1 = (D, m_1)$ and $\mathcal{M}_2 = (D, m_2)$, denoted $\mathcal{M}_1 \uplus \mathcal{M}_2$, is a multiset where the multiplicity of each **element** is the sum of its multiplicities: $m_{\uplus}(x) = m_1(x) + m_2(x)$.

Def B.27. Multiset Intersection: The **intersection** of two multisets $\mathcal{M}_1 = (D, m_1)$ and $\mathcal{M}_2 = (D, m_2)$, denoted $\mathcal{M}_1 \cap \mathcal{M}_2$, is a multiset where the multiplicity of each **element** is the minimum of its multiplicities in the two multisets: $m_{\cap}(x) = \min(m_1(x), m_2(x))$.

Def B.28. Multiset Equality: $\mathcal{M}_1 \equiv \mathcal{M}_2 \equiv \forall x \in D. m_1(x) = m_2(x)$

B.2 Relations

Def B.29. Relation: $R \subseteq A \times B$ is a (binary) relation between A and B

Def B.30. Relation on: a relation $R \subseteq A \times B$ where $A = B$ is called a *relation on* A

Def B.31. Total: A relation $R \subseteq A \times B$ is called *total* (*left total*) iff $\forall x \in A. \exists y \in B. (x, y) \in R$

Def B.32. Converse Relation: $R^{-1} \subseteq B \times A := \{(y, x) \mid (x, y) \in R\}$ is the converse relation of R

Def B.33. Relation Composition: The composition of $R \subseteq A \times B$ and $S \subseteq B \times C$ is defined as $R \circ S := \{(a, c) \in A \times C \mid \exists b \in B. (a, b) \in R \wedge (b, c) \in S\}$

A relation on A : $R \subseteq A \times A$ is called:

- **reflexive** on A , iff $\forall a \in A. (a, a) \in R$
- **irreflexive** on A , iff $\forall a \in A. (a, a) \notin R$
- **symmetric** on A , iff $\forall a, b \in A. (a, b) \in R \Rightarrow (b, a) \in R$
- **asymmetric** on A , iff $\forall a, b \in A. (a, b) \in R \Rightarrow (b, a) \notin R$
- **antisymmetric** on A , iff $\forall a, b \in A. (a, b) \in R \wedge (b, a) \in R \Rightarrow a = b$
- **transitive** on A , iff $\forall a, b, c \in A. (a, b) \in R \wedge (b, c) \in R \Rightarrow (a, c) \in R$
- **equivalence relation** on A , iff R is reflexive symmetric, and transitive

Def B.34. Equality Relation: The equality relation is an equivalence relation on any set.

Def B.35. Divides Relation: The *divides* relation on the integers is defined as $\mid \subseteq \mathbb{Z} \times \mathbb{Z}$, where $\mid = \{(a, b) \in \mathbb{Z} \times \mathbb{Z} \mid \exists k \in \mathbb{Z} : b = a \cdot k\}$. We write $a \mid b$ to denote $(a, b) \in \mid$ (read as "*a divides b*").

For example: $5 \mid 10$ because $\exists k : 10 = 5 \cdot k$ where $k = 2$.

Def B.36. Congruence Modulo: For a fixed $n \in \mathbb{N}$ with $n \geq 1$, the *congruence modulo* n relation on the integers is defined as

$$\equiv_n \subseteq \mathbb{Z} \times \mathbb{Z}, \quad \equiv_n = \{(a, b) \in \mathbb{Z} \times \mathbb{Z} \mid n \mid (a - b)\}.$$

We write $a \equiv b \pmod{n}$ to denote $(a, b) \in \equiv_n$ (read as "*a is congruent to b modulo n*").

Examples:

- $12 \equiv 2 \pmod{5}$ because $5 \mid 10$
- $7 \equiv 22 \pmod{5}$ because $5 \mid 15$
- $9 \equiv 9 \pmod{5}$ because $5 \mid 0$

Note

In the definition of congruence modulo n , we do *not* take the absolute value of the difference:

$$a \equiv b \pmod{n} \iff n \mid (a - b).$$

Since the divisor n can be multiplied by a negative integer, the sign of $(a - b)$ does not matter. For example, $2 \equiv 12 \pmod{5}$ because $2 - 12 = -10$ and $-10 = 5 \cdot (-2)$, so $5 \mid (2 - 12)$.

Exercise: Which of the following binary relations on integers are equivalence relations?

1. Φ . It is not reflexive because we must have $(a, a) \in R$ for all integers a . But Φ has no pairs at all. $\rightarrow$ NO.
2. $\mathbb{Z} \times \mathbb{Z}$
 - It is reflexive since $(a, a) \in R \quad \forall a$
 - It is symmetric because if $(a, b) \in R$ then (b, a) also is
 - It is transitive because every possible pair is in it anyway
 - $\rightarrow$ YES.
3. $m \leq n$
 - It is reflexive since $m \leq m$

- It is not **symmetric**, e.g. $2 \leq 3$ but not $3 \leq 2$

- $\rightarrow$ NO

4. $\exists z \in \mathbb{Z} : m + z = n$

- This is actually the same as $\mathbb{Z} \times \mathbb{Z}$ because for any integers m, n , we can always pick $z = n - m$. $\rightarrow$ YES.

5. $\exists z \in \mathbb{Z} : m \cdot z = n$

- This is the same as $m \mid n$. It is **reflexive** ($m \cdot 1 = m$). However, it is not **symmetric** ($2 \mid 4$ but $4 \nmid 2$). $\rightarrow$ NO.

6. $5 \mid (m - 2)$

- This means m is **congruent** to 2 modulo 5.
- It is **reflexive** ($5 \mid 0$)
- It is **symmetric** (negative or positive does not matter)
- It is **transitive**, if $5 \mid m - n$ and $5 \mid n - k$, then $5 \mid (m - n) + (n - k) = m - k$
- $\rightarrow$ YES.

7. $m \mid n$. It is **reflexive** but not **symmetric**. $\rightarrow$ NO.

Def B.37. Reflexive Closure: The **reflexive closure** of a **relation** R on a **set** A is the **relation** $R \cup \Delta$ where $\Delta = \{\langle a, a \rangle \mid a \in A\}$ is the identity relation. It is defined as:

$$R_r = R \cup \{\langle a, a \rangle \mid a \in A\}$$

Note

Alternative notation: $R_r = R \cup R^0$. Since R^n is the **set** of pairs that are connected by a chain of exactly n hops "reachable in n steps", we can think of R^0 as pairs that are connected by a chain of exactly 0 hops "reachable in 0 steps", i.e. the identity relation.

For example, let $A = \{1, 2, 3\}$, $R = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle\}$, then $R_r = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle, \langle 1, 1 \rangle, \langle 2, 2 \rangle, \langle 3, 3 \rangle\}$

Def B.38. Symmetric Closure: The **symmetric closure** of a **relation** R on a **set** A is the **union** of R and its inverse **relation** R^{-1} . It is defined as:

$$R_s = R \cup R^{-1} = R \cup \{\langle b, a \rangle \mid \langle a, b \rangle \in R\}$$

e.g. $R_s = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle, \langle 2, 1 \rangle, \langle 3, 2 \rangle\}$

Def B.39. Transitive Closure: The **transitive closure** R^+ of a **relation** R on a **set** A is the **smallest transitive relation** containing R . It is defined as:

$$R^+ = \bigcup_{n=1}^{\infty} R^n$$

For example, let $A = \{1, 2, 3, 4\}$, $R = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle, \langle 3, 4 \rangle\}$

- $R^1 = R$
- $R^2 = R \circ R = \{\langle 1, 3 \rangle, \langle 2, 4 \rangle\}$
- $R^3 = R \circ R \circ R = \{\langle 1, 4 \rangle\}$
- No further compositions yields new pairs.

$\rightsquigarrow R^+ = R \cup R^2 \cup R^3 = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle, \langle 3, 4 \rangle, \langle 1, 3 \rangle, \langle 2, 4 \rangle, \langle 1, 4 \rangle\}$

Def B.40. Reflexive Transitive Closure: The **reflexive transitive closure** (A.K.A ??) of a **relation** R on a **set** A is the **smallest relation** that is both **reflexive** and **transitive**. It is defined as:

$$R^* = R^0 \cup R^+ = \bigcup_{n=0}^{\infty} R^n$$

Example: let $A = \{1, 2, 3\}$, $R = \{\langle 1, 2 \rangle, \langle 2, 3 \rangle\}$.

- R^0 is the identity relation on $A = \{\langle 1, 1 \rangle, \langle 2, 2 \rangle, \langle 3, 3 \rangle\}$
- $R^1 = R$
- $R^2 = \{\langle 1, 3 \rangle\}$
- No new pairs can be introduced

$$\rightsquigarrow R^* = \{\langle 1, 1 \rangle, \langle 2, 2 \rangle, \langle 3, 3 \rangle\} \cup \{\langle 1, 2 \rangle, \langle 2, 3 \rangle\} \cup \{\langle 1, 3 \rangle\} = \{\langle 1, 1 \rangle, \langle 2, 2 \rangle, \langle 3, 3 \rangle, \langle 1, 2 \rangle, \langle 2, 3 \rangle, \langle 1, 3 \rangle\}$$

Def B.41. Parital Ordering (non-strict): A relation $R \subseteq A \times A$ is called a **partial ordering** on A , iff R is **reflexive**, **antisymmetric**, and **transitive** on A .

Def B.42. Strict Partial Ordering: A relation $R \subseteq A \times A$ is called a **strict partial ordering** on A , iff R is **irreflexive**, and **transitive** on A .

Note

We usually write **strict partial ordering** with asymmetric symbols like $\prec$ and **non-strict ones** by adding a line that reminds us of equality ($\preceq$).

Def B.43. Comparable: In a **non-strict partial ordering**, two elements a, b are comparable if either $a \preceq b$ or $b \preceq a$. If neither holds, they are incomparable.

Def B.44. Linear Order: A **partial ordering** is called linear (or total) on A , iff all **elements** in A are **comparable**, i.e. if $(x, y) \in R$ or $(y, x) \in R$ for all $x, y \in A$.

For example, the $\leq$ relation is a **linear order** on $\mathbb{N}$. However, the "divides" ($\mid$) relation is a **non-strict partial ordering** on $\mathbb{N}$ that is not **linear** (e.g., neither $2 \mid 3$ nor $3 \mid 2$).

Lemma B.1. **Strict partial orderings** are **asymmetric**.

Proof. By contradiction: if $(a, b) \in R$ and $(b, a) \in R$, then $(a, a) \in R$ by **transitivity**, and that can't be because its **irreflexive**. $\square$

Lemma B.2. If $\preceq$ is a **non-strict partial order**, then $\prec := \{(a, b) \mid a \preceq b \wedge a \neq b\}$ is a **strict partial order**. Conversely, if $\prec$ is a **strict partial order**, then $\preceq := \{(a, b) \mid a \prec b \vee a = b\}$ is a **non-strict partial order**.

Proof. (1) **Non-strict** $\rightarrow$ **Strict**: Assume $\preceq$ is reflexive, antisymmetric, and transitive. Define $a \prec b \iff (a \preceq b \wedge a \neq b)$.

- *Irreflexive:* $a \prec a$ would require $a \neq a$, impossible.
- *Transitive:* If $a \prec b$ and $b \prec c$, then $a \preceq b$, $b \preceq c$, and $a \neq b$, $b \neq c$. By transitivity of $\preceq$, $a \preceq c$. If $a = c$, antisymmetry would force $a = b$, contradiction. So $a \neq c$, hence $a \prec c$.

Thus $\prec$ is a strict partial order.

(2) **Strict** $\rightarrow$ **Non-strict**: Assume $\prec$ is irreflexive and transitive. Define $a \preceq b \iff (a = b \vee a \prec b)$.

- *Reflexive:* $a = a$ gives $a \preceq a$.
- *Antisymmetric:* If $a \preceq b$ and $b \preceq a$, then either $a = b$ or we would have $a \prec b$ and $b \prec a$, contradicting asymmetry.
- *Transitive:* If $a \preceq b$ and $b \preceq c$, then combining cases ($=$ or $\prec$) always gives $a \preceq c$.

Thus $\preceq$ is a non-strict partial order. $\square$

Examples:

- On $\mathbb{N}$, the non-strict partial order $\leq$ induces the strict order $<$ by removing equality. Conversely, $<$ induces $\leq$ by adding equality.
- On sets, the non-strict partial order $\subseteq$ induces the strict order $\subset$. Conversely, $\subset$ induces $\subseteq$ by adding equality.

Exercise: which of the following binary **relation** on integers are **partial ordering**:

1. Φ . It is not **reflexive** because we must have $(a, a) \in R$ for all integers a . But Φ has no pairs at all. $\rightarrow$ NO.
2. $\mathbb{Z} \times \mathbb{Z}$
 - It is **reflexive** since $(a, a) \in R \quad \forall a$
 - It is **symmetric** because if $(a, b) \in R$ then (b, a) also is, $\rightarrow$ it is not **antisymmetric**, $\rightarrow$ NO
3. $m \leq n$

- It is **reflexive** since $m \leq m$
- It is **asymmetric**, e.g. $2 \leq 3$ but not $3 \leq 2$
- It is **transitive**
- $\rightarrow$ YES

4. $\exists z \in \mathbb{Z} : m + z = n$

- This is actually the same as $\mathbb{Z} \times \mathbb{Z}$ because for any integers m, n , we can always pick $z = n - m$. $\rightarrow$ NO.

5. $\exists z \in \mathbb{Z} : m \cdot z = n$

- This is the same as $m \mid n$. It is **reflexive** ($m \cdot 1 = m$). It is **asymmetric** on $\mathbb{N}$ but not on $\mathbb{Z}$ ($2 \mid -2$ and $-2 \mid 2$ but $2 \neq -2$). $\rightarrow$ NO.

6. $5 \mid (m - 2)$

- This means m is **congruent** to 2 modulo 5 .
- It is **reflexive** ($5 \mid 0$)
- It is **symmetric** (not **asymmetric**)
- It is **transitive**, if $5 \mid m - n$ and $5 \mid n - k$, then $5 \mid (m - n) + (n - k) = m - k$
- $\rightarrow$ NO.

7. $m \mid n$. It is **reflexive** but not **asymmetric**. $\rightarrow$ NO.

Def B.45. Partially Ordered Set (Poset): A **Partially Ordered Set** (or **poset**) is a pair $\langle G, \preceq \rangle$ where G is a **set** and $\preceq$ is a **non-strict partial ordering** on G .

Def B.46. Least Element: Let $\langle G, \preceq \rangle$ be a **poset** and $T \subseteq G$. An **element** $t \in T$ is the **least** (or **smallest**, **minimal**, or **$\preceq$ -minimal**) element of T iff $t \preceq t'$ for all $t' \in T$.

Def B.47. Greatest Element: Let $\langle G, \preceq \rangle$ be a **poset** and $T \subseteq G$. An **element** $t \in T$ is the **greatest** (or **largest**, **maximal**, or **$\preceq$ -maximal**) element of T iff $t' \preceq t$ for all $t' \in T$.

Def B.48. Supremum and Infimum: Let $\langle G, \preceq \rangle$ be a **non-strict partial ordering** and $T \subseteq G$, then we call the **least upper bound** ($\sup(T) \in G$) of T the **supremum** of T and the **greatest lower bound** ($\inf(T) \in G$) of T the **infimum** of T (if they exist).

Example: Consider the **non-strict partial ordering** $\langle \mathbb{R}, \preceq \rangle$. Let T be the open interval $(0, 1) \subseteq \mathbb{R}$, then $\sup(T) = 1$, and $\inf(T) = 0$. Note that $0, 1 \in \mathbb{R}$ but $\notin T$. (its different from **min** and **max** of T).

B.3 Functions

Def B.49. Partial function: A [relation](#) $R \subseteq X \times Y$ is a *partial function* from X to Y ($f : X \rightarrow Y$) iff $\forall x \in X, \forall y_1, y_2 \in Y. ((x, y_1) \in R \wedge (x, y_2) \in R \Rightarrow y_1 = y_2)$ (i.e. there is at most one $y \in Y$ with $(x, y) \in f$)

Def B.50. dom, codom: We call X the domain ($\text{dom}(f)$), and Y the codomain ($\text{codom}(f)$)

Def B.51. Function Application: Instead of writing $(x, y) \in f$ we write $f(x) = y$

Def B.52. Undefined: we call a [partial function](#) $f : X \rightarrow Y$ undefined at $x \in X$, iff $\forall y \in Y. (x, y) \notin f$ ($f(x) = \perp$).

Def B.53. Function: A [partial function](#) $f : X \rightarrow Y$ is called a *function* (or *total function*) from X to Y iff it is a [total relation](#). ($\forall x \in X. \exists^1 y \in Y : (x, y) \in f$)

Note

A [partial function](#) guarantees uniqueness, but not existence. So each x has at most one y (some might have none). A [total function](#) guarantees both uniqueness and existence (at most and at least $\rightarrow$ exactly one)

Def B.54. Partial Function Space: Given two sets X and Y , the [set](#) of all possible [partial functions](#) from X to Y is called the partial function space, denoted as $X \rightarrow Y$

Note

If X and Y are finite, then for each element $x \in X$ we have $\Gamma(Y) + 1$ choices (either pick one of the $\Gamma(Y)$ values in Y , or leave it [undefined](#)). Therefore, $\Gamma(X \rightarrow Y) = (\Gamma(Y) + 1)^{\Gamma(X)}$.

Example, $X = Y = \{0, 1\}$

$$X \rightarrow Y = \{\Phi, \{(0, 0)\}, \{(0, 1)\}, \{(1, 0)\}, \{(1, 1)\}, \{(0, 0), (1, 0)\}, \{(0, 1), (1, 1)\}, \{(0, 1), (1, 0)\}, \{(0, 0), (1, 1)\}\}$$

Notice that $\Gamma(X \rightarrow Y) = 9 = (\Gamma(Y) + 1)^{\Gamma(X)} = (2 + 1)^2$

Def B.55. Function Space: Given two sets X and Y , the [set](#) of all possible [functions](#) from X to Y is called the partial function space, denoted as $X \rightarrow Y$

Note

If X and Y are finite, then for each element $x \in X$ we must pick exactly one of the $\Gamma(Y)$ values in Y . Therefore, $\Gamma(X \rightarrow Y) = (\Gamma(Y))^{\Gamma(X)}$.

Example, $X = Y = \{0, 1\}$

$$X \rightarrow Y = \{\{(0, 0), (1, 0)\}, \{(0, 1), (1, 1)\}, \{(0, 1), (1, 0)\}, \{(0, 0), (1, 1)\}\}$$

Notice that $\Gamma(X \rightarrow Y) = 4 = (\Gamma(Y))^{\Gamma(X)} = (2)^2$

A [function](#) $f : X \rightarrow Y$ is called

- **injective** iff $\forall x_1, x_2 \in X. f(x_1) = f(x_2) \Rightarrow x_1 = x_2$
- **surjective** iff $\forall y \in Y. \exists x \in X. f(x) = y$
- **bijective** iff f is [injective](#) and [surjective](#)

Let's look for some examples to make the previous definitions clear.

Assume $X = \{1, 2, 3\}$ $Y = \{a, b\}$ The [relation](#) $R = \{(1, a), (1, b)\}$ is just a [relation](#), it is not a [partial function](#) nor a [total function](#).

If $R = \{(1, a)\}$ it becomes a [partial function](#) but not a [total one](#).

If $R = \{(1, a), (2, a), (3, a)\}$ it becomes a [total function](#) because each $x \in X$ has exactly one $y \in Y$. However, it is not [injective](#) (1, 2, 3 all map to the same value) and not [surjective](#) (b is not reached).

IF $R = \{(1, a), (2, a), (3, b)\}$ then this is a [surjective function](#) but not an [injective one](#) (since 1, 2 map to the same value).

Assume $X = \{1, 2\}$ $Y = \{a, b, c\}$ $f = \{(1, a), (2, b)\}$ then this is an [injective function](#) (different inputs map to different output) but not [surjective](#) since there is no x such that $f(x) = c$.

And finally, assume $X = \{1, 2\}$, $Y = \{a, b\}$, $f = \{(1, a), (2, b)\}$ then this is a **total function** that is both **injective** and **surjective**, i.e. a **bijective function**.

Def B.56. Function Composition: If $f : A \rightarrow B$ and $g : B \rightarrow C$ are **functions**, then we call $g \circ f : A \rightarrow C; x \mapsto g(f(x))$ the composition of g and f (read g after f).

Note

Given $f, g : \mathbb{N} \rightarrow \mathbb{N}$, if f and g both are **injective/bijective**, so is $g \circ f$

Observations:

- If f is **injective**, then the **converse relation** f^{-1} is a **partial function**.

For example: $X = \{1, 2, 3\}, Y = \{a, b, c, d\}, f : X \rightarrow Y = \{(1, a), (2, b), (3, c)\}$. f is clearly **injective** since different inputs map to different outputs. $\text{dom}(f^{-1}) = \text{codom}(f)$ and vice versa. $f^{-1} = \{(a, 1), (b, 2), (c, 3)\}$. It is a **partial function** because for $d \in \text{dom}(f^{-1})$ there is no $f^{-1}(d)$

- If f is **surjective**, then the **converse relation** f^{-1} is a **total relation**.

For example: $X = \{1, 2, 3\}, Y = \{a, b\}, f : X \rightarrow Y = \{(1, a), (2, b), (3, b)\}$. f is clearly **surjective** because both a, b are hit by some x . $f^{-1} = \{(a, 1), (b, 2), (b, 3)\}$. It is a **total function**.

- If f is **bijective**, call the **converse relation** **inverse function**, we also write it as f^{-1} .

- If f is **bijective**, then the **inverse function** f^{-1} is a **total function**.

- If $f : X \rightarrow Y$ is **bijective**, then $f \circ f^{-1} = \text{Id}_Y$.

For example: $X = \{1, 2, 3\}, Y = \{a, b, c\}, f = \{(1, a), (2, b), (3, c)\}, f^{-1} = \{(a, 1), (b, 2), (c, 3)\}$. $(f \circ f^{-1})(y) = f(f^{-1}(y))$

$$- y = a : f^{-1}(a) = 1, f(1) = a$$

$$- y = b : f^{-1}(b) = 2, f(2) = b$$

$$- y = c : f^{-1}(c) = 3, f(3) = c$$

$$\rightarrow f \circ f^{-1} = \text{Id}_Y$$

Def B.57. Image and Preimage: Let $f : X \rightarrow Y$ be a **function**, $X' \subseteq X$ and $Y' \subseteq Y$, then we call

- $f(X') := \{y \in Y \mid \exists x \in X'. (x, y) \in f\}$ the **image** of X' under f ,
- $\text{Im}(f) := f(X)$ the **image** of f , and
- $f^{-1}(Y') := \{x \in X \mid \exists y \in Y'. (x, y) \in f\}$ the **preimage** of Y' under f .

Note

Let $f : \mathbb{N} \rightarrow \mathbb{N}$ be **injective** and let its **partial inverse** be

$$f^{-1} : \mathbb{N} \rightarrow \mathbb{N}, \quad \text{dom}(f^{-1}) = \text{im}(f), \quad f^{-1}(y) = x \text{ iff } f(x) = y.$$

Then $f^{-1} \circ f = \text{Id}_{\mathbb{N}}$ as a **total function** $\mathbb{N} \rightarrow \mathbb{N}$.

Proof. For any $x \in \mathbb{N}$ we have $f(x) \in \text{im}(f) = \text{dom}(f^{-1})$, hence $(f^{-1} \circ f)(x) = f^{-1}(f(x)) = x$. $\square$

For example: Define $f : \mathbb{N} \rightarrow \mathbb{N}$, $f(x) = 2x$. Then f is **injective** with **image** $\text{im}(f) = \{0, 2, 4, 6, \dots\}$. Its **partial inverse** is

$$f^{-1} : \mathbb{N} \rightarrow \mathbb{N}, \quad f^{-1}(y) = \frac{y}{2} \text{ if } y \text{ is even, undefined if } y \text{ is odd.}$$

For all $x \in \mathbb{N}$ we have

$$(f^{-1} \circ f)(x) = f^{-1}(f(x)) = f^{-1}(2x) = \frac{2x}{2} = x,$$

hence $f^{-1} \circ f = \text{Id}_{\mathbb{N}}$. By contrast, $f \circ f^{-1}$ is only the identity on the even numbers.

Another example: Define $f : \mathbb{N} \rightarrow \mathbb{N}$, $f(x) = x + 1$. Then f is **injective** with **image** $\text{im}(f) = \{1, 2, 3, \dots\}$. Its **partial inverse** is

$$f^{-1} : \mathbb{N} \rightarrow \mathbb{N}, \quad f^{-1}(y) = y - 1 \text{ if } y \geq 1, \quad \text{undefined if } y = 0.$$

For all $x \in \mathbb{N}$ we have

$$(f^{-1} \circ f)(x) = f^{-1}(f(x)) = f^{-1}(x+1) = (x+1) - 1 = x,$$

hence again $f^{-1} \circ f = \text{Id}_{\mathbb{N}}$. Here $f \circ f^{-1}$ is only the identity on $\{1, 2, 3, \dots\}$.

Def B.58. Isomorphic: We call two **sets** A and B isomorphic iff there is a **bijection** $f : A \rightarrow B$.

Def B.59. Cardinality: We say that a **set** A is **finite** and has **cardinality** $\Gamma(A) \in \mathbb{N}$, iff there is a **bijection** $f : A \rightarrow \{n \in \mathbb{N} \mid n < \Gamma(A)\}$.

Def B.60. Countably Infinite: We say that a **set** A is countably infinite, iff there is a **bijection** $f : A \rightarrow \mathbb{N}$. The **cardinality** of such set is denoted by $\aleph_0$. Examples: $\mathbb{N}, \mathbb{Z}, \dots$

Def B.61. Uncountably Infinite: We say that a **set** A is **uncountably infinite**, iff A is infinite and there is **no** **bijection** $f : A \rightarrow \mathbb{N}$. The **cardinality** of the such sets is denoted by $\mathfrak{c}$ (for continuum). Examples: $\mathbb{R}, [0, 1], \mathbb{C}, \mathcal{P}(\mathbb{N})$

Def B.62. Countable: A **set** A is called countable, iff it is **finite** or **countably infinite**

Note

The **set** of real numbers $\mathbb{R}$ is not **countable** (there is no **bijection**) $\mathbb{R} \rightarrow \mathbb{N}$

Def B.63. Restriction: Let $f : A \rightarrow B$ and $C \subseteq A$, then we call the **function** $f \upharpoonright_C := \{(c, b) \in f \mid c \in C\}$ the restriction of f to C .

Def B.64. Lambda-Notation: Let X be a **set** and E an expression depending on a variable $x \in X$. The expression $\lambda x \in X. E$ denotes the (**partial**) **function** f from X to Y with $f(x) = E$ for all $x \in X$.

Examples:

- the identity function is written in **lambda notation**: $\lambda n \in \mathbb{N}. n = \{(n, n) \mid n \in \mathbb{N}\} = \text{Id}_{\mathbb{N}} = \{(0, 0), (1, 1), (2, 2), \dots\}$
- the square function: $\lambda x \in \mathbb{N}. x^2 = \{(x, x^2) \mid x \in \mathbb{N}\} = \{(0, 0), (1, 1), (2, 4), (3, 9), \dots\}$
- the "add" operation: $\lambda(x, y) \in \mathbb{N} \times \mathbb{N}. x + y = \{((0, 0), 0), ((1, 0), 1), \dots\}$

Def B.65. Curried Function Type: Let X, Y, Z be sets. An *uncurried* **function** is written as

$$f : X \times Y \rightarrow Z, \quad f(x, y) = E(x, y).$$

The same **function** can equivalently be written in *curried* form as

$$f : X \rightarrow Y \rightarrow Z, \quad f(x)(y) = E(x, y).$$

Thus in the curried notation the **function** takes one argument at a time: for $x \in X$, the value $f(x)$ is itself a **function** $Y \rightarrow Z$, and applying it to $y \in Y$ yields $f(x)(y) \in Z$.

Example. The addition **function** can be written in **uncurried** or **curried** form:

$$\text{add} : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}, \quad \text{add}(m, n) = m + n,$$

equivalently: $\text{add} : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad \text{add}(m)(n) = m + n.$

In the **curried** form, for fixed $m \in \mathbb{N}$, the expression $\text{add}(m) : \mathbb{N} \rightarrow \mathbb{N}$ is the **function** $n \mapsto m + n$, i.e. the **function** that adds m to its argument.

Another Example (Ternary)

$$h : \mathbb{N} \times \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}, \quad h(m, n, o) = m + n + o.$$

Equivalently, in **curried form** we can write

$$h : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad h(m)(n)(o) = m + n + o.$$

In this form, $h(m)$ is a **function**

$$h(m) : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad h(m)(n)(o) = m + n + o,$$

which for fixed m returns the **function** "add m to the sum of two further arguments."

Def B.66. Well-Defined Function: Let X, Y be sets. A **function** declaration $f : X \rightarrow Y$ together with a defining expression $E(x)$ is called *well-defined* iff for every $x \in X$:

- the expression $E(x)$ is meaningful (all variables are bound and the operations are valid), and
- the result $E(x)$ lies in Y .

Equivalently: a **function** is well-defined if its type annotation $f : X \rightarrow Y$ is consistent with the expression that defines it.

Checking well-definedness in practice:

The general definition above says that a **function** declaration must be consistent with its defining expression. Typical ways in which this can fail are:

- **Type mismatch:** the declared type $X \rightarrow Y$ does not agree with the actual input/output structure of the definition.
- **Free variables:** the definition uses symbols that are not bound by the input.
- **Partiality:** the definition does not provide an output for all inputs in the domain.

To illustrate, let us fix two **functions**

$$f : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}, \quad g : \mathbb{N} \rightarrow \mathbb{N},$$

and examine a series of possible definitions of **functions** h . Each case shows whether the definition is **well-defined** and why.

A correct curried definition

$$h : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad h(m)(n) = f(n, g(m)).$$

This is well-defined: for each $m \in \mathbb{N}$ we obtain $g(m) \in \mathbb{N}$, so $(n, g(m)) \in \mathbb{N} \times \mathbb{N}$ and $f(n, g(m)) \in \mathbb{N}$.

Incorrect composition

$$g \circ f : \mathbb{N} \rightarrow \mathbb{N}.$$

This is not well-defined with the given type. Since $f : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}$, the composition $g \circ f$ has domain $\mathbb{N} \times \mathbb{N}$, hence $g \circ f : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}$, not $\mathbb{N} \rightarrow \mathbb{N}$.

A correct nested definition

$$h : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad h(m)(n, o) = f(n, f(m, o)).$$

This is **well-defined**: $f(m, o) \in \mathbb{N}$, so $(n, f(m, o)) \in \mathbb{N} \times \mathbb{N}$, and thus $f(n, f(m, o)) \in \mathbb{N}$.

Type mismatch between **curried** and **uncurried**

$$h : \mathbb{N} \rightarrow \mathbb{N} \rightarrow \mathbb{N}, \quad h = f.$$

This is not **well-defined**, since $f : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}$ has a different type from the declared type of h . Mathematically $f(x, y)$ and $f(x)(y)$ can be identified via **currying**, but in type-theoretic terms they are distinct unless an explicit **curry** operation is applied.

Free variable problem

$$h : \mathbb{N} \rightarrow \mathbb{N}, \quad h(x) = f(x, y).$$

This is not **well-defined**, since y is not bound. Without fixing y in advance, the right-hand side is not a **function** of x alone.

Partial definition problem

$$h : \mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}, \quad h(x, 1) = x.$$

This is not a **total function**, because the definition only specifies values when the second argument is 1. A **function** of type $\mathbb{N} \times \mathbb{N} \rightarrow \mathbb{N}$ must assign a value for every **pair** (x, y) .

Def B.67. Invertible: Let A and B be **sets** and $f : A \rightarrow B$ a **function**, then f is called invertible, iff there is a **function** $g : B \rightarrow A$, such that $f \circ g = \text{Id}_B$. g is called the **inverse function** (or just **inverse**) of f and is written as f^{-1} .

Theorem B.3. A **function** $f : A \rightarrow B$ is **invertible**, iff it is **bijective**.

Proof.

1. If f is **bijective**, then it is **invertible**:

- Assume $f : A \rightarrow B$ is **bijective**
- for each $b \in B$ there exists a unique $a \in A$ such that $f(a) = b$. Define $g : B \rightarrow A$ by sending b to this unique a . Therefore:
 1. $g \circ f = \text{Id}_A$, because if you start with $a \in A$, apply f , then g , you get back a .
 2. $f \circ g = \text{Id}_B$, because if you start with $b \in B$, apply g , then f , you get back b .

- Hence, g is the **inverse** of f . So f is **invertible**

2. If f is **invertible**, then it is **bijjective**:

Suppose there exists a **function** $g : B \rightarrow A$ such that $g \circ f = \text{Id}_A$ and $f \circ g = \text{Id}_B$. We want to show f is **bijjective**.

1. **injectivity**

(a) Assume $f(a) = f(a')$.

(b) Apply g to both sides:

$$a = \text{Id}_A(a) = g(f(a)) = g(f(a')) = a'$$

(c) So, $a = a'$. Hence f is **injective**.

2. **surjectivity**

(a) Take any $b \in B$. Then

$$f(g(b)) = (f \circ g)(b) = \text{Id}_B(b) = b$$

(b) So b is indeed an image of f . Hence f is **surjective**.

□

Lemma B.4. Let A and B be **sets**, then the **set** $\mathcal{R} = \mathcal{P}(A \times B)$ of all **relations** between A and B and the **function space** $A \rightarrow \mathcal{P}(B)$ are **isomorphic**.

Proof. For each $R \in \mathcal{R}$ and each $x \in A$ we define: $S_x^R := \{y \in B \mid (x, y) \in R\}$ (the **set** of all y in B that are related to x by R). Example:

$$A = \{1, 2\} \quad B = \{a, b, c\} \quad R = \{(1, a), (1, c), (2, b)\}$$

Then for $x = 1 : S_1^R = \{a, c\}$ and for $x = 2 : S_2^R = \{b\}$.

We define a the following mapping:

$$\varphi : \mathcal{R} \rightarrow (A \rightarrow \mathcal{P}(B)), \quad \varphi(R) = f_R \quad \text{where } f_R(x) = S_x^R.$$

For our example, this would be $\varphi(R) = f_R = \{1 \mapsto \{a, c\}, 2 \mapsto \{b\}\}$

Then, for any $f \in A \rightarrow \mathcal{P}(B)$, we compute the **inverse** as $\varphi^{-1}(f) = \{(x, y) \mid y \in f(x)\}$. This is exactly the **relation** corresponding to f .

Since both directions match up perfectly, φ is a **bijection**, hence an **isomorphism**. □

Example: the $>$ **relation** on $\mathbb{N}$ can also be represented as the **function** $f_{>}$ that maps a number n to the **set** $\{n+1, \dots\}$, i.e. $f_{>} : \mathbb{N} \rightarrow \mathcal{P}(\mathbb{N}) ; n \mapsto \{m \mid m > n\}$

Another example: let $A = \{x, y\}, B = \{1\}$, the **cartesian product** $A \times B = \{(x, 1), (y, 1)\}$. Hence, the **power set** (all possible relations) is:

$$\mathcal{P}(A \times B) = \{\{\phi\}, \{(x, 1)\}, \{(y, 1)\}, \{(x, 1), (y, 1)\}\}$$

Note: the **cardinality** of the **power set** of the **cartesian product** of two **sets** A, B is $\Gamma(\mathcal{P}(A \times B)) = 2^{\Gamma(A)\Gamma(B)}$.

$\mathcal{P}(B) = \{\{\phi\}, \{1\}\}$, **function space** $A \rightarrow \mathcal{P}(B)$ is the **set** of all **functions** $f : A \rightarrow \mathcal{P}(B)$:

$$A \rightarrow \mathcal{P}(B) = \{f_1, f_2, f_3, f_4\}$$

where:

- $f_1 = \{(x, \{\phi\}), (y, \{\phi\})\}$
- $f_2 = \{(x, \{\phi\}), (y, \{1\})\}$
- $f_3 = \{(x, \{1\}), (y, \{\phi\})\}$
- $f_4 = \{(x, \{1\}), (y, \{1\})\}$

Since $\Gamma(A \rightarrow \mathcal{P}(B)) = \Gamma(\mathcal{P}(A \times B))$, we can define a **bijection**:

- $\varphi(\{\phi\}) \mapsto f_1$
- $\varphi(\{(x, 1)\}) \mapsto f_3$
- $\varphi(\{(y, 1)\}) \mapsto f_2$

- $\varphi(\{(x, 1), (y, 1)\}) \mapsto f_4$

$\rightsquigarrow$ isomorphism.

Note

The bijection $\phi(R) = f$ entails the following properties of R :

- R is a **total relation**, iff $\phi \notin \text{Im}(f)$
- R is a **partial function**, iff $\Gamma(S) \leq 1 \quad \forall S \in \text{Im}(f)$
- $R = \phi$, iff $\text{Im}(f) = \{\phi\}$
- R is **reflexive**, iff $x \in f(x) \quad \forall x \in A$

We have defined **equivalence relations** as **reflexive**, **symmetric**, and **transitive** relations, e.g., **equality** is an **equivalence relation**.

Let S be a **set** of persons, and $=_A \subseteq S^2$, such that $s =_A t$, iff s and t have the same age, then $=_A$ is an **equivalence relation** on S .

Let S and T be **sets** and $S \sim_\Gamma T$, iff there is a **bijection** $f : S \rightarrow T$, then $\sim_\Gamma$ is an **equivalence relation**. For example, let $S = \{1, 2, 3\}$ $T = \{a, b, c\}$. We can define a **bijection** $f = \{1 \mapsto a, 2 \mapsto b, 3 \mapsto c\}$, so, $S \sim_\Gamma T$ (they have the same **size**).

Lemma B.5. If S is a **set** and $R \subseteq S^2$ is an **equivalence relation**, then $= \subseteq R$.

Proof Sketch: because of **reflexivity** ($\forall x \in S, (x, x) \in R$), the **equality relation** $=$ is exactly $\{(x, x) \mid x \in S\}$. Therefore, $= \subseteq R$.

B.4 Equivalence Classes

Def B.68. Equivalence Classes and Quotients: Let S be a [set](#) and R be an [equivalence relation on \$S\$](#) , then for any $x \in S$ we call the set $[x]_R := \{y \in S \mid R(x, y)\}$ the **equivalence class** of x under R , and the [set](#) $S/R := \{[x]_R \mid x \in S\}$ the **quotient space** of S under R , it is often read as S modulo R . The [element](#) x is called the **representative** of $[x]_R \in S/R$.

Example: Let $S = \mathbb{Z}$ and R is the [equivalence relation on \$S\$](#) that describes the same parity ([congruence modulo 2](#)) $x \equiv y \pmod{2} \Leftrightarrow 2 \mid (x - y)$:

$$R = \{(x, y) \in \mathbb{Z} \times \mathbb{Z} \mid x \equiv y \pmod{2}\}$$

- $(x, y) \in R$ if the difference $x - y$ is divisible by 2.
- i.e. either both are even, or both are odd.

All odd-odd pairs:

$$\{(x, y) \mid x, y \in \{\dots, -3, -1, 1, 3, 5, \dots\}\}$$

All even-even pairs:

$$\{(x, y) \mid x, y \in \{\dots, -4, -2, 0, 2, 4, 6, \dots\}\}$$

Example pairs in R : $(1, 3), (5, -1), (7, 7), (2, 4), (0, -6), (-8, 10)$ Then, we have the following two distinct classes:

- $[1]_R = \{\dots, -3, -1, 1, 3, 5, \dots\}$
- $[2]_R = \{\dots, -4, -2, 0, 2, 4, 6, \dots\}$

For each $x \in \mathbb{Z}$ we have:

- the [set](#) $[1]_R$ as the [equivalence class](#) if x is odd
- the [set](#) $[2]_R$ as the [equivalence class](#) if x is even
- 1 and 2 are the [representatives](#)

And the [quotient space set](#) of $\mathbb{Z}$ under R is:

$$\mathbb{Z}/R = \{[1]_R, [2]_R\} = \{\text{odds, evens}\}$$

Def B.69. Canonical Projection: The mapping $\pi_R : S \rightarrow S/R$; $x \mapsto [x]_R$ is called the **canonical projection** or **canonical surjection** of S to S/R .

Given our example, we define the following [canonical projection](#):

$$\pi_R : \mathbb{Z} \rightarrow \{[1]_R, [2]_R\}, \quad x \mapsto [1]_R \text{ if odd, and } x \mapsto [2]_R \text{ if even.}$$

Def B.70. System of Representatives: Let R be an [equivalence relation on \$S\$](#) , a [subset](#) $M \subseteq S$ is called a system of representatives, iff M contains exactly one [representative](#) for each [equivalence class](#) of R .

Given our example, a [system of representatives](#) is any [subset](#) of $\mathbb{Z}$ that picks exactly one [element](#) from each [equivalence class](#), e.g.

- $M = \{1, 2\}$ is a [system of representatives](#)
- $M = \{-3, 0\}$ is a [system of representatives](#)
- $M = \{5, -8\}$ is a [system of representatives](#)

Observation: Often a [quotient space](#) S/R behaves similarly to the original [set](#) S .

Example: $n \equiv_k m \Leftrightarrow k \mid (n - m)$ is an [equivalence relation](#) and $\pi_{\equiv_k} : \mathbb{Z} / \equiv_k \rightarrow \{n \mid 0 \leq n \leq (k - 1)\}$ is [bijective](#).

Another example, the general [congruence modulo](#) $n \equiv_k m \Leftrightarrow k \mid (n - m)$ partitions $\mathbb{Z}$ into exactly k [equivalence classes](#) (all numbers that leave the same remainder when divided by k). e.g. $k = 3$:

- $[0] = \{\dots, -6, -3, 0, 3, 6, \dots\}$
- $[1] = \{\dots, -5, -2, 1, 4, 7, \dots\}$
- $[2] = \{\dots, -4, -1, 2, 5, 8, \dots\}$

We have the **canonical projection** $\pi_{\equiv_k} : \mathbb{Z} \rightarrow \mathbb{Z}/\equiv_k, \quad n \mapsto [n]$ that sends each $n \in \mathbb{Z}$ to its **equivalence class**.

Notice that the **quotient space** $\mathbb{Z}/\equiv_k$ has exactly k classes, so, we have the following **bijection**:

$$\varphi : \mathbb{Z}/\equiv_k \rightarrow \{0, 1, \dots, k-1\}, \quad [n] \mapsto \text{the remainder of } n \text{ mod } k.$$

If we compose the **canonical projection** with the **bijection** we get the remainder, e.g. for $k = 3$:

$$-6 \xrightarrow{\pi_{\equiv_3}} [0] \xrightarrow{\varphi} 0.$$

$$7 \xrightarrow{\pi_{\equiv_3}} [1] \xrightarrow{\varphi} 1.$$

Lemma B.6. Let $= \subseteq S \times S$ be the **equality relation** on a set S , then $[x]_ = \in S/ =$ is always a **singleton** and thus $S \sim_{\Gamma} S/ =$.

Proof. If we take the **equality relation on a set** S , then for any $x \in S$, the **equivalence class** is:

$$[x]_ = = \{y \in S \mid y = x\}$$

But the only y that satisfies $y = x$ is $y = x$ itself. So we get $[x]_ = = \{x\}$. This is a **singleton set**. So every **equivalence class** under $=$ has exactly one element.

The **quotient space** is $S/ = := \{[x]_ = \mid x \in S\}$ and each $[x]_ =$ is the singleton set $\{x\}$, so the **quotient space** becomes $S/ = := \{\{x\} \mid x \in S\}$. In other words, the **quotient space** under the **equality relation** just replaces each **element** x with one-element **set** $\{x\}$.

That means the **set** S is in **bijection** with its **quotient space** under $=$ via the following **bijective** map:

$$\phi : S \rightarrow S/ =, \quad x \mapsto [x]_ = = \{x\}$$

Therefore, we express the **bijection** via same **cardinality**:

$$S \sim_{\Gamma} S/ =$$

□

Fun Example: Let S be the **set** of all cities worldwide and $\sim_c \in S^2$ defined by $a \sim_c b$, iff a and b are in the same country. Then $\sim_c$ is an **equivalence relation on** S , and the **quotient space** $S/ \sim_c$ of S under $\sim_c$ consists of the **equivalence class** $[x]_{\sim_c}$ where $x \in S$. The capital of a country is a good **representative** for an **equivalence class**.

For instance: let's take a **subset** of all worldwide cities for simplicity, consider **domain of discourse**:

$$S = \{\text{Berlin, Munich, Hamburg, Paris, Lyon, Marseille, Rome, Milan}\}.$$

For $a, b \in S$ define the **relation**:

$$a \sim_c b \iff a \text{ and } b \text{ are in the same country.}$$

This **relation** is:

- **reflexive**: any city is in the same country as itself.
- **symmetric**: if Berlin is in the same country as Munich, then Munich is in the same country as Berlin.
- **transitive**: if Berlin and Munich are in the same country, and Munich and Hamburg are in the same country, then Berlin and Hamburg are in the same country.

Hence, $\sim_c$ is an **equivalence relation**.

The **quotient space** $S/ \sim_c$ is the **set** of **equivalence classes**, i.e. the partition of S into countries:

$$[\text{Berlin}]_{\sim_c} = \{\text{Berlin, Munich, Hamburg}\}, \quad [\text{Paris}]_{\sim_c} = \{\text{Paris, Lyon, Marseille}\}, \quad [\text{Rome}]_{\sim_c} = \{\text{Rome, Milan}\}.$$

So

$$S/ \sim_c = \left\{ \{\text{Berlin, Munich, Hamburg}\}, \{\text{Paris, Lyon, Marseille}\}, \{\text{Rome, Milan}\} \right\}.$$

The **set** of chosen **representatives** is: $\{\text{Berlin, Paris, Rome}\}$.

B.5 Algebraic Structures

Def B.71. Mathematical (Algebraic) Structure: A mathematical structure combines multiple [mathematical objects](#) (the components) into a new [object](#). The components usually have names by which they can be referenced. Given a [definition](#) of a mathematical structure S , we say that any [object](#) that conforms to that is an instance of S .

Def B.72. Structure Signature: A structure signature combines the components, their types, and accessor names of a [mathematical structure](#) in a tabular overview.

Def B.73. Group: A group is a [mathematical structure](#) $\langle G, \circ, e, .^{-1} \rangle$ that consists of:

- a base [set](#) G of [objects](#),
- an operation $\circ : G \times G \rightarrow G$, such that $\forall a, b \in G. a \circ b \in G$ and $\forall a, b, c \in G. (a \circ b) \circ c = a \circ (b \circ c)$,
- a unit $e \in G$, such that $\forall a \in G. e \circ a = a \circ e = a$, and
- the [inverse function](#) $.^{-1} : G \rightarrow G$, such that $\forall a \in G. a \circ a^{-1} = a^{-1} \circ a = e$

An example of a [group](#) is the [set](#) $G = \mathbb{Z}$, with the operation $+: \mathbb{Z} \times \mathbb{Z} \rightarrow \mathbb{Z}$; then $e = 0$ and $.^{-1}$ is $\lambda x \in \mathbb{Z}. -x$. We write it as $\langle \mathbb{Z}, +, 0, - \rangle$.

Question: can we do the same for multiplication where $e = 1$? If we think about it, we would be stuck on finding the $.^{-1}$ component of the group. What can we always multiply with any integer and get $e = 1$ back? Take for example 1, we can multiply it with $\frac{1}{1}$ (great!). Take -1 , we can multiply it with $\frac{1}{-1}$ (also great!). But take any other integer, e.g. 2, we must multiply it with $\frac{1}{2}$ to get 1 back, however, $\frac{1}{2} \notin \mathbb{Z}$, it is however $\in \mathbb{Q}$! Moreover, 0 never has a multiplicative inverse in any field! The multiplicative group would only make sense iff $G = \mathbb{Q} \setminus \{0\}$. We write that as $\langle \mathbb{Q} \setminus \{0\}, *, 1, a \mapsto \frac{1}{a} \rangle$.

But we need some [multiplication structure](#) that works for $\mathbb{Z}$!

Def B.74. Monoid: A monoid is a [mathematical structure](#) $\langle M, \circ, e \rangle$ that consists of:

- A base set M of [objects](#),
- An operation $\circ : M \times M \rightarrow M$, such that for all $a, b \in M$, we have $a \circ b \in M$ (closure),
- Associativity: for all $a, b, c \in M$, $(a \circ b) \circ c = a \circ (b \circ c)$, and
- A unit element $e \in M$, such that for all $a \in M. e \circ a = a \circ e = a$.

Hence, the [mathematical structure](#) $\langle \mathbb{Z}, *, 1 \rangle$ is a [monoid](#) (a [group](#) missing an [inverse function](#)).

We saw that $\langle \mathbb{Z}, +, 0, - \rangle$ is a [group](#) under addition. And $\langle \mathbb{Z}, *, 1 \rangle$ is a [monoid](#) under multiplication. Now, we want to do everyday arithmetics with integers; we want expressions where we have additions and [multiplication](#) (e.g. $a * (b + c)$) which evaluates to $(a * b) + (a * c)$.

For that, we need a [mathematical structure](#) where both addition and [multiplication](#) live together.

Def B.75. Ring: A ring is a [mathematical structure](#) $\langle R, +, 0, -, *, 1 \rangle$ consisting of:

- A base [set](#) R .
- An addition operation $+: R \times R \rightarrow R$ such that $\langle R, +, 0, - \rangle$ is a [group](#) (called abelian [group](#)).
- A multiplication operation $*: R \times R \rightarrow R$ such that $\langle R, *, 1 \rangle$ is a [monoid](#).
- Distributivity: for all $a, b, c \in R, a * (b + c) = (a * b) + (a * c), (a + b) * c = (a * c) + (b * c)$

Hence, mixing components from the [group](#) $\langle \mathbb{Z}, +, 0, - \rangle$ and the [monoid](#) $\langle \mathbb{Z}, *, 1 \rangle$ leaves us with a [ring](#) $\langle \mathbb{Z}, +, 0, -, *, 1 \rangle$

Def B.76. Magma: A magma is a [mathematical structure](#) $\langle M, \circ \rangle$ consisting of:

- A base [set](#) M , and
- A binary operation $\circ : M \times M \rightarrow M$.

Example: $\langle \mathbb{Z}, - \rangle$ (integers under subtraction). $\langle \mathbb{N}, + \rangle$ (natural numbers under addition)

Def B.77. Semigroup: A semigroup ([magma](#) with associativity) is a [mathematical structure](#) $\langle S, \circ \rangle$ consisting of:


- A base [set](#) S ,
- A binary operation $\circ : S \times S \rightarrow S$, and
- Associativity: for all $a, b, c \in S, (a \circ b) \circ c = a \circ (b \circ c)$

Example: $\langle \mathbb{Z}, + \rangle$ (integers under addition).

Note

Given collections M and S of instances of **magmas** and **semigroups**, $S \subseteq M$ as there are **magmas** whose operation is not associative.

Mathematical Structures in a Graph:



- nodes are **structures** with **objects** ($\circ: \tau$) and properties ($n \vdash P$).
- edges represent inheritance
 - regular inheritance: $\hookleftarrow$
 - copying substructures: $\xrightarrow{n}$
 - interpretation via φ : $\xrightarrow{\sim}$
- Examples are just interpretations ($\rightsquigarrow$) read backwards
- Any **object** / **property** / **theorem** can be copied along any edge


C Topics in Computer Science

C.1 Graph Theory

Def C.1. Undirected Graph: An undirected graph is a pair $\langle V, E \rangle$ such that

- V is a **set** of **vertices (nodes)**, and
- $E \subseteq \{ \{v, \bar{v}\} \mid v, \bar{v} \in V \wedge v \neq \bar{v} \}$ is the **set** of its **undirected edges**.


Example: An **undirected graph** $G_1 = \langle V_1, E_1 \rangle$, where $V_1 = \{A, B, C, D, E\}$ and $E_1 = \{\{A, B\}, \{A, C\}, \{A, D\}, \{B, D\}, \{B, E\}\}$



Def C.2. Directed Graph: A directed graph (digraph) is a pair $\langle V, E \rangle$ such that

- V is a **set** of **vertices (nodes)**, and
- $E \subseteq V \times V$ is the **set** of its **directed edges**.

Example: A **directed graph** $G_2 = \langle V_2, E_2 \rangle$, where $V_2 = \{1, 2, 3, 4, 5\}$ and $E_2 = \{(1, 1), (1, 2), (2, 3), (3, 2), (2, 4), (5, 4)\}$



Def C.3. Indegree: Given a **directed graph** $\langle V, E \rangle$, the indegree $\text{indeg}(v)$ of a vertex $v \in V$ is defined as $\text{indeg}(v) = \Gamma(\{w \mid (w, v) \in E\})$.

Example: $\text{indeg}(2) = 2$

Def C.4. Outdegree: Given a **directed graph** $\langle V, E \rangle$, the outdegree $\text{outdeg}(v)$ (branching factor) of a vertex $v \in V$ is defined as $\text{outdeg}(v) = \Gamma(\{w \mid (v, w) \in E\})$.

Example: $\text{outdeg}(5) = 1$

Note

For an **undirected graph**, $\text{indeg}(v) = \text{outdeg}(v)$ for all $v \in V$.

Observation: **directed graphs** are nothing else than **relations**.

Def C.5. Initial vs Terminal Node: Let $G = \langle V, E \rangle$ be a **directed graph**, then we call a node $v \in V$



- **initial** (source of G), iff there is no $w \in V$ such that $(w, v) \in E$ (no predecessor)
- **terminal** (sink of G), iff there is no $w \in V$ such that $(v, w) \in E$ (no successor)

Def C.6. Graph Isomorphism: Iff we can find a **bijection** $\psi : V \rightarrow \bar{V}$ between two graphs $G = \langle V, E \rangle$ and $\bar{G} = \langle \bar{V}, \bar{E} \rangle$ then we call them isomorphic. The **bijection** $\psi : V \rightarrow \bar{V}$ is defined as $(a, b) \in E \Leftrightarrow (\psi(a), \psi(b)) \in \bar{E}$ for **directed graph**. And for **undirected graph** it is defined as $\{a, b\} \in E \Leftrightarrow \{\psi(a), \psi(b)\} \in \bar{E}$

Def C.7. Equivalent Graphs: Two graphs G and $\bar{G}$ are **equivalent** iff there is an **isomorphism** ψ between G and $\bar{G}$.

Example: the following two **graphs** are **equivalent** as there exists an **isomorphism** between them given by the **bijection**:


$$\psi := \{a \mapsto 5, b \mapsto 6, c \mapsto 2, d \mapsto 4, e \mapsto 1, f \mapsto 3\}$$



Def C.8. Labeled Graph: A labeled graph G is a quadruple $\langle V, E, L, l \rangle$ where $\langle V, E \rangle$ is a graph and $l : V \cup E \rightarrow L$ is a **partial function** into a **set** L of labels.

Example: $G = \langle V, E, L, l \rangle$ with $V = \{A, B, C, D, E\}$, where

- $E = \{(A, A), (A, B), (B, C), (C, B), (B, D), (E, D)\}$
- $l : V \cup E \rightarrow \{+, -, \phi\} \times \{1, \dots, 9\}$ with
 - $l(A) = 5 \quad l(B) = 3 \quad l(C) = 7 \quad l(D) = 4 \quad l(E) = 8,$
 - $l((A, A)) = -0 \quad l((A, B)) = -2 \quad l((B, C)) = +4,$
 - $l((C, B)) = -4 \quad l((B, D)) = +1 \quad l((E, D)) = -4$



Def C.9. Paths in Graphs: Given a graph $G := \langle V, E \rangle$ we call a $n + 1$ -tuple $p = \langle v_0, \dots, v_n \rangle \in V^{n+1}$ a **path** in G iff $(v_{i-1}, v_i) \in E$ for all $1 \leq i \leq n$ and $n > 0$.

- We say that v_i are nodes on p and that v_0 and v_n are **linked** by p .
- v_0 and v_n are called the **start** and **end** of p (write $\text{start}(p)$ and $\text{end}(p)$), the other v_i are called **inner nodes** of p .
- n is called the **length** of p (write $\text{len}(p)$).
- We denote the **set** of paths in G with $\Pi(G)$.


Note

Not all v_i -s in a **path** are necessarily different.

For a graph $G = \langle V, E \rangle$ and a **path** $p = \langle v_1, \dots, v_n \rangle \in V^{n+1}$, write

- $v \in p$, iff $v \in V$ is a vertex on the **path** ($\exists i. v_i = v$)
- $e \in p$, iff $e = (v, \bar{v}) \in E$ is an edge on the **path** ($\exists i. v_i = v \wedge v_{i+1} = \bar{v}$)

Example: Consider $G = \langle V, E \rangle$ with $V = \{a, b, c, d\}$ and $E = \{(a, b), (b, c)\}$:




An example **path** would be $p = \langle a, b, c \rangle$. Here $n = 2$ and $p \in V^3 = \{(a, a, a), (a, b, d), \dots\}$. Obviously, since $n = 2$ so we have 2 edges and 3 nodes.

- nodes a, b, c are on the **path**
- $\text{start}(p) = a$ and $\text{end}(p) = c$
- b is an inner node, and $\text{len}(p) = 2$
- $e = (a, b), \bar{e} = (b, c) \in p$
- $\Pi(G) = \{(a, b), (b, c), (a, b, c)\}$

Def C.10. Cyclic Graphs: Given a **directed graph** $G = \langle V, E \rangle$, a **path** p is called **cyclic** iff $\text{start}(p) = \text{end}(p)$. A cycle $\langle v_0, \dots, v_n \rangle$ is called **simple**, iff $v_i \neq v_j$ for $1 \leq i, j \leq n$ with $i \neq j$ (all inner nodes are distinct).

Def C.11. DAG: A **directed graph** with no **cycles** is called **directed acyclic graph (DAG)**.

Example: $p = \langle a, b, c, a \rangle$ is simple **cyclic** path. $p = \langle a, b, a, b, a \rangle$ is a **cyclic** path that is not simple.



Def C.12. Node Depth: Let $G = \langle V, E \rangle$ be a **directed graph**, then the depth $\text{dp}(v)$ of a vertex $v \in V$ is defined to be 0, iff v is a source of G and the **supremum** $\sup(\{\text{len}(p) \mid \text{indeg}(\text{start}(p)) = 0 \wedge \text{end}(p) = v\})$ otherwise, i.e. the length of the longest **path** from a source of G to v (can be infinite).

Def C.13. Graph Depth: Given a **directed graph** $G = \langle V, E \rangle$, the **depth** ($\text{dp}(G)$) of G is defined as the **supremum** $\sup(\{\text{len}(p) \mid p \in \Pi(G)\})$, i.e. the maximal path length in G .

Def C.14. Tree: A tree is a **DAG** $G = \langle V, E \rangle$ such that

- There is exactly one **initial node** $v_r \in V$ (called the **root**)
- All nodes but the root have **indegree** of 1.


We call v the **parent** of w , iff $(v, w) \in E$ (w is a child of v). We call a node v a **leaf** of G , iff it is **terminal**, i.e. if it does not have children.

Lemma C.1. For any node $v \in V$ except the root v_r , there is exactly one **path** $p \in \Pi(G)$ with $\text{start}(p) = v_r$ and $\text{end}(p) = v$.

Def C.15. Non-Descendants: Let $\langle X, E \rangle$ be a **DAG**, $x \in X$, and E^* the **reflexive transitive closure** of E . The **non-descendants** of x are the **elements** of the **set** $\text{NonDesc}(x) := \{y \mid (x, y) \notin E^*\} \setminus \text{Parents}(x)$

Example: Consider the following **DAG** $G = \langle X, E \rangle$ where:

- $X = \{A, B, C, D, E\}$
- $E = \{(A, B), (A, C), (B, D), (C, D), (C, E)\}$
- $E^* = \{(A, A), (B, B), (C, C), (D, D), (E, E), (A, B), (A, C), (A, D), (A, E), (B, D), (C, D), (C, E)\}$



Let's find the **non-descendants** of node B :

- First, we find all y such that $(B, y) \in E^*$. Looking at E^* , the only pairs starting with B are (B, B) and (B, D) . Thus, the reachable nodes from B are $\{B, D\}$.
- The set of nodes y where $(B, y) \notin E^*$ is everything else: $\{A, C, E\}$.
- The parents of B is the set $\text{Parents}(B) = \{A\}$.
- By the definition, we exclude the parents: $\text{NonDesc}(B) = \{A, C, E\} \setminus \{A\} = \{C, E\}$.

C.2 Complexity Theory

Def C.16. Algorithm: An algorithm is a series of instructions to control a (computation) process.

Def C.17. Computational Problem: A computational problem is a triple $\langle I, S, \sigma \rangle$, where I is a **set** of inputs, S is a **set** of solutions, and $\sigma : I \rightarrow \mathcal{P}(S)$ assigns to each input the set of valid solutions. The problem is called **computable** if there exists an algorithm $\alpha : I \rightarrow S$ such that for every $i \in I$, $\alpha(i) \in \sigma(i)$ and $\alpha(i)$ is defined, i.e. the **algorithm** always **halts**.

Def C.18. Partially Computable Problem: A **computational problem** $\langle I, S, \sigma \rangle$ is called **partially computable** if there exists a **partial algorithm** $\alpha : I \rightarrow S$ such that for every $i \in I$, if $\alpha(i)$ is defined then $\alpha(i) \in \sigma(i)$. In this case α may diverge on some inputs.

Def C.19. Resource: A resource is a measurable quantity consumed by an **algorithm** during execution. The most common resources are **time** (number of computation steps) and **space** (amount of memory used).

Def C.20. Complexity Theory: (Computational) complexity theory focuses on classifying **computational problems** according to their resource usage, and relating these classes to each other.

Def C.21. Complexity Measure: A complexity measure is a mapping $m : I \rightarrow \mathbb{N}$ that assigns to each input $i \in I$ a natural number describing the amount of some **resource** consumed by an **algorithm** on input i .

Def C.22. Input Size: The size of an input $i \in I$, denoted $|i|$, is a natural number measuring the length of the encoding of i (for instance, the number of bits required to represent i).

Def C.23. Asymptotic Analysis: Asymptotic analysis studies the growth of a **complexity measure** as the **input size** n tends to infinity. It abstracts from constant factors and lower-order terms to capture the long-term behavior of an **algorithm**.

Def C.24. Running Time: If an **algorithm** α terminates within $t(n)$ steps on every input of **size** n , then we say that α has **running time** $T(\alpha) := t$.

Def C.25. Time Complexity: Let $S \subseteq \mathbb{N} \rightarrow \mathbb{N}$ be a **set** of natural number **functions**. An **algorithm** α with **running time** $T(\alpha) = t$ is said to have **time complexity** in S (written $T(\alpha) \in S$, or colloquially $T(\alpha) = S$) if $t \in S$.

Def C.26. Space Complexity: Let $S \subseteq \mathbb{N} \rightarrow \mathbb{N}$ be a **set** of natural number **functions**. An **algorithm** α is said to have **space complexity** in S if, on every input of **size** n , α uses at most $s(n)$ units of memory for some $s \in S$.

Def C.27. Best-Case Complexity: The best-case complexity of an **algorithm** α is the function

$$T_{\alpha}^{\text{best}}(n) = \min\{\text{running time of } \alpha(i) \mid i \in I, |i| = n\}.$$

It describes the minimal resource usage among all inputs of size n .

Def C.28. Average-Case Complexity: The average-case complexity of an **algorithm** α is the expected **resource** usage over all inputs of size n , with respect to some probability distribution D on I . Formally,

$$T_{\alpha}^{\text{avg}}(n) = \mathbb{E}_{i \sim D, |i|=n}[\text{running time of } \alpha(i)].$$

Def C.29. Worst-Case Complexity: The worst-case complexity of an **algorithm** α is the function

$$T_{\alpha}^{\text{worst}}(n) = \max\{\text{running time of } \alpha(i) \mid i \in I, |i| = n\}.$$

It describes the maximal **resource** usage among all inputs of size n .

Note

We are mostly interested in **worst-case complexity**.

Def C.30. Big- $\mathcal{O}$ Notation: In practice, **asymptotic analysis** is most often applied to the **worst-case complexity**. Let $f, g : \mathbb{N} \rightarrow \mathbb{R}_{\geq 0}$. We say that f is **asymptotically bounded** by g , written $f \in \mathcal{O}(g)$ or $f \leq_a g$, if there exist constants $c > 0$ and $n_0 \in \mathbb{N}$ such that

$$\forall n \geq n_0 : f(n) \leq c \cdot g(n).$$

Equivalently, the **Landau set** $\mathcal{O}(g)$ is the set of all such **functions**:

$$\mathcal{O}(g) = \{f \mid \exists c > 0, \exists n_0 \in \mathbb{N}, \forall n \geq n_0 : f(n) \leq c \cdot g(n)\}.$$

Example: Let $f(n) = 3n^2 + 2n + 7$ and $g(n) = n^2$. We want to show that $f \in \mathcal{O}(g)$.

By the [definition of Big- \$\mathcal{O}\$](#) , this means we must find constants $c > 0$ and $n_0 \in \mathbb{N}$ such that

$$\forall n \geq n_0 : f(n) \leq c \cdot g(n).$$

Let's take $n_0 = 1$, then $\forall n \geq 1$ we have:

$$2n \leq 2n^2 \quad \text{and} \quad 7 \leq 7n^2.$$

Therefore,

$$f(n) = 3n^2 + 2n + 7 \leq 3n^2 + 2n^2 + 7n^2 = 12n^2.$$

Here we see that $f(n) \leq 12 \cdot g(n)$ for all $n \geq 1$. Thus the definition holds with $c = 12$ and $n_0 = 1$. We conclude that $f(n) \in \mathcal{O}(n^2)$.

Lemma C.2. Growth Ranking. For $\bar{k} > 2$ and $k > 1$ we have:

$$\mathcal{O}(1) \subset \mathcal{O}(\log_2(n)) \subset \mathcal{O}(n) \subset \mathcal{O}(n^2) \subset \mathcal{O}(n^{\bar{k}}) \subset \mathcal{O}(k^n)$$

The following [sets](#) are often used for S in $T(\alpha)$:

Landau set	class name	rank	Landau set	class name	rank
$\mathcal{O}(1)$	constant	1	$\mathcal{O}(n^2)$	quadratic	4
$\mathcal{O}(\log_2 n)$	logarithmic	2	$\mathcal{O}(n^k)$	polynomial	5
$\mathcal{O}(n)$	linear	3	$\mathcal{O}(k^n)$	exponential	6

Table 6: Common complexity classes

Computing with [Landau sets](#) is quite simple:

- $\mathcal{O}(c \cdot f) = \mathcal{O}(f)$ for any constant $c \in \mathbb{N}$ (*drop constant factors*)
- $\mathcal{O}(f) \subseteq \mathcal{O}(g) \Rightarrow \mathcal{O}(f + g) = \mathcal{O}(g)$ (*drop low-complexity summands*)
- $\mathcal{O}(f \cdot g) = \mathcal{O}(f) \cdot \mathcal{O}(g)$ (*distribute over products*)

Example: $\mathcal{O}(4n^3 + 3n + 7^{1000n}) = \mathcal{O}(2^n)$ (*exponential*)

Note

$$\mathcal{O}(n) \subset \mathcal{O}(n \cdot \log_2(n)) \subset \mathcal{O}(n^2) \text{ (linearithmic)}$$

Note

Updated Growth Ranking & Common Classes

To account for all algorithmic patterns discussed (such as nested loops with logarithmic guards), we refine the growth hierarchy. For $a > 2$ and $b > 1$, the asymptotic ranking is :

$$\mathcal{O}(1) \subset \mathcal{O}(\log_2 n) \subset \mathcal{O}(n) \subset \mathcal{O}(n \log_2 n) \subset \mathcal{O}(n^2) \subset \mathcal{O}(n^a) \subset \mathcal{O}(b^n)$$

Determining the Time/Space Complexity of Algorithms

Assume we have

- Γ : a mapping that tells us the complexity of variables.
- α : a program fragment.
- $T_\Gamma(\alpha)$: the [time complexity](#) of α under Γ .
- $C_\Gamma(\alpha)$: the context update, i.e. what new cost α introduces.

We use [induction](#) on program structure:

1. Constant

If $\alpha = \delta$ (where δ is a literal, like 5 or `true`), then $T_\Gamma(\alpha) \in \mathcal{O}(1)$

2. Variable

If $\alpha = v$ and $v \in \text{dom}(\Gamma)$, then $T_\Gamma(\alpha) \in \mathcal{O}(\Gamma(v))$

3. Function Call

If $\alpha = \varphi(\psi)$, with $T_\Gamma(\alpha) \in \mathcal{O}(f)$, and $T_{\Gamma \cup C_\Gamma(\varphi)}(\psi) \in \mathcal{O}(g)$, then $T_\Gamma(\alpha) \in \mathcal{O}(f \circ g)$

4. Assignment

If $\alpha = v := \varphi$, with $T_\Gamma(\varphi) \in S$, then $T_\Gamma(\alpha) \in S$, $C_\Gamma(\alpha) = \Gamma \cup (v, S)$

5. Composition (sequence)

If $\alpha = \varphi; \psi$, with $T_\Gamma(\varphi) \in P$, and $T_{\Gamma \cup C_\Gamma(\varphi)}(\psi) \in Q$, then $T_\Gamma(\alpha) \in \max\{P, Q\}$

6. Branching

If $\alpha = \text{if } \gamma \text{ then } \varphi \text{ else } \psi \text{ end}$, with $T_\Gamma(\gamma) \in C$, $T_{\Gamma \cup C_\Gamma(\gamma)}(\varphi) \in P$, and $T_{\Gamma \cup C_\Gamma(\gamma)}(\psi) \in Q$, then $T_\Gamma(\alpha) \in \max\{C, P, Q\}$

7. Loop

If $\alpha = \text{while } \gamma \text{ do } \varphi \text{ end}$, with $T_\Gamma(\gamma) \in \mathcal{O}(f)$, and $T_{\Gamma \cup C_\Gamma(\gamma)}(\varphi) \in \mathcal{O}(g)$, then $T_\Gamma(\alpha) \in \mathcal{O}(f(n) \cdot g(n))$

Note

By default, the overall time complexity is $T(\alpha) = T_\phi(\alpha)$

Sequence Example

```
x := readArray(n)    -- O(n)
y := sum(x)           -- O(n)
max(O(n), O(n)) = O(n)
```

Conditional Example

```
if isSorted(x) then   -- O(n)
  binarySearch(x,k)   -- O(log n)
else
  linearSearch(x,k)   -- O(n)
end
max(O(n), O(log(n), O(n))) = O(n)
```

While loop example

```
i := 1
while i <= n do
  i := i + 1
end
```

Guard (loop condition check) runs $\mathcal{O}(n)$ times, body is $\mathcal{O}(1)$. Total $\mathcal{O}(n \cdot 1) = \mathcal{O}(n)$

Function Call Example

Suppose `buildArray(n)` $\in \mathcal{O}(n^2)$ returns an array of size n^2 and `sort(m)` $\in \mathcal{O}(m \cdot \log(m))$

```
sort(buildArray(n))
O((n^2) · log(n^2)) = O(n^2 · log(n))
```

Nested For Loop Example

```
for i = 1 to n do      --  $\mathcal{O}(n)$ 
  for j = 1 to n do    --  $\mathcal{O}(n)$ 
    x := x + 1          --  $\mathcal{O}(1)$ 
```

Inner loop: $\mathcal{O}(n \cdot 1)$, outer loop: $\mathcal{O}(n \cdot n \cdot 1)$, total: $\mathcal{O}(n^2)$

Nested (one logarithmic loop)

```
for i = 1 to n do
  k := 1
  while k < n do
    k := 2 * k
```

Inner loop starts with $k = 1$, each iteration k doubles and the loop stops when $k \geq n$. How many times can we double before reaching n ? Solve $2^t \geq n \rightsquigarrow t = \lceil \log_2(n) \rceil$. So the inner loop runs $\mathcal{O}(\log(n))$ iterations. Each iteration is $\mathcal{O}(1)$, therefore total $\mathcal{O}(\log(n))$.

Outer loop runs exactly n times with each time the cost is $\mathcal{O}(\log(n))$. Total cost is $n \cdot \mathcal{O}(\log(n)) = \mathcal{O}(n \cdot \log(n))$

Note

Take home message: not every time you see two nested loops you assume $\mathcal{O}(n^2)$

Here is another example: assume $\llbracket \text{body} \rrbracket \in \mathcal{O}(1)$:

```
i, j := 0
while i <= n-1 do
  i := i+1
  while j <= 10 do
    j := j+1
     $\llbracket \text{body} \rrbracket$ 
  end
end
```

Outer Loop: i increments by 1 each time until $i > n - 1$ we stop, so it runs n times.

Inner Loop: the condition is $j \leq 10$, but note that j is initialized once before the loop and not reset for each i ! That means, first outer iteration $j = 0$, runs until $j = 11$, i.e. 11 iterations. After that, $j = 11$ always and the condition $j \leq 10$ is false, so the inner loop never runs again. Inner loop total cost is $\mathcal{O}(11) = \mathcal{O}(1)$

Therefore, all in all we have linear time ($\mathcal{O}(n)$).

C.3 Formal Languages and Grammars

Def C.31. Alphabet: An alphabet is a [finite set](#); we call each element $a \in A$ a **character**, and an n -tuple $s \in A^n$ a **string** (of length n over A).

Note

We often write a string $\langle c_1, \dots, c_n \rangle$ as " $c_1 \dots c_n$ ", for instance "**abc**" for $\langle a, b, c \rangle$

Example: Let $A = \{0, 1\}$ be an [alphabet](#)

- "0", "1" are strings of length 1 ($A^1 = \{0, 1\}$)
- "00", "01", "10", "11" are strings of length 2 in (A^2)
- "001" $\in A^3$

Def C.32. Empty String: Let A be an [alphabet](#), then $A^0 = \{\langle \rangle\}$, where $\langle \rangle$ is the unique 0-tuple. We consider $\langle \rangle$ as the string of length 0 and call it the **empty string** and denote it with ϵ .

Note

$\{1, 2, 3\} = \{3, 2, 1\}$ BUT $\langle 1, 2, 3 \rangle \neq \langle 3, 2, 1 \rangle$ ([sets](#) vs strings)

Def C.33. String Length: Given a string s , we denote its length with $|s|$.

Def C.34. String Concatenation: The concatenation $\text{conc}(s, t)$ of two strings $s = \langle s_1, \dots, s_n \rangle \in A^n$ and $t = \langle t_1, \dots, t_m \rangle \in A^m$ is defined as $s + t$ or simply st

Example: $\text{conc}(\text{"text"}, \text{"book"}) = \text{"text"} + \text{"book"} = \text{"textbook"}$.

Def C.35. Kleene Plus and Kleene Star: Let A be an [alphabet](#), then we define the [sets](#):

- $A^+ := \bigcup_{i \in \mathbb{N}^+} A^i$ (*nonempty strings*), and
- $A^* := A^+ \cup \{\epsilon\}$ (*strings*).

Def C.36. Formal Language: A [set](#) $L \subseteq A^*$ is called a *formal language* over A .

Example: If $A = \{a, b\}$ then:

$$A^* = \bigcup_{n=0}^{\infty} A^n$$

where $A^0 = \epsilon$, $A^1 = \{a, b\}$, $A^2 = \{a, b, ab, ba\}$, and so on ...

$$\begin{aligned} A^* &= \{\epsilon, a, b, aa, ab, ba, bb, aaa, aab, aba, abb, \dots\} \\ A^+ &= \{a, b, aa, ab, ba, bb, aaa, aab, aba, abb, \dots\} \end{aligned}$$

A [formal language](#) might be $L = \{a, b, aab, bbb, \dots\}$

Def C.37. Repeating Chars: We use $c^{[n]}$ for the string that consists of the character c repeated n times.

Examples:

- Let $A = \{a, b\}$, $a^{[5]} = \text{"aaaaa"}$.
- The [set](#) $M := \{ba^{[n]} \mid n \in \mathbb{N}\}$ of strings that start with character b followed by an arbitrary number of a 's is a [formal language](#) over $A = \{a, b\}$

Note

A [formal language](#) might be infinite and even undecidable even if the [alphabet](#) A is [finite](#). For example, let $A = \{a, b\}$ then the [formal language](#) $L = \{a, ab, bba\}$ is [finite](#) but the [formal language](#) $L = \{a^{[n]} \mid n \in \mathbb{N}\}$ is infinite.

Since we cannot list all strings in a [formal language](#) L because there are infinitely many, we need a **finite description** that can generate all strings in L . We need **Grammars**.

Def C.38. Phrase Structure Grammar: A phrase structure grammar (*type 0 grammar, unrestricted grammar, grammar*) is a tuple $\langle N, \Sigma, P, S \rangle$ where

- N is a **finite set** of **nonterminal symbols**,
- Σ is a **finite set** of **terminal symbols**. (members of $N \cup \Sigma$ are called symbols).
- P is a **finite set** of **production rules**: pairs $p := h \rightarrow b$ (also written as $h \Rightarrow b$), where $h \in (\Sigma \cup N)^* N (\Sigma \cup N)^*$ and $b \in (\Sigma \cup N)^*$. The string h is called the **head** of p and b the **body**.
- $S \in N$ is a distinguished symbol called the **start symbol** (also **sentence symbol**).

The **sets** N and Σ are assumed to be **disjoint**. Any word $w \in \Sigma^*$ is called a **terminal word**.

Note

Production rules map strings with at least one **nonterminal** to arbitrary other strings

Notation: If we have n **production rule** $h \rightarrow b_i$ sharing a head, we often write $h \rightarrow b_1 \mid \dots \mid b_n$ instead.

Example: A simple **grammar** for english sentences:

$$\begin{aligned} S &::= NP \text{ Vi} \\ NP &::= \text{Article } N \\ \text{Article} &::= \text{the} \mid \text{a} \mid \text{an} \\ N &::= \text{dog} \mid \text{teacher} \mid \dots \\ \text{Vi} &::= \text{sleeps} \mid \text{smells} \mid \dots \end{aligned}$$

Here S is the **start symbol**, NP , Article , N , and Vi are **nonterminals**, the , a , dog , smells , ... are **terminals**. Let's generate a valid sentence from the **grammar**:

NP		Vi
$\text{Article } N$		Vi
$\text{the} \quad N$		Vi
$\text{the} \quad \text{dog}$		Vi
$\text{the} \quad \text{dog}$	sleeps	

Def C.39. Lexical: A **production rule** whose head is a single **nonterminal** and whose body consists of a single **terminal** is called **lexical** or a **lexical insertion rule**.

Example: Both of the following **rules** are **lexical**

$$\begin{aligned} \text{Noun} &\rightarrow \text{"dog"} \\ \text{Verb} &\rightarrow \text{"run"} \end{aligned}$$

Def C.40. Lexicon of a Grammar: The **subset** of **lexical rules** of a **grammar** G is called the **lexicon** of G

Def C.41. Vocabulary: The **set** of body symbols are called the vocabulary (or the alphabet).

Def C.42. Lexical Categories of a Grammar: The **nonterminals** that appear in the heads of **lexical** rules are called lexical categories of the **grammar**.

Def C.43. Structural Rules: The non **lexicon production rules** are called structural.

Def C.44. Phrasal categories: The **nonterminals** in the heads of **production rules** that expand into other **nonterminals** are called phrasal (syntactic) categories.

Def C.45. G -Derivation: Given a **phrase structure grammar** $G := \{N, \Sigma, P, S\}$, we say G **derives** $t \in (\Sigma \cup N)^*$ from $s \in (\Sigma \cup N)^*$ in one step, iff there is a **production rule** $p \in P$ with $p = h \rightarrow b$ and there are $u, v \in (\Sigma \cup N)^*$, such that $s = uhv$ and $t = ubv$. We write $s \rightarrow_G^p t$ (or $s \rightarrow_G t$ if p is clear from the context) and use $\rightarrow_G^*$ for the **reflexive transitive** closure of $\rightarrow_G$. We call $s \rightarrow_G^* t$ a G -derivation of t from s .

Let's unravel that. Take the following **grammar**:

- $N = \{S\}$
- $\Sigma = \{a, b\}$
- $P = \{S \rightarrow aSb, S \rightarrow ab\}$
- $S = S$

The following derivations holds:

- $S \xrightarrow{p}_G \text{aSb}$ where $p : S \rightarrow \text{aSb}$ (*one step*)
- $\text{aSb} \xrightarrow{q}_G \text{aabb}$ where $q : S \rightarrow \text{ab}$ (*one step*)
- $S \xrightarrow{S \rightarrow \text{aSb}}_G \text{aSb} \xrightarrow{S \rightarrow \text{ab}}_G \text{aabb}$ (*multiple steps*). So, $S \xrightarrow{*}_G \text{aabb}$

Def C.46. Sentential Form: Given a **phrase structure grammar** $G := \{N, \Sigma, P, S\}$, we say that $s \in (\Sigma \cup N)^*$ is a **sentential form** of G , iff $S \xrightarrow{*}_G s$

Def C.47. Sentence: A **sentential form** that does not contain **nonterminals** is called a sentence of G , we also say that G **accepts** s . We say that G **rejects** s , iff it is not a sentence of G .

Examples

1. "Article teacher Vi" is a **sentential form**

$$\begin{aligned} S &\rightarrow_G \text{NP Vi} \\ &\rightarrow_G \text{Article N Vi} \\ &\rightarrow_G \text{Article teacher Vi} \end{aligned}$$

2. "the teacher sleeps" is a **sentence**

$$\begin{aligned} S &\xrightarrow{*}_G \text{Article teacher Vi} \\ &\rightarrow_G \text{the teacher Vi} \\ &\rightarrow_G \text{the teacher sleeps} \end{aligned}$$

Def C.48. Language Generation: The language $L(G)$ of G is the **set** of its **sentences**. We say that $L(G)$ is **generated** by G .

Def C.49. Equivalent Grammars: We call two **grammars** **equivalent**, iff they have the same **languages**.

Def C.50. Universal Grammar: A **grammar** G is said to be **universal** if $L(G) = \Sigma^*$

Def C.51. Syntactic Analysis: Syntactic analysis (parsing / syntax analysis) is the process of analyzing a string of symbols, either in a **formal** or a **natural language** by means of a **grammar**.

Note

The shape of the **grammar** determines the size of its **language**

Def C.52. Context-Sensitive: We call a **grammar** context-sensitive (or type 1), if the bodies of **production rules** have no less symbols than the heads.

Example: The **language** $\{a^{[n]}b^{[n]}c^{[n]}\}$ is **accepted** by

$$\begin{aligned} S &\rightarrow a b c \mid A \\ A &\rightarrow a A B c \mid a b c \\ c B &\rightarrow B c \\ b B &\rightarrow b b \end{aligned}$$

Def C.53. Context-Free: We call a **grammar** context-free (or type 2), iff the heads have exactly one symbol.

Example: The **language** $\{a^{[n]}b^{[n]}\}$ is **accepted** by $S \rightarrow a S b \mid \epsilon$

Def C.54. Regular: We call a **grammar** regular (or type 3, or REG), if additionally the bodies are empty or consist of a **nonterminal**, optionally followed by a **terminal symbol**.

Example: The **language** $\{a^{[n]}\}$ is **accepted** by $S \rightarrow S a \mid \epsilon$

Note

By extension, a **formal language** L is called *context-sensitive* / *context-free* / *regular*, iff it is the **language** of a respective **grammar**. **Context-free grammars** are sometimes **CFGs** and context-free languages **CFLs**.

Observation: **Natural languages** are probably **context sensitive** but parsable in real time!